

# Carnegie Mellon University

Department of Statistics & Data Science

## Thesis Proposal

Draft as of March 29, 2026

Simulation-based Time Series Analysis:  
Testing for Structure, Likelihood-free Estimation,  
and Prediction with Transformers using Synthetic Data

Michael Wieck-Sosa

April 3, 2026

**Committee:**

Aaditya Ramdas, Advisor

Cosma Shalizi

Yuting Wei, University of Pennsylvania

Wei Biao Wu, University of Chicago

**Abstract**

Technological advances have enabled the collection of large-scale time series data in many fields. Traditional time series methods often struggle to capture the complex dynamics present in such data because of their restrictive assumptions, such as stationarity, linear dynamics, or Gaussian noise. Consequently, there is growing demand for more flexible time series methods. At the same time, recent developments in machine learning, notably the transformer architecture, have demonstrated remarkable success in sequence modeling.

In this thesis proposal, we present simulation-based methods for problems in multivariate time series analysis. In the first part, we introduce a bootstrap-based framework for conditional independence testing with nonstationary time series that utilizes nonlinear regression. In the second part, we develop a framework for simulation-based parameter estimation based on matching random features, which we apply to nonlinear non-Gaussian state-space models, discrete-time dynamical systems and ODEs with observational noise, discretized Lévy-driven SDEs, and more. In the third part, we introduce a simulation-and-regression framework for time series prediction with transformers that accounts for parameter uncertainty.

## Contents

|          |                                                       |           |
|----------|-------------------------------------------------------|-----------|
| <b>1</b> | <b>Introduction</b>                                   | <b>4</b>  |
| 1.1      | Thesis Agenda . . . . .                               | 4         |
| <b>2</b> | <b>Conditional Independence Testing</b>               | <b>5</b>  |
| 2.1      | Overview . . . . .                                    | 5         |
| 2.2      | Setting . . . . .                                     | 5         |
| 2.3      | Time-varying Regression and Error Processes . . . . . | 5         |
| 2.4      | Testing Procedure . . . . .                           | 6         |
| <b>3</b> | <b>Simulation-based Estimation</b>                    | <b>8</b>  |
| 3.1      | Overview . . . . .                                    | 8         |
| 3.2      | Notation and setting . . . . .                        | 8         |
| 3.3      | Random Fourier features . . . . .                     | 9         |
| 3.4      | Estimators . . . . .                                  | 10        |
| 3.5      | Examples . . . . .                                    | 11        |
| <b>4</b> | <b>Prediction with Transformers via Simulation</b>    | <b>16</b> |
| 4.1      | Background . . . . .                                  | 16        |
| 4.2      | State-space Modeling . . . . .                        | 18        |
| 4.3      | Confidence Distributions . . . . .                    | 20        |
| 4.4      | Proposed Work . . . . .                               | 23        |
| 4.5      | Algorithms . . . . .                                  | 25        |
| <b>5</b> | <b>Timeline</b>                                       | <b>28</b> |

# 1 Introduction

## 1.1 Thesis Agenda

We provide a concise overview of the organization of this thesis proposal, along with some motivations for each topic.

**I. Conditional Independence Testing (Section 2).** Given nonstationary time series  $X_{1:n}$ ,  $Y_{1:n}$ ,  $Z_{1:n}$ , we develop a nonlinear regression-based test of the null hypothesis that  $X_t$  is conditionally independent of  $Y_t$  given  $Z_t$  at all times  $t$ , with any desired leads and lags incorporated into these variables. Tests of conditional independence are used in many problems, such as variable selection, causal discovery, the validation of theory-based causal graphs, and the evaluation of predictor fairness; see Hardt et al. (2016).

**II. Simulation-based Estimation (Section 3).** Next, we introduce a simulation-based framework for estimating the unknown parameter in a parametric statistical model for a time series  $X_{1:n}$ . Such methods are desirable for models where the likelihood is analytically or computationally intractable, but simulating from the model is easy. For example, nonlinear non-Gaussian state-space models, discrete-time dynamical systems and ODEs with observational noise, and discretized Lévy-driven SDEs.

**III. Prediction with Transformers (Section 4).** Lastly, we develop a simulation-and-regression framework for training transformers for the tasks of forecasting and smoothing. Forecasting aims to characterize the future values of a stochastic process given the current information. It underpins monetary policy decisions in economics, portfolio optimization in finance, public health policy decisions during epidemics, and early warning systems for extreme weather events. Smoothing aims to characterize the past values of a latent stochastic process given the current information. Applications include estimating sentiment in the social sciences, stochastic volatility in finance, infectivity in epidemiology, and atmospheric variables in the Earth Sciences.

## 2 Conditional Independence Testing

### 2.1 Overview

In this part of the thesis, we introduce a framework for conditional independence testing with nonstationary time series. To the best of our knowledge, this is the first conditional independence testing framework for nonstationary time series outside of the linear-Gaussian setting. In particular, we use a theoretical framework for nonstationary time series based on the foundational work of Wu (2005); Zhou and Wu (2009); Dahlhaus et al. (2019). We propose a bootstrap-based testing procedure, which we justify with a distribution-uniform version of the strong Gaussian approximation from Mies and Steland (2023). We prove that our tests have Type I error control, *uniformly* over a large collection of null distributions. The main ideas are presented here, and we refer readers to the paper (Wieck-Sosa et al., 2025) for more details.

### 2.2 Setting

Let  $X_{1:n}$ ,  $Y_{1:n}$ ,  $Z_{1:n}$  be time series taking values in  $\mathbb{R}^{d_X}$ ,  $\mathbb{R}^{d_Y}$ ,  $\mathbb{R}^{d_Z}$ , respectively, for some  $n, d_X, d_Y, d_Z \in \mathbb{N}$ . To simplify the notation, let these variables incorporate any desired leads and lags. The only requirement is that the conditioning set  $Z_t$  is known at time  $t$ . In this section, we present a test for the null hypothesis that

$$X_t \perp\!\!\!\perp Y_t \mid Z_t \text{ for all } t \in [n]. \quad (1)$$

We allow the time series to have nonlinear temporal dynamics and general forms of nonstationarity. The paper (Wieck-Sosa et al., 2025) includes the details about the asymptotic framework, the assumptions about the nonstationarity and decay of temporal dependence using the physical dependence measure of Wu (2005), and the rates at which the number of dimensions and time-offsets can grow with  $n$ .

### 2.3 Time-varying Regression and Error Processes

We can always decompose

$$X_t = f_t(Z_t) + \varepsilon_t, \quad Y_t = g_t(Z_t) + \xi_t, \quad (2)$$

where  $f_t(z) = \mathbb{E}(X_t \mid Z_t = z)$  and  $g_t(z) = \mathbb{E}(Y_t \mid Z_t = z)$  are the regression functions, which may vary over time. The error processes  $\varepsilon_{1:n}$  and  $\xi_{1:n}$  can

also be nonstationary and temporally dependent. Denote the vectorized outer product of the errors at time  $t$  by

$$R_t = \text{vec}(\varepsilon_t \xi_t^\top). \quad (3)$$

Next, let  $\hat{f}_t$  and  $\hat{g}_t$  be estimates of  $f_t$  and  $g_t$ , respectively, and let

$$\hat{\varepsilon}_t = X_t - \hat{f}_t(Z_t), \quad \hat{\xi}_t = Y_t - \hat{g}_t(Z_t), \quad (4)$$

be the corresponding residuals. Denote the vectorized outer product of these residuals at time  $t$  by

$$\hat{R}_t = \text{vec}(\hat{\varepsilon}_t \hat{\xi}_t^\top). \quad (5)$$

Crucially, our testing framework can be used with *any regression estimator*. The required rates of convergence for the regression estimators are discussed in the paper (Wieck-Sosa et al., 2025). In our experiments, we use the sieve estimator from Ding and Zhou (2021), though there are many other regression estimators for nonstationary time series to choose from (Vogt, 2012; Zhang and Wu, 2015; Yousuf and Ng, 2021; Chen et al., 2022; Kurisu et al., 2025). Our test possesses a property known as *rate double robustness*, which means that we only require modest convergence rates for the products of the estimation errors, rather than for each estimation error individually.

## 2.4 Testing Procedure

Our proposed testing procedure can be viewed as an extension of the *generalized covariance measure* (GCM) test from Shah and Peters (2020) from the iid setting to the nonstationary time series setting, so we refer to it as the *dynamic generalized covariance measure* (dGCM) test. The core idea is that, under the null hypothesis of conditional independence from (1), the covariance of the errors from (2) will be zero at all times. Note that these covariances can be zero at all times, even under alternatives in which the corresponding conditional dependencies always hold. Consequently, we can only hope to have power against alternatives in which these covariances are non-zero for at least *some* times.

Roughly speaking, our test statistic uses the empirical covariances among the residuals from (4) to try to detect non-zero covariances among the errors from (2), for at least *some* times. To estimate the desired quantile of the test statistic, we use a bootstrap procedure justified by a *distribution-uniform* version of the strong Gaussian approximation from Mies and Steland (2023) applied to the process of error products  $R_t$  from (3). This is possible because, under the null hypothesis from (1), the error products  $R_t$  from (3) will

have mean zero at all times  $t$ . The approximating Gaussian process has a time-varying covariance structure, which we estimate using a rolling-window approach; see the paper (Wieck-Sosa et al., 2025) for how to select the window size. We prove that the following procedure has Type I error control, uniformly over a large collection of null distributions.

### Dynamic Generalized Covariance Measure (dGCM) Test.

- **Inputs:** Data  $X_{1:n}, Y_{1:n}, Z_{1:n}$ , significance level  $\alpha$ , number of simulations  $s, p \in [2, \infty]$  for norm in test statistic, and method for selecting the window size  $L$ .
- Regress  $X_t$  on  $Z_t, t \in [n]$ , and obtain residuals  $\hat{\varepsilon}_t, t \in [n]$ .
- Regress  $Y_t$  on  $Z_t, t \in [n]$ , and obtain residuals  $\hat{\xi}_t, t \in [n]$ .
- Obtain residual products  $\hat{R}_t = \text{vec}(\hat{\varepsilon}_t \hat{\xi}_t^\top), t \in [n]$ .
- Calculate test statistic on time series of residual products

$$T(\hat{R}_{1:n}) = \max_{j \in [n]} \left\| \frac{1}{\sqrt{n}} \sum_{t \leq j} \hat{R}_t \right\|_p.$$

- Select the window size  $L$  for covariance estimation using  $\hat{R}_{1:n}$ .
- For  $t \in [n]$ , obtain estimates of the time-varying covariances as

$$\hat{\Sigma}_t = \frac{1}{L} \left( \sum_{j=t-L+1}^t \hat{R}_j \right) \left( \sum_{j=t-L+1}^t \hat{R}_j \right)^\top.$$

- For each  $r \in [s]$  and  $t \in [n]$ , simulate *independent* Gaussians

$$\tilde{R}_t^{(r)} \sim N(0, \hat{\Sigma}_t).$$

- For each  $r \in [s]$ , calculate test statistic on simulated Gaussian process

$$T(\tilde{R}_{1:n}^{(r)}) = \max_{j \in [n]} \left\| \frac{1}{\sqrt{n}} \sum_{t \leq j} \tilde{R}_t^{(r)} \right\|_p.$$

- Calculate the  $1 - \alpha$  empirical quantile  $\hat{q}_{1-\alpha}^{\text{boot}}$  of the simulated test statistics  $T(\tilde{R}_{1:n}^{(1)}), \dots, T(\tilde{R}_{1:n}^{(s)})$ .
- Reject null hypothesis from (1) at level  $\alpha$  if  $T(\hat{R}_{1:n}) > \hat{q}_{1-\alpha}^{\text{boot}}$ , else retain.

### 3 Simulation-based Estimation

#### 3.1 Overview

In this part of the thesis, we develop a simulation-based estimation framework for parametric models of stochastic processes. Given an  $\mathbb{R}^d$ -valued time series  $X_{1:n}$  generated by such a model, we aim to estimate the unknown parameter  $\theta_0 \in \mathbb{R}^p$  for the model. Unfortunately, the likelihood function is intractable, but simulating from the model with different parameter values is straightforward.

Our core idea is to estimate  $\theta_0$  by tuning the parameter values so that a small number of *random features* of the simulated data closely match those of the observed data. Results from the “geometry from a time series” literature in nonlinear dynamics indicate that models with a  $p$ -dimensional parameter space can typically be identified using only  $2p + 1$  random features, regardless of the dimension of the observations. Our approach avoids both user-selected summary statistics and costly neural network-based procedures for learning them. This leads to a user-friendly method that can be readily applied in many fields, such as ecology and epidemiology.

We introduce two estimators. First, a time-average estimator for processes that are stationary, or at least asymptotically mean stationary. Second, a rolling-window estimator for processes with general forms of nonstationarity. We prove that both estimators are consistent and asymptotically normal under nonrestrictive conditions. Specifically, we use a theoretical framework for nonstationary time series based on Wu (2005); Zhou and Wu (2009); Dahlhaus et al. (2019); Mies and Steland (2023). Our experiments suggest that simulation-based estimation can be made more automatic by matching random features, with little loss in statistical efficiency.

#### 3.2 Notation and setting

For real numbers  $x, y \in \mathbb{R}$ , denote  $x \wedge y = \min(x, y)$  and  $x \vee y = \max(x, y)$ . For a vector  $x \in \mathbb{R}^d$ , denote the  $\ell^p$  norm by  $\|x\|_p$  and the Euclidean norm by  $\|x\| = \|x\|_2$ . For a random vector  $X$ ,  $\|X\|_{\mathcal{L}^q} = (\mathbb{E}\|X\|^q)^{1/q}$  denotes the  $\mathcal{L}^q$  norm of the Euclidean norm. For any natural number  $j \in \mathbb{N}$ , denote  $[j] = \{1, 2, \dots, j\}$ . For a sequence of random vectors  $X_t$ ,  $t = 1, \dots, n$ , each taking values in  $\mathbb{R}^d$ , we denote a subsequence as  $X_{t_1:t_2} = (X_{t_1}, \dots, X_{t_2})^\top$ , which takes values in  $\mathbb{R}^{(t_2-t_1+1) \times d}$  for some  $t_1, t_2 \in [n]$ .

We observe a time series  $X_t^{\text{obs}}$ ,  $t = 1, \dots, n$ , taking values in  $\mathbb{R}^d$  for some fixed dimension  $d \in \mathbb{N}$ . The observed time series is assumed to arise



from a known generative model with an unknown  $p$ -dimensional parameter  $\theta_0 \in \Theta \subset \mathbb{R}^p$ ,  $p \in \mathbb{N}$ . Each value of  $\theta \in \Theta$  determines a law of a stochastic process. For any value of  $\theta$ , we can use the generative model to simulate  $s \in \mathbb{N}$  realizations of a length  $n$  time series  $(X_{1:n}^{(r)}(\theta))_{r \in [s]}$ . When we do not need to refer to a particular realization, we drop the superscript.

### 3.3 Random Fourier features

The central idea is that we should estimate the  $p$ -dimensional parameter  $\theta_0$  by matching the values of a small number of random features of the simulated and observed data. Consider  $k = 2p + 1$  randomly drawn functions  $\varphi_1, \dots, \varphi_k$  from a distribution  $D$  over a suitably nice function class  $\mathcal{F}$ . The  $k$  functions should be sampled independently of one another and of the data. Each function  $\varphi_i : \mathbb{R}^{(m+1) \times d} \rightarrow \mathbb{R}$  takes as input a length  $m + 1$  subsequence of the  $d$ -dimensional time series, where  $m$  is chosen based on the dynamics of the generative model. For instance, if the model is an order  $m^*$  Markov process, then set  $m = m^*$ .

Let  $x_1, \dots, x_{m+1}$  be vectors in  $\mathbb{R}^d$ , and define  $x = (x_1, \dots, x_{m+1})^\top \in \mathbb{R}^{(m+1) \times d}$ . In this paper, we focus on random Fourier features, though our framework can be used with different random features. Specifically, we consider random Fourier features of the form

$$\varphi_i(x) = \cos \left( \sum_{j=1}^{m+1} \Omega_{i,j} \cdot x_j + \alpha_i \right), \quad (6)$$

where  $\Omega_{i,j} \stackrel{\text{iid}}{\sim} N(0, I_d)$  and  $\alpha_i \stackrel{\text{iid}}{\sim} U(-\pi, \pi)$  for  $i = 1, \dots, k$  and  $j = 1, \dots, m+1$ . Denote all  $k$  of these functions by

$$\varphi = (\varphi_1, \dots, \varphi_k), \quad (7)$$

so that  $\varphi : \mathbb{R}^{(m+1) \times d} \rightarrow \mathbb{R}^k$ . Define the  $i$ -th random feature of the observed and  $r$ -th simulated time series at time  $t = m + 1, \dots, n$  as

$$f_{t,i}^{\text{obs}} = \varphi_i(X_{t-m:t}^{\text{obs}}), \quad f_{t,i}^{(r)}(\theta) = \varphi_i(X_{t-m:t}^{(r)}(\theta)).$$

Write all  $k$  random features at time  $t$  as

$$f_t^{\text{obs}} = (f_{t,1}^{\text{obs}}, \dots, f_{t,k}^{\text{obs}}), \quad f_t^{(r)}(\theta) = (f_{t,1}^{(r)}(\theta), \dots, f_{t,k}^{(r)}(\theta)). \quad (8)$$

### 3.4 Estimators

**Time-average estimator.** For processes that are asymptotically mean stationary, it suffices to consider the time-average of  $k = 2p + 1$  random features. Define the observed and simulated time-averages of the random features by

$$F^{\text{obs}} = \frac{1}{n-m} \sum_{t=m+1}^n f_t^{\text{obs}}, \quad \bar{F}^{\text{sim}}(\theta) = \frac{1}{n-m} \sum_{t=m+1}^n \bar{f}_t^{\text{sim}}(\theta), \quad (9)$$

where  $\bar{f}_t^{\text{sim}}(\theta) = \frac{1}{s} \sum_{r=1}^s f_t^{(r)}(\theta)$ . The time-average estimator is given by

$$\hat{\theta}^{\text{TA}} = \underset{\theta \in \Theta}{\operatorname{argmin}} \left\| \hat{Q}_n^{\text{TA}}(\theta) \right\|, \quad (10)$$

where the sample discrepancy function is defined as

$$\hat{Q}_n^{\text{TA}}(\theta) = F^{\text{obs}} - \bar{F}^{\text{sim}}(\theta). \quad (11)$$

In practice, an optimization procedure is used to find the minimizer of  $\left\| \hat{Q}_n^{\text{TA}}(\cdot) \right\|$ .

**Rolling-window estimator.** For many nonstationary processes, a different approach is required because the time-average of the time-varying means of the random features may not converge to a limiting mean. Even when it does converge, the limiting mean may not uniquely identify each  $\theta$ . Thus, we introduce the following estimator.

Define the observed and simulated rolling-window random features at time  $t$  by

$$F_t^{\text{obs}} = \frac{1}{w \wedge t} \sum_{j=(t-w) \vee 1}^t f_j^{\text{obs}}, \quad \bar{F}_t^{\text{sim}}(\theta) = \frac{1}{w \wedge t} \sum_{j=(t-w) \vee 1}^t \bar{f}_j^{\text{sim}}(\theta), \quad (12)$$

where  $\bar{f}_t^{\text{sim}}(\theta) = \frac{1}{s} \sum_{r=1}^s f_t^{(r)}(\theta)$  and  $w \in \mathbb{N}$  is the window size. The rolling-window estimator is given by

$$\hat{\theta}^{\text{RW}} = \underset{\theta \in \Theta}{\operatorname{argmin}} \left| \hat{Q}_n^{\text{RW}}(\theta) \right|, \quad (13)$$

where the sample discrepancy function is defined as

$$\begin{aligned} \hat{Q}_n^{\text{RW}}(\theta) &= \frac{1}{n-m} \sum_{t=m+\tau+L}^n \left[ \left\| F_t^{\text{obs}} - \bar{F}_t^{\text{sim}}(\theta) \right\|^2 + K_t(\theta) \right], \\ K_t(\theta) &= 2 \left( F_{t-L}^{\text{obs}} - \bar{F}_{t-L}^{\text{sim}}(\theta) \right)^\top \left( \left[ f_t^{\text{obs}} - \bar{f}_t^{\text{sim}}(\theta) \right] - \left[ F_{t-L}^{\text{obs}} - \bar{F}_{t-L}^{\text{sim}}(\theta) \right] \right), \end{aligned} \quad (14)$$

where  $\tau \in \mathbb{N}$  is an initial time-offset and  $L \in \mathbb{N}$  is a lag. In practice, an optimization procedure is used to find the minimizer of  $\left| \hat{Q}_n^{\text{RW}}(\cdot) \right|$ . We discuss the selection of  $w$ ,  $\tau$ , and  $L$  in the paper.

**Extensions.** Several extensions are possible. First, both estimators can be generalized by replacing the Euclidean norm with a weighted norm, using a weight matrix given by the inverse of a suitable estimator of the long-run covariance matrix; see Gouriéroux and Monfort (1996). Second, the sample averages used in both estimators can be replaced with different mean estimators. Third, one may extend both estimators to allow for early stopping, so that only a fraction of the time series is used.

### 3.5 Examples

We assess the finite-sample performance of the proposed estimators on several data generating processes. We optimize the objective functions using the differential evolution solver `differential_evolution` from the `scipy.optimize` module in the SciPy Python library, although our experiments indicate that other optimizers perform comparably. For the ODE examples, we used the Runge–Kutta 5(4) solver RK45 from the `scipy.integrate` module in SciPy. Each density plot is constructed from 1,000 independent estimates of the  $p$ -dimensional parameter using  $2p + 1$  random Fourier features and 10 simulations per parameter value. To be clear, each parameter estimate is based on an independent realization of the process and an independent draw of  $2p + 1$  random features.

**Gaussian model.** For  $t = 1, \dots, n$ , we observe

$$X_t = \mu + \epsilon_t,$$

where  $\epsilon_t \stackrel{\text{iid}}{\sim} N(0, 1)$ . The unknown parameter is  $\mu$ , so we use  $2p + 1 = 3$  random features.





Figure 1: Three random features averaged over  $n = 1000$  times and  $s = 10$  simulations, for each parameter value on a grid in  $[-3, 3]$ . Color denotes  $\mu$ .

We aim to estimate

$$\mu_0 = 0.5,$$

using the time-average estimator.



**Logistic map model.** For  $t = 1, \dots, n$ , define the state by

$$Z_t = \rho Z_{t-1}(1 - Z_{t-1}).$$


We observe

$$X_t = Z_t + \sigma \epsilon_t,$$

where  $\epsilon_t \stackrel{\text{iid}}{\sim} N(0, 1)$  and the unknown initial value  $Z_0$  is sampled from a  $U[0, 1]$  distribution, so that it is fixed before running the procedure. Our goal is to estimate

$$\rho_0 = 3.9, \quad \sigma_0 = 0.1,$$

using the time-average estimator. We emphasize that the logistic map is in a chaotic regime when  $\rho = 3.9$ ; see Devaney (2018) for more information.



**State-space model.** For times  $t = 1, \dots, n$ , define the state by

$$Z_t = Z_{t-1} + \psi(\mu - Z_{t-1}) + \sigma\epsilon_t,$$


where  $\epsilon_t \stackrel{\text{iid}}{\sim} N(0, 1)$ . For dimensions  $j = 1, \dots, d \equiv 25$ , we observe

$$X_t^{(j)} = Z_t + \lambda\xi_t^{(j)},$$

where each  $\xi_t^{(j)} \stackrel{\text{iid}}{\sim} N(0, 1)$ . We aim to estimate

$$\psi_0 = 0.25, \quad \sigma_0 = 0.08, \quad \mu_0 = 2.5, \quad \lambda_0 = 0.05,$$

using the time-average estimator.



**SIR model.** For time horizon  $T = 50$ , let  $(S_v, I_v, R_v)_{v \in [0, T]}$ , be defined by the susceptible-infected-recovered (SIR) system of ODEs

$$\begin{aligned} \frac{dS_v}{dv} &= -\frac{\beta}{N} S_v I_v, \\ \frac{dI_v}{dv} &= \frac{\beta}{N} S_v I_v - \gamma I_v, \\ \frac{dR_v}{dv} &= \gamma I_v, \end{aligned}$$

where  $S_v/N$ ,  $I_v/N$ , and  $R_v/N$  denote the fractions of the population that are susceptible, infected, and recovered (or removed), respectively,  $\beta > 0$  is the transmission rate, and  $\gamma > 0$  is the recovery rate. Let the known initial values be  $S_0 = 980,000$ ,  $I_0 = 20,000$ ,  $R_0 = 0$ , where the number of individuals is  $N = S_0 + I_0 + R_0 = 1,000,000$ .


For each  $t = 1, \dots, n$ , we test  $\text{TE}_t = \lfloor N/100 \rfloor$  people with equal probability using a test with sensitivity  $\text{SE}_t = 0.95$  and specificity  $\text{SP}_t = 0.99$ . We observe the proportion of positive tests

$$X_t = Z_t/\text{TE}_t, \quad Z_t \sim \text{Binomial}(\text{TE}_t, \pi_t),$$

where  $\pi_t = \text{SE}_t I_{Tt/n}/N + (1 - \text{SP}_t)(1 - I_{Tt/n}/N)$ . This model allows the number of tests  $\text{TE}_t$ , sensitivity  $\text{SE}_t$ , and specificity  $\text{SP}_t$  to change with  $t$ , although we keep them constant for simplicity. We aim to estimate

$$\beta_0 = 0.5, \quad \gamma_0 = 0.2,$$

using the rolling-window estimator. Note that the bounds for the optimizer were set to  $[0.01, 1]$  for  $\beta$  and  $[0.01, 0.5]$  for  $\gamma$ .



**Lotka–Volterra model.** For time horizon  $T = 20$ , let  $(Z_v^{(1)}, Z_v^{(2)})_{v \in [0, T]}$  be defined by the Lotka–Volterra predator–prey system of ODEs

$$\begin{aligned} \frac{dZ_v^{(1)}}{dv} &= \alpha Z_v^{(1)} - \beta Z_v^{(1)} Z_v^{(2)}, \\ \frac{dZ_v^{(2)}}{dv} &= \delta Z_v^{(1)} Z_v^{(2)} - \gamma Z_v^{(2)}, \end{aligned}$$

where  $Z_v^{(1)}$  and  $Z_v^{(2)}$  respectively represent the prey and predator population densities, i.e. the number of individuals per unit area. The known initial values are  $Z_0^{(1)} = 21$  and  $Z_0^{(2)} = 9$ . The components of the parameter are the prey growth rate  $\alpha > 0$ , predation rate  $\beta > 0$ , predator growth rate  $\delta > 0$ , predator death rate  $\gamma > 0$ .


The predator-prey populations are observed through noise. For  $t = 1, \dots, n$  and  $j \in \{1, 2\}$ , define

$$X_t^{(j)} = Z_{Tt/n}^{(j)} + \sigma^{(j)} \epsilon_t^{(j)},$$

where each  $\epsilon_t^{(j)} \stackrel{\text{iid}}{\sim} N(0, 1)$ .  $X_t^{(1)}$ ,  $X_t^{(2)}$  denote the observed prey and predator populations, and  $\epsilon_t^{(1)}$ ,  $\epsilon_t^{(2)}$  represent the measurement noise with the noise level controlled by  $\sigma^{(1)}$ ,  $\sigma^{(2)}$ . We aim to estimate

$$\begin{aligned} \alpha_0 &= 0.8, & \beta_0 &= 0.08, & \delta_0 &= 0.06, \\ \gamma_0 &= 1.2, & \sigma_0^{(1)} &= 0.5, & \sigma_0^{(2)} &= 0.7, \end{aligned}$$

using the rolling-window estimator.



**Structural time series model.** We consider a two-dimensional structural time series model with cycles, trends, and an abrupt change-point. For  $u \in [0, 1]$ , define

$$\begin{aligned} Z_u^{(1)} &= \mu^{(1)}u + \alpha \mathbf{1}_{u \geq \tau} + \cos(2\pi\beta^{(1)}u), \\ Z_u^{(2)} &= \mu^{(2)}u - \alpha \mathbf{1}_{u \geq \tau} + \sin(2\pi\beta^{(2)}u). \end{aligned}$$

For  $t = 1, \dots, n$  and  $j \in \{1, 2\}$ , we observe


$$X_t^{(j)} = Z_{t/n}^{(j)} + \sigma^{(j)}\epsilon_t^{(j)},$$

where  $\epsilon_t^{(j)} \stackrel{\text{iid}}{\sim} N(0, 1)$  and each  $\sigma^{(j)} > 0$  controls the level of the noise. We aim to estimate

$$\begin{aligned} \alpha_0 &= 1.8, & \tau_0 &= 0.28, & \mu_0^{(1)} &= -1.5, & \mu_0^{(2)} &= 0.4, \\ \beta_0^{(1)} &= 2.8, & \beta_0^{(2)} &= 2.8, & \sigma_0^{(1)} &= 0.9, & \sigma_0^{(2)} &= 0.9, \end{aligned}$$

using the rolling-window estimator.

Note that our theoretical guarantees only apply to models in which the distributions change smoothly with the parameter. The abrupt change-point in this model violates this assumption. Nevertheless, in practice, our method still reliably estimates the location of the change-point. On the other hand, our theoretical framework allows for *known* abrupt change-points and other general forms of nonstationarity. We emphasize that our theoretical framework does provide guarantees of consistency and asymptotic normality for unknown *smooth* change-points, which can be parametrized using logistic functions, for instance.



## 4 Prediction with Transformers via Simulation

### 4.1 Background

We provide a background on topics that will be discussed in this section.

**Transformers and Time Series Foundation Models (TSFMs).** Drawing inspiration from the remarkable success of transformer-based foundation models for language and vision tasks, many researchers are now seeking to develop foundation models for time series using transformers (Vaswani et al., 2017). *Time series foundation models* (TSFMs) are large models, with millions or even billions of parameters, that are pre-trained on massive collections of real-world and synthetic time series. There has recently been a growing emphasis on pre-training TSFMs using synthetic time series (Ansari et al., 2024, 2025; Liu et al., 2025). In fact, many TSFMs are pre-trained *exclusively* on synthetic time series (Dooley et al., 2023; Bhethanabhotla et al., 2024; Verdenius et al., 2024; Taga et al., 2025; Oreshkin et al., 2026). The synthetic time series are often generated by *parametric* time series



models with a variety of different parameter values. Many TSFMs are also pre-trained on *convex combinations* of time series realizations; see TSMixup from Ansari et al. (2024).

Liu et al. (2025) argue that adaptively fine-tuning TSFMs on carefully constructed synthetic data remains underexplored. That is, post-hoc training of a pre-trained transformer using synthetic time series that are tailored to the user’s known prediction targets. Currently, the statistical foundations of so-called adaptive fine-tuning procedures are not well understood, particularly in the setting of nonstationary time series. This gap presents a natural entry point for statisticians. We develop a method for adaptively fine-tuning TSFMs with relevant synthetic time series data. At the same time, we link it to classical simulation-and-regression.

**Simulation-and-Regression.** We take inspiration from the *simulation-and-regression* methods of Longstaff and Schwartz (2001) and Moussa et al. (2026). The *least squares Monte Carlo* (LSMC) method from Longstaff and Schwartz (2001) was originally developed for pricing American options. For the purposes of this discussion, we will gloss over the details involving finance. In the LSMC algorithm, simulation is used to estimate the regression function  $f$  for two random variables  $X, Y$  by drawing  $X^{(r)}, Y^{(r)}, r = 1, \dots, s$ , and minimizing the squared loss

$$\frac{1}{s} \sum_{r=1}^s (X^{(r)} - f(Y^{(r)}))^2,$$

over some function class to obtain an estimate  $\hat{f}$  of the regression function  $f$ , where the simulations are drawn from the model with some selected parameter.

Similarly, the *extremum Monte Carlo* (XMC) method from Moussa et al. (2026) uses regression (e.g., random forests (Breiman, 2001)) or quantile regression (e.g., quantile regression forests (Meinshausen and Ridgeway, 2006)) to perform forecasting and smoothing for state-space models when the parameter is assumed to be known. In practice, the parameter is estimated (e.g., using simulation-based estimation) and treated as known. Consequently, the estimators in the XMC method minimize a simulation-based empirical risk with respect to some loss  $L$  (e.g., squared loss or pinball loss), where the simulations are drawn from the distribution corresponding to the estimated parameter. That is, they aim to find the function  $f$  in some function class that minimizes the Monte Carlo approximation to the risk

$$R(f) = \mathbb{E}_{P_{\hat{\theta}}} [L(X - f(Y))],$$

for some loss  $L$ , where the expectation is with respect to the distribution  $P_{\hat{\theta}}$  of  $(X, Y)$  corresponding to the estimated parameter  $\hat{\theta}$ . Typically, the estimate  $\hat{\theta}$  is not equal to the true parameter  $\theta_0$ .

## 4.2 State-space Modeling

We give a brief overview of forecasting and smoothing for state-space models. We emphasize that our proposed work on transformer-based time series prediction is not restricted to state-space models, or even to individual time series models. For example, we can use convex combinations of time series models, as in TSMixup from Ansari et al. (2024). Nevertheless, the following definitions and examples will be helpful to have in mind for the discussions in Section 4.4.

*State-space models* decompose observed time series into latent states, control inputs, and noise. Consider the (possibly nonstationary) state-space model for an  $\mathbb{R}^{d_Y}$ -valued observed process  $Y_{1:n}$ , an  $\mathbb{R}^{d_X}$ -valued latent state process  $X_{1:n}$ , and  $\mathbb{R}^{d_Z}$ -valued control inputs  $Z_{1:n}$ . For times  $t = 1, \dots, n$ , let

$$\begin{aligned} Y_t &= M_t(X_t, Z_t, \varepsilon_t^Y, \theta), & (\varepsilon_t^X, \varepsilon_t^Y) &\stackrel{\text{iid}}{\sim} P_\theta(\varepsilon_t^X, \varepsilon_t^Y), \\ X_t &= S_t(X_{t-1}, Z_t, \varepsilon_t^X, \theta), & X_0 &\sim P_\theta(X_0), \end{aligned} \quad (15)$$

where  $M_t$  is the measurement function at time  $t$ ,  $S_t$  is the state transition function at time  $t$ , and  $\theta \in \Theta \subset \mathbb{R}^p$  is a  $p$ -dimensional parameter.

We explain the ideas of smoothing and forecasting in the context of regression for simplicity. The same ideas apply for quantile regression by replacing the conditional expectations with the corresponding conditional quantiles. In the context of state-space modeling, *smoothing* is the task of estimating the past latent states  $X_{1:t}$  using the information up to time  $t$ . This can be done via simulation-based estimates of the regression functions  $\mathbb{E}(X_j \mid Y_{1:t} = y_{1:t}, Z_{1:t} = z_{1:t})$  for  $j = 1, \dots, t$ . *Forecasting* is the task of predicting the future observations or states multiple time-steps ahead. In the state-space setting, a future path of the control inputs must be specified, either deterministically or stochastically, or the model must be specified so that no controls are used. Given the observations  $Y_{1:n}$  and (optionally) control inputs  $Z_{1:n+H}$ , we aim to make point forecasts via simulation-based estimates of  $\mathbb{E}(X_{n+h} \mid Y_{1:n} = y_{1:n}, Z_{1:n+h} = z_{1:n+h})$  or  $\mathbb{E}(Y_{n+h} \mid Y_{1:n} = y_{1:n}, Z_{1:n+h} = z_{1:n+h})$  for  $h = 1, \dots, H$  for some maximal forecasting horizon  $H \in \mathbb{N}$ .

We present some examples. We will use versions of these models in the papers that apply our proposed methods. We use the model from (16) to

estimate a certain notion of latent infectivity in epidemiology via smoothing based on observed daily case counts, vaccination coverage, mobility metrics, and variant relative frequency. We use the model from (17) for generating synthetic time series for calibrating multi-quantile forecasts of time series foundation models.

First, consider a discrete-valued process driven by a latent autoregressive process and controls. We condition on the controls, so that we may view each control as a discretely observed function of continuous time  $Z_t \equiv Z(t/n)$  that is suitably nice (e.g., Hölder continuous in time). We define a (conditional) state-space model for  $(X_{1:n}, Y_{1:n})$  given  $Z_{1:n}$  as follows

$$\begin{aligned} Y_t \mid X_t, Z_t &\sim \text{Pois} \left( \exp \left( \kappa X_t + \beta^\top Z_t \right) \right), \\ X_t &= \mu + \phi(X_{t-1} - \mu) + \varepsilon_t^X, \end{aligned} \quad (16)$$

for some effect sizes  $\kappa \in \mathbb{R}$ ,  $\beta \in \mathbb{R}^{d_Z}$ , autoregressive coefficient  $\phi \in (-1, 1)$ , and mean  $\mu \in \mathbb{R}$ , with  $X_0 \sim N(\mu, \frac{\sigma_X^2}{1-\phi^2})$  and  $\varepsilon_t^X \stackrel{\text{iid}}{\sim} N(0, \sigma_X^2)$  independent of  $Z_{1:n}$ , for some  $\sigma_X > 0$ .

Second, consider a multiplicative seasonality-trend-noise decomposition with no controls

$$\begin{aligned} Y_t &= \text{Se} \left( t, \theta^{Se} \right) \text{Tr} \left( t, \theta^{Tr} \right) \exp \left( \kappa X_t - \frac{\kappa^2 \sigma_X^2}{2(1-\phi^2)} \right) \exp \left( \varepsilon_t^Y - \frac{\sigma_Y^2}{2} \right), \\ X_t &= \phi X_{t-1} + \varepsilon_t^X, \end{aligned} \quad (17)$$

for some scaling factor  $\kappa > 0$  and autoregressive coefficient  $\phi \in (-1, 1)$ , with  $X_0 \sim N(0, \frac{\sigma_X^2}{1-\phi^2})$  and  $(\varepsilon_t^X, \varepsilon_t^Y) \stackrel{\text{iid}}{\sim} N(\mathbf{0}, \Sigma)$ , where

$$\Sigma = \begin{pmatrix} \sigma_X^2 & 0 \\ 0 & \sigma_Y^2 \end{pmatrix},$$

for some  $\sigma_X, \sigma_Y > 0$ . The mean function decomposes into a product of seasonality and trend functions with parameters  $\theta^{Se}$  and  $\theta^{Tr}$ . For pre-training their prior-data fitted network for time series forecasting, Dooley et al. (2023) considers a seasonality function based on a linear combination of sines and cosines (with weekly, monthly, and yearly periods), and a trend function based on multiplying a linear trend and exponential trend.

We make a few remarks about the exponentiated autoregressive process from (17). Recall that the stationary distribution of  $X_t$  is  $N(0, \frac{\sigma_X^2}{1-\phi^2})$ , so the stationary distribution of  $\exp(\kappa X_t)$  is log-normal with mean  $\exp(\frac{1}{2} \kappa^2 \frac{\sigma_X^2}{1-\phi^2})$ .

Hence, the stationary distribution of  $\exp(\kappa X_t - \frac{1}{2}\kappa^2 \frac{\sigma_X^2}{1-\phi^2})$  is log-normal with mean  $\exp(0) = 1$ , and the (local) expectation of the observation  $Y_t$  at time  $t$  is the seasonality times the trend. Therefore, we can interpret this multiplicative decomposition with the exponentiated autoregressive process as a way to stochastically stretch and compress the deterministic mean function, in a manner that is compatible with the framework of locally stationary time series, in the sense of Dahlhaus et al. (2019). When using this model in the context of training transformers for forecasting, this model can be seen as an extension of the generative model used by Dooley et al. (2023).

To estimate the parameters of these models, we use the rolling-window estimator for nonstationary time series from Section 3. It is straightforward to show that these models satisfy the sufficient conditions for our estimator to be consistent and asymptotically normal (i.e. within a certain framework for nonstationary time series, which builds on Dahlhaus et al. (2019); Mies (2023)). Based on Lemma 2.4 and Corollary 2.2 in Kokoszka et al. (2025), it is straightforward to prove that there is a decay of the physical dependence measure for these nonstationary state-space models with exponentiated autoregressive processes.

### 4.3 Confidence Distributions

Our proposed method can be viewed as a frequentist, likelihood-free analogue of the Laplace approximation to the posterior distribution in Bayesian statistics. Specifically, when using the Laplace approximation to the posterior distribution to obtain an approximate posterior predictive distribution. Recall that the Laplace approximation uses a Gaussian distribution centered at the maximum a posteriori (MAP) estimate with covariance given by the inverse observed Fisher information to approximate the posterior distribution. The Bernstein–von Mises theorem is used to justify the Laplace approximation. In contrast, to justify our approach, we use asymptotic theory for likelihood-free parameter estimators for parametric time series models (e.g., Section 3), nonstationary time series theory (Mies, 2023), and the theory of confidence distributions.

In general, our frequentist procedure can be used with any *confidence distribution* over  $\Theta \subset \mathbb{R}^p$  which expresses our uncertainty about the unknown parameter  $\theta_0 \in \mathbb{R}^p$ . Ideally, this distribution is easy to simulate from. We give the definition of confidence distributions for the  $p = 1$  case, and the generalization to the  $p \in \mathbb{N}$  case. In the definition,  $\Theta \subset \mathbb{R}^p$  is the parameter space and  $\mathcal{X}_n$  is the sample space corresponding to the data  $X_{1:n}$ .

**Definition 1** (Confidence Distribution,  $p = 1$ , (Xie and Singh, 2013)). A function  $\widehat{\text{CD}}_n(\cdot) = \widehat{\text{CD}}_n(X_{1:n}, \cdot)$  on  $\mathcal{X}_n \times \Theta \rightarrow [0, 1]$  is called a confidence distribution (CD) for a parameter  $\theta$ , if:

1. For each given  $X_{1:n}$ ,  $\widehat{\text{CD}}_n(\cdot) = \widehat{\text{CD}}_n(X_{1:n}, \cdot)$  is a cumulative distribution function on  $\Theta$ .
2. At the true parameter value  $\theta = \theta_0$ ,  $\widehat{\text{CD}}_n(\theta_0) \equiv \widehat{\text{CD}}_n(X_{1:n}, \theta_0)$ , as a function of the sample  $X_{1:n}$ , follows the uniform distribution  $U[0, 1]$ .

Also, the function  $\widehat{\text{CD}}_n(\cdot)$  is an asymptotic confidence distribution (aCD), if the  $U[0, 1]$  requirement is true only asymptotically.

The definition of confidence distributions for the general  $p \in \mathbb{N}$  case involves depth functions. As discussed in Liu et al. (1999), a data depth is a way of measuring how deep or central a point  $x \in \mathbb{R}^d$  is with respect to a probability distribution  $P$  on  $\mathbb{R}^d$ ,  $d \in \mathbb{N}$ , or with respect to a given data cloud  $X_{1:n}$ . For example, the Mahalanobis depth ( $D_M$ ) (Mahalanobis, 1936) at  $x \in \mathbb{R}^d$  with respect to  $P$  is defined as

$$D_M(P, x) = \frac{1}{1 + (x - \mu_P)^\top \Sigma_P^{-1} (x - \mu_P)},$$

where  $\mu_P$  and  $\Sigma_P$  are the mean vector and the dispersion matrix of  $P$ , respectively. We obtain the sample version of  $D_M$  by replacing  $\mu_P$  and  $\Sigma_P$  by the corresponding sample estimates  $\hat{\mu}_P$  and  $\hat{\Sigma}_P$ . See Liu et al. (1999) for more examples of data depth. In general, a point with a larger depth value (i.e. closer to 1) means that it is closer to the center of the distribution or more central in a data cloud, while points with smaller depth values (i.e. closer to 0) are farther away from the center, so data depth is a center-outward ordering of multivariate observations.

**Definition 2** (Depth Confidence Distribution,  $p \in \mathbb{N}$ , (Liu et al., 2022)). A function  $\widehat{\text{CD}}_n(\cdot) \equiv \widehat{\text{CD}}_n(X_{1:n}, \cdot)$  on  $\Theta \subset \mathbb{R}^p$  is called a depth confidence distribution (depth-CD) associated with depth function  $D$  for a vector-valued parameter  $\theta$ , if:

1. It is a distribution function on the parameter space  $\Theta$  for any fixed sample set  $X_{1:n}$ .
2. The  $(1 - \alpha)$  “central region” of the distribution  $\widehat{\text{CD}}_n(\cdot)$ ,  $\mathcal{R}_{1-\alpha}(\widehat{\text{CD}}_n) = \{\theta \in \Theta : C_{\widehat{\text{CD}}_n}(\theta) \geq \alpha\}$ , is a confidence region for  $\theta$  with a coverage probability of  $(1 - \alpha)$ . Here, the centrality function associated with depth  $D$  and confidence distribution  $\widehat{\text{CD}}_n(\cdot)$ , that is,  $C_{\widehat{\text{CD}}_n}(\theta)$ , is also referred to as a depth confidence curve (depth-CV).

If the statements in (2) hold only asymptotically, then we refer to  $\widehat{\text{CD}}_n$  and  $C_{\widehat{\text{CD}}_n}$  as asymptotic depth-CD and asymptotic depth-CV, respectively.

In this work, we focus on asymptotic Gaussian depth confidence distributions of the form  $\widehat{\text{CD}}_n = N(\hat{\theta}, \hat{\Sigma}/n)$  for a multi-dimensional parameter  $\theta \in \mathbb{R}^p$ ,  $p \in \mathbb{N}$ . Let  $q_{1-\alpha}$  be the  $1-\alpha$  quantile of  $\chi_p^2$ , i.e.,  $\mathbb{P}(\chi_p^2 \leq q_{1-\alpha}) = 1-\alpha$ . As usual, the asymptotic  $1-\alpha$  confidence region for  $\theta_0$  is

$$\mathcal{R}_{1-\alpha}(\widehat{\text{CD}}_n) = \{\theta \in \Theta : (\theta - \hat{\theta})^\top (\hat{\Sigma}/n)^{-1} (\theta - \hat{\theta}) \leq q_{1-\alpha}\},$$

so applying the  $\chi_p^2$  CDF  $F_{\chi_p^2}$ , subtracting by 1, dividing by  $-1$ , yields

$$\mathcal{R}_{1-\alpha}(\widehat{\text{CD}}_n) = \{\theta \in \Theta : C_{\widehat{\text{CD}}_n}(\theta) = 1 - F_{\chi_p^2}[(\theta - \hat{\theta})^\top (\hat{\Sigma}/n)^{-1} (\theta - \hat{\theta})] \geq \alpha\}.$$

See Example 3 in Liu et al. (2022) for an illustration with a bivariate normal distribution.

Estimating the asymptotic covariance matrix of the parameter, as in Theorem 3.3 of Pakes and Pollard (1989), will involve estimating (1) the Jacobian of the population discrepancy function at  $\theta_0$ , and (2) the asymptotic covariance matrix of the  $\sqrt{n}$ -scaled sample discrepancy function at  $\theta_0$ . Note that we refer to the discrepancy function as the function of the parameter value, whose norm induces the objective function. The former can be estimated using numerical differentiation tools, and the later can be estimated using covariance estimation results from nonstationary time series theory Mies (2023); Mies and Steland (2023). In both estimators, we will use a plug-in approach with the estimated parameter  $\hat{\theta}$  in place of the true parameter  $\theta_0$ .

We justify this Gaussianity through the asymptotic normality of the parameter estimator (e.g., Section 3), nonstationary time series theory (Mies, 2023), and Theorem 3 (Theorem 2 in Liu et al. (2022)).

**Theorem 3** (Depth-CDs from Pivot Statistics (Liu et al., 2022)). *Assume that  $A_n(X_{1:n})$  is a nonsingular matrix such that  $A_n(X_{1:n})(\hat{\theta} - \theta)$  follows a distribution  $Q_n$  that is free of all unknown parameters. Also assume that  $\eta_n$  is a random vector, independent of the sample  $X_{1:n}$ , following the distribution  $Q_n$ . Then, conditional on  $X_{1:n}$ , the distribution function of  $(\hat{\theta} - A_n(X_{1:n})^{-1}\eta_n)$  is a depth-CD for  $\theta$ , under the following conditions:*

1. *The depth  $D$  is affine-invariant.*
2. *The depth contours  $\{\eta_n : D_{Q_n}(\eta_n) = \epsilon\}$  all have probability zero with respect to  $Q_n$  for any  $\epsilon > 0$ .*

Alternatively, a parametric bootstrap approach can be used (Efron, 1979; Efron and Tibshirani, 1994). To use a parametric bootstrap approach, we would estimate the parameter on the observed data, simulate from the model with the estimated parameter, then re-estimate the parameter for each simulated realization. The sampling distribution of the re-estimates is used as the depth confidence distribution. Theorem 1 in Liu et al. (2022) shows that this bootstrap distribution is an asymptotic depth confidence distribution under certain regularity conditions. For more information about the connections between bootstrap distributions and confidence distributions, see Example 3 in Xie and Singh (2013) and the related discussions in Efron (1998).

We opt for an asymptotic Gaussian depth confidence distribution, rather than a bootstrap depth confidence distribution, to improve the computational efficiency of our method. This computational efficiency is particularly important, because our proposed method is intended to be applicable to models with intractable likelihoods, where parameter estimation typically relies on simulation-based procedures that are substantially more computationally demanding than methods like maximum likelihood estimation. Therefore, our procedure shares some of the motivations behind the Krinsky and Robb (1986) method.

#### 4.4 Proposed Work

Motivated by the topics discussed in Section 4.1, we study simulation-based time series prediction with transformers. Specifically, we focus on two time series prediction tasks: (1) forecasting future observations and latent states, and (2) smoothing past latent states.

Our method inherits the computational benefits of sequence-to-sequence modeling. In the context of smoothing, our proposed method only requires training one sequence-to-sequence model, as opposed to separate models for each time as in Moussa et al. (2026). Similarly, we only require one sequence-to-sequence model to obtain multi-step forecasts.

We propose a transformer-based simulation-and-regression framework that incorporates uncertainty about the true parameter. We suggest minimizing a simulation-based empirical risk using a sequence-to-sequence loss, where the simulations are generated by first drawing a parameter from an estimated *confidence distribution* over the parameter space  $\Theta \subset \mathbb{R}^p$  (see Definition 2) and then simulating a time series from the model at that parameter. That is, we aim to find the sequence-to-sequence function  $f$  in some function class (i.e., the class of transformer networks) that minimizes the Monte Carlo

approximation to the risk

$$\begin{aligned} R_{\widehat{P}_{\text{mix}}}(f) &= \mathbb{E}_{\widehat{P}_{\text{mix}}} \left[ \sum_{t \in \mathcal{T}} L(X_t - [f(Y)]_t) \right] \\ &= \mathbb{E}_{\theta \sim \widehat{\text{CD}}_n} \left[ \mathbb{E}_{P_\theta} \left[ \sum_{t \in \mathcal{T}} L(X_t - [f(Y)]_t) \right] \right], \end{aligned} \quad (18)$$

where the second equality holds by Fubini's theorem. Here,  $L$  is some loss,

$$\widehat{P}_{\text{mix}} = \int_{\Theta} P_\theta d\widehat{\text{CD}}_n(\theta),$$

is the mixture distribution,  $\widehat{\text{CD}}_n$  is the estimated confidence distribution over  $\Theta$ ,  $P_\theta$  is the distribution of the sequences corresponding to  $\theta$ ,  $\mathcal{T}$  is a set of times, and  $[f(Y)]_t$  is the sequence-to-sequence function at time  $t$ . The details of the procedure are presented in Section 4.5.

In contrast, the simulation-and-regression methods from Longstaff and Schwartz (2001) and Moussa et al. (2026), discussed in Section 4.1, only simulate from the model at the estimated parameter. The corresponding sequence-to-sequence version of the XMC method from Moussa et al. (2026) would minimize the Monte Carlo approximation to the risk

$$R_{\widehat{P}_\theta}(f) = \mathbb{E}_{\widehat{P}_\theta} \left[ \sum_{t \in \mathcal{T}} L(X_t - [f(Y)]_t) \right], \quad (19)$$

over some sequence-to-sequence function class. Here,  $L$  is some loss,  $\widehat{P}_\theta$  is the distribution of the sequences corresponding to the estimated parameter  $\hat{\theta}$ ,  $\mathcal{T}$  is a set of times, and  $[f(Y)]_t$  is the sequence-to-sequence function at time  $t$ .

When the asymptotic variance of the parameter estimator is “relatively small”, there is more concentration of the simulation budget in regions of the parameter space near the parameter estimate. On the other hand, when the asymptotic variance of the parameter estimator is “relatively large”, there is less concentration of the simulation budget in regions of the parameter space near the parameter estimate. Hence, roughly speaking, the extremes are a point mass at the parameter estimate and a uniform distribution over the parameter space. In this sense, our method can be situated between two existing paradigms in simulation-based time series prediction: (1) simulation-and-regression for time series using parameter point estimates as in Moussa et al. (2026), and (2) time series foundation models pre-trained on synthetic



time series as in Dooley et al. (2023); Bhethanabhotla et al. (2024); Verdenius et al. (2024); Taga et al. (2025); Oreshkin et al. (2026).

When used with an untrained transformer, our procedure can be viewed as a *robust simulation-and-regression method*, as it accounts for parameter uncertainty during training. In contrast, classical methods just use the point estimate. When used with a pre-trained transformer, our method produces a *contextualized time series foundation model*, as it performs post-hoc training (i.e. fine-tuning) with relevant synthetic time series. This fills a gap in the TSFM literature, i.e., the “missed opportunity” of “adaptively fine-tuning” TSFMs on “carefully designed synthetic data” discussed in Liu et al. (2025).

We instantiate our framework in the context of sequence-to-sequence quantile regression, so that the transformer outputs multiple quantile predictions. Typically, the set of quantiles  $\tau \in \{0.1, 0.2, \dots, 0.9\}$  is used. For example, see Das et al. (2024); Auer et al. (2025); Wang et al. (2025); Adler et al. (2025). In this case, the loss function  $L$  is chosen to be the sum of multiple quantile losses (with a quantile crossing penalty). Extensive details are provided in the Appendix of the proposal slides. We aim to answer the following question about calibrating multi-quantile predictions across quantile levels.

**Question.** Does our method (1) improve the calibration of pre-trained transformers, and (2) yield well-calibrated transformers when trained from scratch?

## 4.5 Algorithms

The proposed forecasting and smoothing methods are presented below. We present the forecasting method in the context of forecasting future latent states. However, the same method can be used to forecast future observations, simply by replacing the future latent states with future observations.

The proposed procedures utilize an estimator  $\hat{\theta}$  of the true parameter  $\theta_0$  of a (possibly nonstationary) parametric time series model. The estimator is assumed to be consistent and asymptotically normal. For example, the estimators from Section 3 can be used. For the cross-validation scheme, we introduce an aggregation function  $\text{Agg}(\cdot)$  for more flexibility. For example,  $\text{Agg}(\cdot)$  can be taken to be the average over the  $K$  folds, or just the error in the validation fold if using an 80/20 train/validation split. As usual, for the loss function  $L$ , the squared loss is used for regression and the pinball loss is used for quantile regression.

### Forecasting with Transformers.

- **Inputs:** Observations  $Y_{1:n}^{\text{obs}}$ , state-space model (15), number of simulations  $s$ , number of folds  $K$  for cross-validation, proportion of simulations in each fold  $c_1, \dots, c_K$ , cross-validation error aggregation function  $\text{Agg}(\cdot)$ , estimator  $\hat{\theta}$ , estimation method for asymptotic covariance of  $\hat{\theta}$ , maximal forecasting horizon  $H$ , and loss function  $L$ .
- Obtain estimate  $\hat{\theta}$  of the true parameter  $\theta_0$ .
- Obtain estimate  $\hat{\Sigma}$  of the asymptotic covariance  $\Sigma$  of estimator  $\hat{\theta}$ .
- For each  $r \in [s]$ , draw  $\theta^{(r)} \sim N(\hat{\theta}, \hat{\Sigma}/n)$  from the confidence distribution, and simulate from the model (15),

$$X_{1:n+H}^{(r)}, Y_{1:n+H}^{(r)} \sim P_{n+H, \theta^{(r)}}.$$

- Denote the  $r$ -th sequence-to-sequence loss by

$$\ell_r(f) = \sum_{h=1}^H L(X_{n+h}^{(r)} - [f(Y_{1:n}^{(r)})]_{n+h}),$$

where  $[f(Y_{1:n}^{(r)})]_t$  is the  $r$ -th sequence-to-sequence prediction at time  $t$ .

- Partition the simulation indices  $[s]$  into  $K$  folds  $F_1, \dots, F_K$ , where,

$$F_i = \{r \in [s] : \sum_{j=1}^{i-1} c_j s < r \leq \sum_{j=1}^i c_j s\}, \quad i = 1, \dots, K.$$

- For each  $k \in [K]$  and each value of the tuning parameter  $\lambda \in \Lambda$ , obtain the minimizer  $\hat{f}_\lambda^{-k}$  of the average loss

$$\frac{1}{|F_{-k}|} \sum_{r \in F_{-k}} \ell_r(f),$$

over the function space of transformers, where  $F_{-k} = \bigcup_{j \in [K] \setminus \{k\}} F_j$ .

- Record the corresponding error on the  $k$ -th validation set

$$e_k(\lambda) = \frac{1}{|F_k|} \sum_{r \in F_k} \ell_r(\hat{f}_\lambda^{-k}).$$

- Compute the aggregated error over the  $K$  folds

$$\text{CV}(\lambda) = \text{Agg}(e_1(\lambda), \dots, e_K(\lambda)).$$

- Find the value  $\lambda^*$  of the tuning parameter  $\lambda$  that minimizes  $\text{CV}(\lambda)$ .
- Obtain the estimate  $\hat{f}$  by minimizing the average loss using all  $s$  simulations with the selected tuning parameter  $\lambda^*$  to

$$\frac{1}{s} \sum_{r \in [s]} \ell_r(f).$$

- Apply the estimated function  $\hat{f}(Y_{1:n}^{\text{obs}})$  to make the multi-step forecast of the future observations at times  $n+1, \dots, n+H$ .

### Smoothing with Transformers.

- **Inputs:** Observations  $Y_{1:n}^{\text{obs}}$ , state-space model (15), number of simulations  $s$ , number of folds  $K$  for cross-validation, proportion of simulations in each fold  $c_1, \dots, c_K$ , cross-validation error aggregation function  $\text{Agg}(\cdot)$ , estimator  $\hat{\theta}$ , estimation method for asymptotic covariance of  $\hat{\theta}$ , and loss function  $L$ .
- Obtain estimate  $\hat{\theta}$  of the true parameter  $\theta_0$ .
- Obtain estimate  $\hat{\Sigma}$  of the asymptotic covariance  $\Sigma$  of estimator  $\hat{\theta}$ .
- For each  $r \in [s]$ , draw  $\theta^{(r)} \sim N(\hat{\theta}, \hat{\Sigma}/n)$  from the confidence distribution, and simulate from the model (15),

$$X_{1:n}^{(r)}, Y_{1:n}^{(r)} \sim P_{n, \theta^{(r)}}.$$

- Denote the  $r$ -th sequence-to-sequence loss by

$$\ell_r(f) = \sum_{t \in [n]} L(X_t^{(r)} - [f(Y_{1:n}^{(r)})]_t),$$

where  $[f(Y_{1:n}^{(r)})]_t$  is the  $r$ -th sequence-to-sequence prediction at time  $t$ .

- Partition the simulation indices  $[s]$  into  $K$  folds  $F_1, \dots, F_K$ , where,

$$F_i = \{r \in [s] : \sum_{j=1}^{i-1} c_j s < r \leq \sum_{j=1}^i c_j s\}, \quad i = 1, \dots, K.$$

- For each  $k \in [K]$  and each value of the tuning parameter  $\lambda \in \Lambda$ , obtain the minimizer  $\hat{f}_\lambda^{-k}$  of the average loss

$$\frac{1}{|F_{-k}|} \sum_{r \in F_{-k}} \ell_r(f),$$

over the function space of transformers, where  $F_{-k} = \bigcup_{j \in [K] \setminus \{k\}} F_j$ .

- Record the corresponding error on the  $k$ -th validation set

$$e_k(\lambda) = \frac{1}{|F_k|} \sum_{r \in F_k} \ell_r(\hat{f}_\lambda^{-k}).$$

- Compute the aggregated error over the  $K$  folds

$$\text{CV}(\lambda) = \text{Agg}(e_1(\lambda), \dots, e_K(\lambda)).$$

- Find the value  $\lambda^*$  of the tuning parameter  $\lambda$  that minimizes  $\text{CV}(\lambda)$ .
- Obtain the estimate  $\hat{f}$  by minimizing the average loss using all  $s$  simulations with the selected tuning parameter  $\lambda^*$  to

$$\frac{1}{s} \sum_{r \in [s]} \ell_r(f).$$

- Apply the estimated function  $\hat{f}(Y_{1:n}^{\text{obs}})$  to predict the latent states at times  $1, \dots, n$ .

## 5 Timeline

**Step 1a: Proposal.** I plan to propose my thesis on April 3rd, 2026.

**Step 1b: Revise and Submit.** I will submit the revision of the conditional independence testing paper corresponding to Section 2, which has received a revise decision from the journal, by May 3rd, 2026. Soon afterwards, I will submit the simulation-based parameter estimation paper corresponding to Section 3.

**Step 2: Calibrating Transformers.** I will answer the proposed question from Section 4.4 using the procedures from Section 4.5. My *primary objectives*, from now until the defense of my thesis, are to: develop the method, evaluate performance through extensive experiments on transformer-based forecasting and smoothing, write the paper, and the submit paper. This paper will involve training (relatively small) transformers from scratch and fine-tuning pre-trained transformers by simulating from parametric time series models.

**Possible Extension: Application to Epidemiology.** *After submitting the main paper*, if time allows, a possible extension is to apply the method to epidemiological time series. The goal is to smooth a latent infectivity process by simulating from a version of the model from (16). I have already prepared the data (variant relative frequency, vaccine coverage, mobility metrics, and daily case counts). However, the main paper *must be submitted* before proceeding to this applied paper, so this should be viewed as a stretch goal.

**Step 3: Defense.** I plan to defend my thesis before the spring defense deadline set by the Department, April 15, 2027, and submit the paperwork by May 1, 2027.

## References

- Adler, C., Chang, Y., Draxler, F., Abdi, S., and Smyth, P. (2025). Beyond accuracy: Are time series foundation models well-calibrated? arXiv preprint arXiv:2510.16060.
- Ansari, A. F., Shchur, O., Küken, J., Auer, A., Han, B., Mercado, P., Rangapuram, S. S., et al. (2025). Chronos-2: From univariate to universal forecasting. arXiv preprint arXiv:2510.15821.
- Ansari, A. F., Stella, L., Turkmen, C., Zhang, X., Mercado, P., Shen, H., Shchur, O., et al. (2024). Chronos: Learning the language of time series. arXiv preprint arXiv:2403.07815.
- Auer, A., Podest, P., Klotz, D., Böck, S., Klambauer, G., and Hochreiter, S. (2025). Tirez: Zero-shot forecasting across long and short horizons with enhanced in-context learning. arXiv preprint arXiv:2505.23719.

- Bhethanabhotla, S. K., Swelam, O., Siems, J., Salinas, D., and Hutter, F. (2024). Mamba4cast: Efficient zero-shot time series forecasting with state space models. arXiv preprint arXiv:2410.09385.
- Breiman, L. (2001). Random forests. *Machine Learning*, 45:5–32.
- Chen, L., Smetanina, E., and Wu, W. B. (2022). Estimation of nonstationary nonparametric regression model with multiplicative structure. *The Econometrics Journal*, 25(1):176–214.
- Dahlhaus, R., Richter, S., and Wu, W. B. (2019). Towards a general theory for nonlinear locally stationary processes. *Bernoulli*, 25(2):1013–1044.
- Das, A., Kong, W., Sen, R., and Zhou, Y. (2024). A decoder-only foundation model for time-series forecasting. *International Conference on Machine Learning*, 41.
- Devaney, R. L. (2018). *A first course in chaotic dynamical systems: theory and experiment*. CRC Press.
- Ding, X. and Zhou, Z. (2021). Simultaneous sieve inference for time-inhomogeneous nonlinear time series regression. arXiv preprint arXiv:2112.08545.
- Dooley, S., Khurana, G. S., Mohapatra, C., Naidu, S. V., White, C., et al. (2023). Forecastpfm: Synthetically-trained zero-shot forecasting. *Advances in Neural Information Processing Systems*, 36:2403–2426.
- Efron, B. (1979). Bootstrap methods: Another look at the jackknife. *Annals of Statistics*, 7(1):1–26.
- Efron, B. (1998). Ra fisher in the 21st century. *Statistical Science*, pages 95–114.
- Efron, B. and Tibshirani, R. J. (1994). *An introduction to the bootstrap*. Chapman and Hall/CRC.
- Gourieroux, C. and Monfort, A. (1996). *Simulation-based econometric methods*. Oxford University Press, Oxford, England.
- Hardt, M., Price, E., and Srebro, N. (2016). Equality of opportunity in supervised learning. *Advances in Neural Information Processing Systems*, 29:3323–3331.

- Kokoszka, P., Mohammadi, N., and Wang, H. (2025). Dependence properties of stochastic volatility models. *Journal of Time Series Analysis*, 46(3):421–431.
- Krinsky, I. and Robb, A. L. (1986). On approximating the statistical properties of elasticities. *The Review of Economics and Statistics*, pages 715–719.
- Kurusu, D., Fukami, R., and Koike, Y. (2025). Adaptive deep learning for nonlinear time series models. *Bernoulli*, 31(1):240–270.
- Liu, D., Liu, R. Y., and ge Xie, M. (2022). Nonparametric fusion learning for multiparameters: Synthesize inferences from diverse sources using data depth and confidence distribution. *Journal of the American Statistical Association*, 117(540):2086–2104.
- Liu, R. Y., Parelius, J. M., and Singh, K. (1999). Multivariate analysis by data depth: descriptive statistics, graphics and inference. *The Annals of Statistics*, 27(3):783–858.
- Liu, X., Aksu, T., Liu, J., Wen, Q., Liang, Y., Xiong, C., Savarese, S., Sahoo, D., Li, J., and Liu, C. (2025). Empowering time series analysis with synthetic data: A survey and outlook in the era of foundation models. arXiv preprint arXiv:2503.11411.
- Longstaff, F. A. and Schwartz, E. S. (2001). Valuing american options by simulation: A simple least-squares approach. *The Review of Financial Studies*, 14(1):113–147.
- Mahalanobis, P. C. (1936). On the generalized distance in statistics. *Proceedings of the National Institute of Science of India*, 2:49–55.
- Meinshausen, N. and Ridgeway, G. (2006). Quantile regression forests. *Journal of Machine Learning Research*, 7(6):983–999.
- Mies, F. (2023). Functional estimation and change detection for nonstationary time series. *Journal of the American Statistical Association*, 118(542):1011–1022.
- Mies, F. and Steland, A. (2023). Sequential gaussian approximation for nonstationary time series in high dimensions. *Bernoulli*, 29(4):3114–3140.
- Moussa, K., Blasques, F., and Koopman, S. J. (2026). Extremum monte carlo filters: signal extraction via simulation and regression. *Journal of Business and Economic Statistics*, pages 1–13.

- Oreshkin, B. N., Jauhari, M., Selvam, R. K., Wolff, M., Pan, W., Ramasubramanian, S., Olivares, K. G., et al. (2026). Zero-shot forecasting by simulation alone. arXiv preprint arXiv:2601.00970.
- Pakes, A. and Pollard, D. (1989). Simulation and the asymptotics of optimization estimators. *Econometrica*, pages 1027–1057.
- Shah, R. D. and Peters, J. (2020). The hardness of conditional independence testing and the generalised covariance measure. *Annals of Statistics*, 48(3):1514–1538.
- Taga, E. O., Ildiz, M. E., and Oymak, S. (2025). Timepfn: Effective multivariate time series forecasting with synthetic data. *Proceedings of the AAAI conference on artificial intelligence*, 39(19):20761–20769.
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Łukasz Kaiser, and Polosukhin, I. (2017). Attention is all you need. *Advances in Neural Information Processing Systems*, 30:3323–3331.
- Verdenius, S., Zerio, A., and Wang, R. L. (2024). Lat-pfn: A joint embedding predictive architecture for in-context time-series forecasting. arXiv preprint arXiv:2405.10093.
- Vogt, M. (2012). Nonparametric regression for locally stationary time series. *Annals of Statistics*, 50(5):2601–2633.
- Wang, X., Zhou, T., Gao, J., Ding, B., and Zhou, J. (2025). Output scaling: Yinglong-delayed chain of thought in a large pretrained time series forecasting model. arXiv preprint arXiv:2506.11029.
- Wieck-Sosa, M., Haddad, M. F. C., and Ramdas, A. (2025). Conditional independence testing with a single realization of a multivariate nonstationary nonlinear time series. arXiv preprint arXiv:2504.21647.
- Wu, W. B. (2005). Nonlinear system theory: another look at dependence. *Proceedings of the National Academy of Sciences*, 102(40):14150–14154.
- Xie, M. and Singh, K. (2013). Confidence distribution, the frequentist distribution estimator of a parameter: A review. *International Statistical Review*, 81(1):3–39.
- Yousuf, K. and Ng, S. (2021). Boosting high dimensional predictive regressions with time varying parameters. *Journal of Econometrics*, 224(1):60–87.



- Zhang, T. and Wu, W. B. (2015). Time-varying nonlinear regression models: nonparametric estimation and model selection. *Annals of Statistics*, 43(2):741–768.
- Zhou, Z. and Wu, W. B. (2009). Local linear quantile estimation for nonstationary time series. *The Annals of Statistics*, 37(5B):2696–2729.