

Simulation-based Time Series Analysis:

Likelihood-free Estimation, Testing for Structure, and Prediction with Transformers using Synthetic Data

Michael Wieck-Sosa

Department of Statistics & Data Science, Carnegie Mellon University

April 3, 2026

Overview of presentation

- ▶ Introduction to time series

Overview of presentation

- ▶ Introduction to time series
- ▶ Conditional independence testing

Overview of presentation

- ▶ Introduction to time series
- ▶ Conditional independence testing
- ▶ Simulation-based estimation

Overview of presentation

- ▶ Introduction to time series
- ▶ Conditional independence testing
- ▶ Simulation-based estimation
- ▶ Prediction with transformers trained on synthetic data

Time series setting

- ▶ We observe X_t , $t = 1, \dots, n$, for some sample size $n \in \mathbb{N}$

Time series setting

- ▶ We observe X_t , $t = 1, \dots, n$, for some sample size $n \in \mathbb{N}$
- ▶ Each observation X_t takes values in $\mathbb{R}^d$ for some dimension $d \in \mathbb{N}$

Time series setting

- ▶ We observe X_t , $t = 1, \dots, n$, for some sample size $n \in \mathbb{N}$
- ▶ Each observation X_t takes values in $\mathbb{R}^d$ for some dimension $d \in \mathbb{N}$
- ▶ **Temporal dependence:** X_t may depend on the past $X_{t-1}, X_{t-2}, \dots$


Time series setting

- ▶ We observe X_t , $t = 1, \dots, n$, for some sample size $n \in \mathbb{N}$
- ▶ Each observation X_t takes values in $\mathbb{R}^d$ for some dimension $d \in \mathbb{N}$
- ▶ **Temporal dependence:** X_t may depend on the past $X_{t-1}, X_{t-2}, \dots$
- ▶ **Nonstationarity:** the joint distribution of $X_{t_1+\tau}, \dots, X_{t_m+\tau}$ at $m \in \mathbb{N}$ distinct time points $t_1, \dots, t_m$ may change with the time-offset τ

Time series setting

- ▶ We observe X_t , $t = 1, \dots, n$, for some sample size $n \in \mathbb{N}$
- ▶ Each observation X_t takes values in $\mathbb{R}^d$ for some dimension $d \in \mathbb{N}$
- ▶ **Temporal dependence:** X_t may depend on the past $X_{t-1}, X_{t-2}, \dots$
- ▶ **Nonstationarity:** the joint distribution of $X_{t_1+\tau}, \dots, X_{t_m+\tau}$ at $m \in \mathbb{N}$ distinct time points $t_1, \dots, t_m$ may change with the time-offset τ
- ▶ Example: time series in epidemiology

COVID-19 time series May 2020–2022 ($n=365 \times 2=730$, $d=4$)



Conditional independence testing

- ▶ Let $X_{1:n}$, $Y_{1:n}$, $Z_{1:n}$ be nonstationary time series taking values in $\mathbb{R}^{d_x}$, $\mathbb{R}^{d_y}$, $\mathbb{R}^{d_z}$

Conditional independence testing

- ▶ Let $X_{1:n}$, $Y_{1:n}$, $Z_{1:n}$ be nonstationary time series taking values in $\mathbb{R}^{d_x}$, $\mathbb{R}^{d_y}$, $\mathbb{R}^{d_z}$
- ▶ Let these variables incorporate any desired leads and lags (to simplify notation)

Conditional independence testing

- ▶ Let $X_{1:n}$, $Y_{1:n}$, $Z_{1:n}$ be nonstationary time series taking values in $\mathbb{R}^{d_x}$, $\mathbb{R}^{d_y}$, $\mathbb{R}^{d_z}$
- ▶ Let these variables incorporate any desired leads and lags (to simplify notation)
- ▶ The conditioning set of variables Z_t must be known at time t

Conditional independence testing

- ▶ Let $X_{1:n}$, $Y_{1:n}$, $Z_{1:n}$ be nonstationary time series taking values in $\mathbb{R}^{d_x}$, $\mathbb{R}^{d_y}$, $\mathbb{R}^{d_z}$
- ▶ Let these variables incorporate any desired leads and lags (to simplify notation)
- ▶ The conditioning set of variables Z_t must be known at time t
- ▶ We present a test for the null hypothesis that

$$X_t \perp\!\!\!\perp Y_t \mid Z_t \text{ for all } t \in [n]$$

Conditional independence testing

- ▶ Let $X_{1:n}$, $Y_{1:n}$, $Z_{1:n}$ be nonstationary time series taking values in $\mathbb{R}^{d_x}$, $\mathbb{R}^{d_y}$, $\mathbb{R}^{d_z}$
- ▶ Let these variables incorporate any desired leads and lags (to simplify notation)
- ▶ The conditioning set of variables Z_t must be known at time t
- ▶ We present a test for the null hypothesis that

$$X_t \perp\!\!\!\perp Y_t \mid Z_t \text{ for all } t \in [n]$$

- ▶ Motivations: understanding structure, variable selection, and causal discovery

Conditional independence testing

- ▶ Let $X_{1:n}$, $Y_{1:n}$, $Z_{1:n}$ be nonstationary time series taking values in $\mathbb{R}^{d_X}$, $\mathbb{R}^{d_Y}$, $\mathbb{R}^{d_Z}$
- ▶ Let these variables incorporate any desired leads and lags (to simplify notation)
- ▶ The conditioning set of variables Z_t must be known at time t
- ▶ We present a test for the null hypothesis that

$$X_t \perp\!\!\!\perp Y_t \mid Z_t \text{ for all } t \in [n]$$

- ▶ Motivations: understanding structure, variable selection, and causal discovery
- ▶ Example: links in the global economic system

Real data application: links in the global economic system

- ▶ S&P 500 (US), FTSE 100 (UK), Hang Seng (HK), and Nikkei 225 (JP)

Real data application: links in the global economic system

- ▶ S&P 500 (US), FTSE 100 (UK), Hang Seng (HK), and Nikkei 225 (JP)
- ▶ Daily log returns based on adjusted closing prices


Real data application: links in the global economic system

- ▶ S&P 500 (US), FTSE 100 (UK), Hang Seng (HK), and Nikkei 225 (JP)
- ▶ Daily log returns based on adjusted closing prices
- ▶ Trading days from January 2022 to April 2025 ($n = 860$)

Real data application: links in the global economic system

- ▶ S&P 500 (US), FTSE 100 (UK), Hang Seng (HK), and Nikkei 225 (JP)
- ▶ Daily log returns based on adjusted closing prices
- ▶ Trading days from January 2022 to April 2025 ($n = 860$)
- ▶ We show a subset of the results based on BH-adjusted p-values

Closing prices of each stock index over time



Independence testing reveals same day returns are dependent

- For each pair $X, Y \in \{\text{S\&P}, \text{FTSE}, \text{HangSeng}, \text{Nikkei}\}$, test

$$X_t \perp\!\!\!\perp Y_t \text{ for all times } t \in [n]$$

Independence testing reveals same day returns are dependent

- ▶ For each pair $X, Y \in \{\text{S\&P}, \text{FTSE}, \text{HangSeng}, \text{Nikkei}\}$, test

$$X_t \perp\!\!\!\perp Y_t \text{ for all times } t \in [n]$$

- ▶ We reject *every* null hypothesis at the significance level $\alpha = 0.05$

Conditional independence testing reveals structure

- For each triplet $X, Y, Z \in \{\text{S\&P}, \text{FTSE}, \text{HangSeng}, \text{Nikkei}\}$, test

$$X_t \perp\!\!\!\perp Y_t \mid Z_t \text{ for all times } t \in [n]$$

Conditional independence testing reveals structure

- ▶ For each triplet $X, Y, Z \in \{\text{S\&P}, \text{FTSE}, \text{HangSeng}, \text{Nikkei}\}$, test

$$X_t \perp\!\!\!\perp Y_t \mid Z_t \text{ for all times } t \in [n]$$

- ▶ **Retain:** S\&P(t) is independent of HangSeng(t) and Nikkei(t) given FTSE(t)

Conditional independence testing reveals structure

- ▶ For each triplet $X, Y, Z \in \{\text{S\&P}, \text{FTSE}, \text{HangSeng}, \text{Nikkei}\}$, test

$$X_t \perp\!\!\!\perp Y_t \mid Z_t \text{ for all times } t \in [n]$$

- ▶ **Retain:** S\&P(t) is independent of HangSeng(t) and Nikkei(t) given FTSE(t)
- ▶ **Reject:** S\&P(t) is independent of FTSE(t) given HangSeng(t) or Nikkei(t)

Conditional independence testing reveals structure

- ▶ For each triplet $X, Y, Z \in \{\text{S\&P}, \text{FTSE}, \text{HangSeng}, \text{Nikkei}\}$, test

$$X_t \perp\!\!\!\perp Y_t \mid Z_t \text{ for all times } t \in [n]$$

- ▶ **Retain:** S\&P(t) is independent of HangSeng(t) and Nikkei(t) given FTSE(t)
- ▶ **Reject:** S\&P(t) is independent of FTSE(t) given HangSeng(t) or Nikkei(t)
- ▶ UK has info about Asia-Pacific markets on the same day, but not vice versa

Conditional independence testing reveals structure

- ▶ For each triplet $X, Y, Z \in \{\text{S\&P}, \text{FTSE}, \text{HangSeng}, \text{Nikkei}\}$, test

$$X_t \perp\!\!\!\perp Y_t \mid Z_t \text{ for all times } t \in [n]$$

- ▶ **Retain:** S&P(t) is independent of HangSeng(t) and Nikkei(t) given FTSE(t)
- ▶ **Reject:** S&P(t) is independent of FTSE(t) given HangSeng(t) or Nikkei(t)
- ▶ UK has info about Asia-Pacific markets on the same day, but not vice versa
- ▶ Why? Earliest to latest closing exchanges: Tokyo, Hong Kong, London, New York

Conditional independence testing reveals structure

- ▶ For each triplet $X, Y, Z \in \{\text{S\&P}, \text{FTSE}, \text{HangSeng}, \text{Nikkei}\}$, test

$$X_t \perp\!\!\!\perp Y_t \mid Z_t \text{ for all times } t \in [n]$$

- ▶ **Retain:** S&P(t) is independent of HangSeng(t) and Nikkei(t) given FTSE(t)
- ▶ **Reject:** S&P(t) is independent of FTSE(t) given HangSeng(t) or Nikkei(t)
- ▶ UK has info about Asia-Pacific markets on the same day, but not vice versa
- ▶ Why? Earliest to latest closing exchanges: Tokyo, Hong Kong, London, New York
- ▶ Next, we present the testing procedure

Algorithm (1/2)

- **Inputs:** Data $X_{1:n}$, $Y_{1:n}$, $Z_{1:n}$, significance level α , number of simulations s , $p \in [2, \infty]$ for norm in test statistic, and method for selecting the window size L

Algorithm (1/2)

- ▶ **Inputs:** Data $X_{1:n}$, $Y_{1:n}$, $Z_{1:n}$, significance level α , number of simulations s , $p \in [2, \infty]$ for norm in test statistic, and method for selecting the window size L
- ▶ Regress X_t on Z_t , $t \in [n]$, and obtain residuals $\hat{\varepsilon}_t$, $t \in [n]$

Algorithm (1/2)

- ▶ **Inputs:** Data $X_{1:n}$, $Y_{1:n}$, $Z_{1:n}$, significance level α , number of simulations s , $p \in [2, \infty]$ for norm in test statistic, and method for selecting the window size L
- ▶ Regress X_t on Z_t , $t \in [n]$, and obtain residuals $\hat{\varepsilon}_t$, $t \in [n]$
- ▶ Regress Y_t on Z_t , $t \in [n]$, and obtain residuals $\hat{\xi}_t$, $t \in [n]$

Algorithm (1/2)

- ▶ **Inputs:** Data $X_{1:n}$, $Y_{1:n}$, $Z_{1:n}$, significance level α , number of simulations s , $p \in [2, \infty]$ for norm in test statistic, and method for selecting the window size L
- ▶ Regress X_t on Z_t , $t \in [n]$, and obtain residuals $\hat{\varepsilon}_t$, $t \in [n]$
- ▶ Regress Y_t on Z_t , $t \in [n]$, and obtain residuals $\hat{\xi}_t$, $t \in [n]$
- ▶ Obtain residual products $\hat{R}_t = \text{vec}(\hat{\varepsilon}_t \hat{\xi}_t^\top)$, $t \in [n]$

Algorithm (1/2)

- ▶ **Inputs:** Data $X_{1:n}$, $Y_{1:n}$, $Z_{1:n}$, significance level α , number of simulations s , $p \in [2, \infty]$ for norm in test statistic, and method for selecting the window size L
- ▶ Regress X_t on Z_t , $t \in [n]$, and obtain residuals $\hat{\varepsilon}_t$, $t \in [n]$
- ▶ Regress Y_t on Z_t , $t \in [n]$, and obtain residuals $\hat{\xi}_t$, $t \in [n]$
- ▶ Obtain residual products $\hat{R}_t = \text{vec}(\hat{\varepsilon}_t \hat{\xi}_t^\top)$, $t \in [n]$
- ▶ Calculate test statistic on time series of residual products

$$T(\hat{R}_{1:n}) = \max_{j \in [n]} \left\| \frac{1}{\sqrt{n}} \sum_{t \leq j} \hat{R}_t \right\|_p$$

Algorithm (2/2)

- ▶ Select the window size L for covariance estimation using $\hat{R}_{1:n}$

Algorithm (2/2)

- ▶ Select the window size L for covariance estimation using $\hat{R}_{1:n}$
- ▶ For $t \in [n]$, obtain estimates of the time-varying covariances as

$$\hat{\Sigma}_t = \frac{1}{L} \left(\sum_{j=t-L+1}^t \hat{R}_j \right) \left(\sum_{j=t-L+1}^t \hat{R}_j \right)^\top$$

Algorithm (2/2)

- ▶ Select the window size L for covariance estimation using $\hat{R}_{1:n}$
- ▶ For $t \in [n]$, obtain estimates of the time-varying covariances as

$$\hat{\Sigma}_t = \frac{1}{L} \left(\sum_{j=t-L+1}^t \hat{R}_j \right) \left(\sum_{j=t-L+1}^t \hat{R}_j \right)^\top$$

- ▶ For each $r \in [s]$ and $t \in [n]$, simulate *independent* Gaussians $\tilde{R}_t^{(r)} \sim N(0, \hat{\Sigma}_t)$

Algorithm (2/2)

- ▶ Select the window size L for covariance estimation using $\hat{R}_{1:n}$
- ▶ For $t \in [n]$, obtain estimates of the time-varying covariances as

$$\hat{\Sigma}_t = \frac{1}{L} \left(\sum_{j=t-L+1}^t \hat{R}_j \right) \left(\sum_{j=t-L+1}^t \hat{R}_j \right)^\top$$

- ▶ For each $r \in [s]$ and $t \in [n]$, simulate *independent* Gaussians $\tilde{R}_t^{(r)} \sim N(0, \hat{\Sigma}_t)$
- ▶ For each $r \in [s]$, calculate test statistic on simulated Gaussian process

$$T(\tilde{R}_{1:n}^{(r)}) = \max_{j \in [n]} \left\| \frac{1}{\sqrt{n}} \sum_{t \leq j} \tilde{R}_t^{(r)} \right\|_p$$

Algorithm (2/2)

- ▶ Select the window size L for covariance estimation using $\hat{R}_{1:n}$
- ▶ For $t \in [n]$, obtain estimates of the time-varying covariances as

$$\hat{\Sigma}_t = \frac{1}{L} \left(\sum_{j=t-L+1}^t \hat{R}_j \right) \left(\sum_{j=t-L+1}^t \hat{R}_j \right)^\top$$

- ▶ For each $r \in [s]$ and $t \in [n]$, simulate *independent* Gaussians $\tilde{R}_t^{(r)} \sim N(0, \hat{\Sigma}_t)$
- ▶ For each $r \in [s]$, calculate test statistic on simulated Gaussian process

$$T(\tilde{R}_{1:n}^{(r)}) = \max_{j \in [n]} \left\| \frac{1}{\sqrt{n}} \sum_{t \leq j} \tilde{R}_t^{(r)} \right\|_p$$

- ▶ Calculate the $1 - \alpha$ empirical quantile $\hat{q}_{1-\alpha}^{\text{boot}}$ of $T(\tilde{R}_{1:n}^{(1)}), \dots, T(\tilde{R}_{1:n}^{(s)})$

Algorithm (2/2)

- ▶ Select the window size L for covariance estimation using $\hat{R}_{1:n}$
- ▶ For $t \in [n]$, obtain estimates of the time-varying covariances as

$$\hat{\Sigma}_t = \frac{1}{L} \left(\sum_{j=t-L+1}^t \hat{R}_j \right) \left(\sum_{j=t-L+1}^t \hat{R}_j \right)^\top$$

- ▶ For each $r \in [s]$ and $t \in [n]$, simulate *independent* Gaussians $\tilde{R}_t^{(r)} \sim N(0, \hat{\Sigma}_t)$
- ▶ For each $r \in [s]$, calculate test statistic on simulated Gaussian process

$$T(\tilde{R}_{1:n}^{(r)}) = \max_{j \in [n]} \left\| \frac{1}{\sqrt{n}} \sum_{t \leq j} \tilde{R}_t^{(r)} \right\|_p$$

- ▶ Calculate the $1 - \alpha$ empirical quantile $\hat{q}_{1-\alpha}^{\text{boot}}$ of $T(\tilde{R}_{1:n}^{(1)}), \dots, T(\tilde{R}_{1:n}^{(s)})$
- ▶ Reject null hypothesis at level α if $T(\hat{R}_{1:n}) > \hat{q}_{1-\alpha}^{\text{boot}}$, else retain

The hardness of conditional independence (CI) testing

No-free-lunch in CI testing (Shah and Peters 2020)

The hardness of conditional independence (CI) testing

No-free-lunch in CI testing (Shah and Peters 2020)

- ▶ If a CI test has Type-I error control (i.e. size α), then it cannot have power against any alternative (i.e. power against each alternative $Q \in \mathcal{Q}$ is at most α)

The hardness of conditional independence (CI) testing

No-free-lunch in CI testing (Shah and Peters 2020)

- ▶ If a CI test has Type-I error control (i.e. size α), then it cannot have power against any alternative (i.e. power against each alternative $Q \in \mathcal{Q}$ is at most α)
- ▶ If a CI test has power β against an alternative $Q \in \mathcal{Q}$, then there exists a null distribution $P \in \mathcal{P}$ s.t. the test will reject with prob. greater than or equal to β

The hardness of conditional independence (CI) testing

No-free-lunch in CI testing (Shah and Peters 2020)

- ▶ If a CI test has Type-I error control (i.e. size α), then it cannot have power against any alternative (i.e. power against each alternative $Q \in \mathcal{Q}$ is at most α)
- ▶ If a CI test has power β against an alternative $Q \in \mathcal{Q}$, then there exists a null distribution $P \in \mathcal{P}$ s.t. the test will reject with prob. greater than or equal to β
- ▶ This result has been extended to the time series setting (Bodik and Pasche 2024)

The hardness of conditional independence (CI) testing

No-free-lunch in CI testing (Shah and Peters 2020)

- ▶ If a CI test has Type-I error control (i.e. size α), then it cannot have power against any alternative (i.e. power against each alternative $Q \in \mathcal{Q}$ is at most α)
- ▶ If a CI test has power β against an alternative $Q \in \mathcal{Q}$, then there exists a null distribution $P \in \mathcal{P}$ s.t. the test will reject with prob. greater than or equal to β
- ▶ This result has been extended to the time series setting (Bodik and Pasche 2024)
- ▶ **Implication:** We can only hope to achieve Type-I error control on a subset of null distributions $\mathcal{P}_n^{\text{CI}}$, so must restrict the null hypothesis to make CI testing feasible

The hardness of conditional independence (CI) testing

No-free-lunch in CI testing (Shah and Peters 2020)

- ▶ If a CI test has Type-I error control (i.e. size α), then it cannot have power against any alternative (i.e. power against each alternative $Q \in \mathcal{Q}$ is at most α)
- ▶ If a CI test has power β against an alternative $Q \in \mathcal{Q}$, then there exists a null distribution $P \in \mathcal{P}$ s.t. the test will reject with prob. greater than or equal to β
- ▶ This result has been extended to the time series setting (Bodik and Pasche 2024)
- ▶ **Implication:** We can only hope to achieve Type-I error control on a subset of null distributions $\mathcal{P}_n^{\text{CI}}$, so must restrict the null hypothesis to make CI testing feasible
- ▶ **Key assumption:** uniform decay of physical dependence measure (Wu 2005)

Key assumption: Representation in terms of noise inputs (Wu 2005)

- ▶ Let $(\epsilon_i)_{i \in \mathbb{Z}}$ be an iid sequence of $U[0, 1]$ random seeds

Key assumption: Representation in terms of noise inputs (Wu 2005)

- ▶ Let $(\epsilon_i)_{i \in \mathbb{Z}}$ be an iid sequence of $U[0, 1]$ random seeds
- ▶ Define the sequence of noise inputs up to time t by $\epsilon_t = (\epsilon_t, \epsilon_{t-1}, \dots)$

Key assumption: Representation in terms of noise inputs (Wu 2005)

- ▶ Let $(\epsilon_i)_{i \in \mathbb{Z}}$ be an iid sequence of $U[0, 1]$ random seeds
- ▶ Define the sequence of noise inputs up to time t by $\epsilon_t = (\epsilon_t, \epsilon_{t-1}, \dots)$
- ▶ Let Θ be a parameter space, which can be infinite-dimensional

Key assumption: Representation in terms of noise inputs (Wu 2005)

- ▶ Let $(\epsilon_i)_{i \in \mathbb{Z}}$ be an iid sequence of $U[0, 1]$ random seeds
- ▶ Define the sequence of noise inputs up to time t by $\epsilon_t = (\epsilon_t, \epsilon_{t-1}, \dots)$
- ▶ Let Θ be a parameter space, which can be infinite-dimensional
- ▶ Let $G_t^{(n)} : \mathbb{R}^\infty \times \Theta \rightarrow \mathbb{R}^d$ be some measurable function, $t, n, d \in \mathbb{N}$

Key assumption: Representation in terms of noise inputs (Wu 2005)

- ▶ Given a parameter $\theta \in \Theta$, index $n \in \mathbb{N}$ (linked to sample size and approximation), and noise inputs ϵ_t , a time series $X_{1:n}$ is *generated by nature* as

$$X_t = G_t^{(n)}(\epsilon_t, \theta), \quad t = 1, \dots, n$$

Key assumption: Representation in terms of noise inputs (Wu 2005)

- ▶ Given a parameter $\theta \in \Theta$, index $n \in \mathbb{N}$ (linked to sample size and approximation), and noise inputs ϵ_t , a time series $X_{1:n}$ is *generated by nature* as

$$X_t = G_t^{(n)}(\epsilon_t, \theta), \quad t = 1, \dots, n$$

- ▶ We **do not know** the generative model $G_t^{(n)}$, noise inputs ϵ_t , parameter θ

Key assumption: Representation in terms of noise inputs (Wu 2005)

- ▶ Given a parameter $\theta \in \Theta$, index $n \in \mathbb{N}$ (linked to sample size and approximation), and noise inputs ϵ_t , a time series $X_{1:n}$ is *generated by nature* as

$$X_t = G_t^{(n)}(\epsilon_t, \theta), \quad t = 1, \dots, n$$

- ▶ We **do not know** the generative model $G_t^{(n)}$, noise inputs ϵ_t , parameter θ
- ▶ Let $P_{\theta,n}$ be the (finite-dimensional) distribution of $(G_t^{(n)}(\epsilon_t, \theta))_{t \in [n]}$

Key assumption: Representation in terms of noise inputs (Wu 2005)

- ▶ Given a parameter $\theta \in \Theta$, index $n \in \mathbb{N}$ (linked to sample size and approximation), and noise inputs ϵ_t , a time series $X_{1:n}$ is *generated by nature* as

$$X_t = G_t^{(n)}(\epsilon_t, \theta), \quad t = 1, \dots, n$$

- ▶ We **do not know** the generative model $G_t^{(n)}$, noise inputs ϵ_t , parameter θ
- ▶ Let $P_{\theta,n}$ be the (finite-dimensional) distribution of $(G_t^{(n)}(\epsilon_t, \theta))_{t \in [n]}$
- ▶ Let P_θ be the law of $(G_t(\epsilon_t, \theta))_{t \in \mathbb{N}}$ for some limiting mapping G_t

Key assumption: Representation in terms of noise inputs (Wu 2005)

- ▶ Given a parameter $\theta \in \Theta$, index $n \in \mathbb{N}$ (linked to sample size and approximation), and noise inputs ϵ_t , a time series $X_{1:n}$ is *generated by nature* as

$$X_t = G_t^{(n)}(\epsilon_t, \theta), \quad t = 1, \dots, n$$

- ▶ We **do not know** the generative model $G_t^{(n)}$, noise inputs ϵ_t , parameter θ
- ▶ Let $P_{\theta,n}$ be the (finite-dimensional) distribution of $(G_t^{(n)}(\epsilon_t, \theta))_{t \in [n]}$
- ▶ Let P_θ be the law of $(G_t(\epsilon_t, \theta))_{t \in \mathbb{N}}$ for some limiting mapping G_t
- ▶ Denote the collections by $\mathcal{P}_n = \{P_{\theta,n} : \theta \in \Theta\}$ and $\mathcal{P} = \{P_\theta : \theta \in \Theta\}$

Key assumption: Decaying influence of noise inputs (Wu 2005)

- Let $(\epsilon_i^*)_{i \in \mathbb{Z}}$ be an iid copy of $(\epsilon_i)_{i \in \mathbb{Z}}$. For $\ell \in \mathbb{Z}$, $j \in \mathbb{N}_0$, denote

$$\tilde{\epsilon}_{\ell,j} = (\epsilon_\ell, \dots, \epsilon_{\ell-j+1}, \epsilon_{\ell-j}^*, \epsilon_{\ell-j-1}, \dots)$$

Key assumption: Decaying influence of noise inputs (Wu 2005)

- Let $(\epsilon_i^*)_{i \in \mathbb{Z}}$ be an iid copy of $(\epsilon_i)_{i \in \mathbb{Z}}$. For $\ell \in \mathbb{Z}$, $j \in \mathbb{N}_0$, denote

$$\tilde{\epsilon}_{\ell,j} = (\epsilon_\ell, \dots, \epsilon_{\ell-j+1}, \epsilon_{\ell-j}^*, \epsilon_{\ell-j-1}, \dots)$$

- There exist $\Psi, \rho > 0$ such that, for some $q > 2$ and all $j \in \mathbb{N}$, we have

$$\sup_{\theta \in \Theta} \sup_{\ell \in \mathbb{Z}} \sup_{i \in \mathbb{N}} \sup_{t \in \mathbb{N}} \left\| G_t^{(i)}(\epsilon_\ell, \theta) - G_t^{(i)}(\tilde{\epsilon}_{\ell,j}, \theta) \right\|_{L^q(\theta)} \leq \Psi(j \vee 1)^{-\rho}$$

$$\sup_{\theta \in \Theta} \sup_{\ell \in \mathbb{Z}} \sup_{i \in \mathbb{N}} \sup_{t \in \mathbb{N}} \left\| G_t^{(i)}(\epsilon_\ell, \theta) \right\|_{L^q(\theta)} \leq \Psi$$

where $\|\cdot\|_{L^q(\theta)} = \mathbb{E}_\theta (\|\cdot\|_2^q)^{1/q}$, $\mathbb{E}_\theta(\cdot)$ is expectation w.r.t. θ -dependent distribution

Key assumption: Decaying influence of noise inputs (Wu 2005)

- ▶ Let $(\epsilon_i^*)_{i \in \mathbb{Z}}$ be an iid copy of $(\epsilon_i)_{i \in \mathbb{Z}}$. For $\ell \in \mathbb{Z}$, $j \in \mathbb{N}_0$, denote

$$\tilde{\epsilon}_{\ell,j} = (\epsilon_\ell, \dots, \epsilon_{\ell-j+1}, \epsilon_{\ell-j}^*, \epsilon_{\ell-j-1}, \dots)$$

- ▶ There exist $\Psi, \rho > 0$ such that, for some $q > 2$ and all $j \in \mathbb{N}$, we have

$$\sup_{\theta \in \Theta} \sup_{\ell \in \mathbb{Z}} \sup_{i \in \mathbb{N}} \sup_{t \in \mathbb{N}} \left\| G_t^{(i)}(\epsilon_\ell, \theta) - G_t^{(i)}(\tilde{\epsilon}_{\ell,j}, \theta) \right\|_{L^q(\theta)} \leq \Psi(j \vee 1)^{-\rho}$$

$$\sup_{\theta \in \Theta} \sup_{\ell \in \mathbb{Z}} \sup_{i \in \mathbb{N}} \sup_{t \in \mathbb{N}} \left\| G_t^{(i)}(\epsilon_\ell, \theta) \right\|_{L^q(\theta)} \leq \Psi$$

where $\|\cdot\|_{L^q(\theta)} = \mathbb{E}_\theta (\|\cdot\|_2^q)^{1/q}$, $\mathbb{E}_\theta(\cdot)$ is expectation w.r.t. θ -dependent distribution

- ▶ We make this assumption for $X_{1:n}$, $Y_{1:n}$, and $Z_{1:n}$

We prove our test has uniformly asymptotic Type I error control

- For each $n \in \mathbb{N}$, let $\mathcal{P}_n^* \subset \mathcal{P}_n^{\text{CI}}$ be a collection of joint distributions for $X_{1:n}$, $Y_{1:n}$, and $Z_{1:n}$ for which the null hypothesis holds *and all assumptions are satisfied*

We prove our test has uniformly asymptotic Type I error control

- ▶ For each $n \in \mathbb{N}$, let $\mathcal{P}_n^* \subset \mathcal{P}_n^{\text{CI}}$ be a collection of joint distributions for $X_{1:n}$, $Y_{1:n}$, and $Z_{1:n}$ for which the null hypothesis holds *and all assumptions are satisfied*
- ▶ Need sufficiently fast convergence rates for the regression estimators

We prove our test has uniformly asymptotic Type I error control

- ▶ For each $n \in \mathbb{N}$, let $\mathcal{P}_n^* \subset \mathcal{P}_n^{\text{CI}}$ be a collection of joint distributions for $X_{1:n}$, $Y_{1:n}$, and $Z_{1:n}$ for which the null hypothesis holds *and all assumptions are satisfied*
- ▶ Need sufficiently fast convergence rates for the regression estimators
- ▶ For all distributions $P \in \mathcal{P}_n^*$, the **product** of $L^2(P)$ norms of the estimation errors must be $o\left(\frac{1}{\sqrt{n} \text{polylog}(n)}\right)$, but each $L^2(P)$ norm only needs to be $o\left(\frac{1}{\text{polylog}(n)}\right)$

We prove our test has uniformly asymptotic Type I error control

- ▶ For each $n \in \mathbb{N}$, let $\mathcal{P}_n^* \subset \mathcal{P}_n^{\text{CI}}$ be a collection of joint distributions for $X_{1:n}$, $Y_{1:n}$, and $Z_{1:n}$ for which the null hypothesis holds *and all assumptions are satisfied*
- ▶ Need sufficiently fast convergence rates for the regression estimators
- ▶ For all distributions $P \in \mathcal{P}_n^*$, the **product** of $L^2(P)$ norms of the estimation errors must be $o\left(\frac{1}{\sqrt{n} \text{polylog}(n)}\right)$, but each $L^2(P)$ norm only needs to be $o\left(\frac{1}{\text{polylog}(n)}\right)$

Theorem (Informal)

Recall the test statistic $T(\hat{R}_{1:n})$ and estimated quantile $\hat{q}_{1-\alpha}^{\text{boot}}$. We have

$$\limsup_{n \rightarrow \infty} \sup_{P \in \mathcal{P}_n^*} \mathbb{P}_P \left(T(\hat{R}_{1:n}) > \hat{q}_{1-\alpha}^{\text{boot}} \right) \leq \alpha.$$

Simulation-based estimation and inference

- ▶ Suppose we observe a time series $X_{1:n}^{\text{obs}}$ taking values in $\mathbb{R}^d$

Simulation-based estimation and inference

- ▶ Suppose we observe a time series $X_{1:n}^{\text{obs}}$ taking values in $\mathbb{R}^d$
- ▶ We assume the observed time series has been *generated by nature* as

$$X_t^{\text{obs}} = G_t^{(n)}(\epsilon_t, \theta_0), \quad t = 1, \dots, n$$

at some unknown parameter $\theta_0 \in \Theta \subset \mathbb{R}^p$ with some unknown noise inputs ϵ_t

Simulation-based estimation and inference

- ▶ Suppose we observe a time series $X_{1:n}^{\text{obs}}$ taking values in $\mathbb{R}^d$
- ▶ We assume the observed time series has been *generated by nature* as

$$X_t^{\text{obs}} = G_t^{(n)}(\epsilon_t, \theta_0), \quad t = 1, \dots, n$$

at some unknown parameter $\theta_0 \in \Theta \subset \mathbb{R}^p$ with some unknown noise inputs ϵ_t

- ▶ We want to estimate θ_0 , get a confidence set, or test goodness-of-fit

Simulation-based estimation and inference

- ▶ Suppose we observe a time series $X_{1:n}^{\text{obs}}$ taking values in $\mathbb{R}^d$
- ▶ We assume the observed time series has been *generated by nature* as

$$X_t^{\text{obs}} = G_t^{(n)}(\epsilon_t, \theta_0), \quad t = 1, \dots, n$$

at some unknown parameter $\theta_0 \in \Theta \subset \mathbb{R}^p$ with some unknown noise inputs ϵ_t

- ▶ We want to estimate θ_0 , get a confidence set, or test goodness-of-fit
- ▶ The likelihood is intractable, so we use a simulation-based approach

Simulation-based estimation and inference

- ▶ Suppose we observe a time series $X_{1:n}^{\text{obs}}$ taking values in $\mathbb{R}^d$
- ▶ We assume the observed time series has been *generated by nature* as

$$X_t^{\text{obs}} = G_t^{(n)}(\epsilon_t, \theta_0), \quad t = 1, \dots, n$$

at some unknown parameter $\theta_0 \in \Theta \subset \mathbb{R}^p$ with some unknown noise inputs ϵ_t

- ▶ We want to estimate θ_0 , get a confidence set, or test goodness-of-fit
- ▶ The likelihood is intractable, so we use a simulation-based approach
- ▶ **Key assumption:** We know $G_t^{(n)}$ and can simulate from it at any value of $\theta \in \Theta$

How simulation-based estimation is usually done

- ▶ **Rough idea:** Tune the parameters θ until some features F “look right”

How simulation-based estimation is usually done

- ▶ **Rough idea:** Tune the parameters θ until some features F “look right”
- ▶ For (asymptotically) stationary time series, often use minimum distance estimator

$$\hat{\theta} = \operatorname{argmin}_{\theta \in \Theta} \left\| F^{\text{obs}} - \frac{1}{s} \sum_{r=1}^s F^{(r)}(\theta) \right\|^2$$

How simulation-based estimation is usually done

- ▶ **Rough idea:** Tune the parameters θ until some features F “look right”
- ▶ For (asymptotically) stationary time series, often use minimum distance estimator

$$\hat{\theta} = \operatorname{argmin}_{\theta \in \Theta} \left\| F^{\text{obs}} - \frac{1}{s} \sum_{r=1}^s F^{(r)}(\theta) \right\|^2$$

- ▶ For consistency $\hat{\theta} \xrightarrow{P} \theta_0$, need $F(\theta) \xrightarrow{P} \Phi(\theta)$ for all θ

How simulation-based estimation is usually done

- ▶ **Rough idea:** Tune the parameters θ until some features F “look right”
- ▶ For (asymptotically) stationary time series, often use minimum distance estimator

$$\hat{\theta} = \operatorname{argmin}_{\theta \in \Theta} \left\| F^{\text{obs}} - \frac{1}{s} \sum_{r=1}^s F^{(r)}(\theta) \right\|^2$$

- ▶ For consistency $\hat{\theta} \xrightarrow{P} \theta_0$, need $F(\theta) \xrightarrow{P} \Phi(\theta)$ for all θ
- ▶ Hope Φ^{-1} exists and is nice, so we can translate from $\Phi(\theta)$ back to θ

How simulation-based estimation is usually done

- ▶ **Rough idea:** Tune the parameters θ until some features F “look right”
- ▶ For (asymptotically) stationary time series, often use minimum distance estimator

$$\hat{\theta} = \operatorname{argmin}_{\theta \in \Theta} \left\| F^{\text{obs}} - \frac{1}{s} \sum_{r=1}^s F^{(r)}(\theta) \right\|^2$$

- ▶ For consistency $\hat{\theta} \xrightarrow{P} \theta_0$, need $F(\theta) \xrightarrow{P} \Phi(\theta)$ for all θ
- ▶ Hope Φ^{-1} exists and is nice, so we can translate from $\Phi(\theta)$ back to θ
- ▶ For nonstationary time series, need rolling-window estimator (not discussed here)

Which features should we use?

- ▶ **Method of simulated moments:** F based on functions of moments

Which features should we use?

- ▶ **Method of simulated moments:** F based on functions of moments
 - ▶ Sample mean, variance, skewness, autocorrelation, and so on

Which features should we use?

- ▶ **Method of simulated moments:** F based on functions of moments
 - ▶ Sample mean, variance, skewness, autocorrelation, and so on
- ▶ **Indirect inference:** F from parameter estimates in auxiliary model

Which features should we use?

- ▶ **Method of simulated moments:** F based on functions of moments
 - ▶ Sample mean, variance, skewness, autocorrelation, and so on
- ▶ **Indirect inference:** F from parameter estimates in auxiliary model
 - ▶ To get $\hat{\theta}$ for MA(1): $X_t = \epsilon_t + \theta_0 \epsilon_{t-1}$

Which features should we use?

- ▶ **Method of simulated moments:** F based on functions of moments
 - ▶ Sample mean, variance, skewness, autocorrelation, and so on
- ▶ **Indirect inference:** F from parameter estimates in auxiliary model
 - ▶ To get $\hat{\theta}$ for MA(1): $X_t = \epsilon_t + \theta_0 \epsilon_{t-1}$
 - ▶ Use $F = \hat{\beta}$ from AR(1): $X_t = \beta X_{t-1} + \xi_t$

Which features should we use?

- ▶ **Method of simulated moments:** F based on functions of moments
 - ▶ Sample mean, variance, skewness, autocorrelation, and so on
- ▶ **Indirect inference:** F from parameter estimates in auxiliary model
 - ▶ To get $\hat{\theta}$ for MA(1): $X_t = \epsilon_t + \theta_0 \epsilon_{t-1}$
 - ▶ Use $F = \hat{\beta}$ from AR(1): $X_t = \beta X_{t-1} + \xi_t$
- ▶ **Synthetic likelihood:** F from user-chosen summary statistics

Which features should we use?

- ▶ **Method of simulated moments:** F based on functions of moments
 - ▶ Sample mean, variance, skewness, autocorrelation, and so on
- ▶ **Indirect inference:** F from parameter estimates in auxiliary model
 - ▶ To get $\hat{\theta}$ for MA(1): $X_t = \epsilon_t + \theta_0 \epsilon_{t-1}$
 - ▶ Use $F = \hat{\beta}$ from AR(1): $X_t = \beta X_{t-1} + \xi_t$
- ▶ **Synthetic likelihood:** F from user-chosen summary statistics
 - ▶ Coefficients from polynomial regression of X_t on $X_{t-1}, \dots, X_{t-\ell}$

Which features should we use?

- ▶ **Method of simulated moments:** F based on functions of moments
 - ▶ Sample mean, variance, skewness, autocorrelation, and so on
- ▶ **Indirect inference:** F from parameter estimates in auxiliary model
 - ▶ To get $\hat{\theta}$ for MA(1): $X_t = \epsilon_t + \theta_0 \epsilon_{t-1}$
 - ▶ Use $F = \hat{\beta}$ from AR(1): $X_t = \beta X_{t-1} + \xi_t$
- ▶ **Synthetic likelihood:** F from user-chosen summary statistics
 - ▶ Coefficients from polynomial regression of X_t on $X_{t-1}, \dots, X_{t-\ell}$
- ▶ **Neural summary statistics:** F from neural network

Which features should we use?

- ▶ **Method of simulated moments:** F based on functions of moments
 - ▶ Sample mean, variance, skewness, autocorrelation, and so on
- ▶ **Indirect inference:** F from parameter estimates in auxiliary model
 - ▶ To get $\hat{\theta}$ for MA(1): $X_t = \epsilon_t + \theta_0 \epsilon_{t-1}$
 - ▶ Use $F = \hat{\beta}$ from AR(1): $X_t = \beta X_{t-1} + \xi_t$
- ▶ **Synthetic likelihood:** F from user-chosen summary statistics
 - ▶ Coefficients from polynomial regression of X_t on $X_{t-1}, \dots, X_{t-\ell}$
- ▶ **Neural summary statistics:** F from neural network
 - ▶ Extract early layer of neural network

Which features should we use?

- ▶ Good features are sensitive to small parameter changes

Which features should we use?

- ▶ Good features are sensitive to small parameter changes
- ▶ Takes time to find good features

Which features should we use?

- ▶ Good features are sensitive to small parameter changes
- ▶ Takes time to find good features
- ▶ Can we just use *random features* of the data instead?

Which features should we use?

- ▶ Good features are sensitive to small parameter changes
- ▶ Takes time to find good features
- ▶ Can we just use *random features* of the data instead?
- ▶ Automatic feature selection, computationally efficient, little knowledge needed

How many random features?

- **Answer:** Need $2p + 1$ random features for a p -dimensional parameter space Θ

How many random features?

- ▶ **Answer:** Need $2p + 1$ random features for a p -dimensional parameter space Θ
- ▶ Probabilistic Whitney (1936) embedding theorem due to Sauer, J. A. Yorke, and Casdagli (1991)

How many random features?

- ▶ **Answer:** Need $2p + 1$ random features for a p -dimensional parameter space Θ
- ▶ Probabilistic Whitney (1936) embedding theorem due to Sauer, J. A. Yorke, and Casdagli (1991)
- ▶ For “almost every” C^1 smooth function Φ from $\mathbb{R}^p$ to $\mathbb{R}^{2p+1}$:
 1. Φ **is one-to-one** on compact subsets of $\mathbb{R}^p$
 2. Φ **has a nice derivative** on certain compact subsets of $\mathbb{R}^p$
 - Details: the $(2p + 1) \times p$ Jacobian matrix of Φ has full rank p on each compact subset of a smooth manifold contained in $\mathbb{R}^p$

How many random features?

- ▶ Let P_θ be the joint distribution of $[X_t(\theta), \dots, X_{t-m}(\theta)]^\top$ for some $m \in \mathbb{N}$

How many random features?

- ▶ Let P_θ be the joint distribution of $[X_t(\theta), \dots, X_{t-m}(\theta)]^\top$ for some $m \in \mathbb{N}$
- ▶ Under smoothness conditions on $\theta \mapsto P_\theta$, the expectations of $2p + 1$ “generic” continuous, bounded functions $\varphi : \mathbb{R}^{(m+1) \times d} \rightarrow \mathbb{R}^{2p+1}$ will be C^1 smooth in θ

How many random features?

- ▶ Let P_θ be the joint distribution of $[X_t(\theta), \dots, X_{t-m}(\theta)]^\top$ for some $m \in \mathbb{N}$
- ▶ Under smoothness conditions on $\theta \mapsto P_\theta$, the expectations of $2p+1$ “generic” continuous, bounded functions $\varphi : \mathbb{R}^{(m+1) \times d} \rightarrow \mathbb{R}^{2p+1}$ will be C^1 smooth in θ
- ▶ We define a C^1 smooth function $\Phi : \mathbb{R}^p \rightarrow \mathbb{R}^{2p+1}$ by *extending*

$$\theta \mapsto \mathbb{E}_{P_\theta} \left[\varphi \left([X_t(\theta), \dots, X_{t-m}(\theta)]^\top \right) \right]$$

from the compact subset $\Theta \subset \mathbb{R}^p$ to $\mathbb{R}^p$

How many random features?

- ▶ Let P_θ be the joint distribution of $[X_t(\theta), \dots, X_{t-m}(\theta)]^\top$ for some $m \in \mathbb{N}$
- ▶ Under smoothness conditions on $\theta \mapsto P_\theta$, the expectations of $2p+1$ “generic” continuous, bounded functions $\varphi : \mathbb{R}^{(m+1) \times d} \rightarrow \mathbb{R}^{2p+1}$ will be C^1 smooth in θ
- ▶ We define a C^1 smooth function $\Phi : \mathbb{R}^p \rightarrow \mathbb{R}^{2p+1}$ by *extending*

$$\theta \mapsto \mathbb{E}_{P_\theta} \left[\varphi \left([X_t(\theta), \dots, X_{t-m}(\theta)]^\top \right) \right]$$

from the compact subset $\Theta \subset \mathbb{R}^p$ to $\mathbb{R}^p$

- ▶ Unless we are very unlucky in our selection of φ , then Φ will be one-to-one on Θ and have a nice derivative

How many random features?

- ▶ Let P_θ be the joint distribution of $[X_t(\theta), \dots, X_{t-m}(\theta)]^\top$ for some $m \in \mathbb{N}$
- ▶ Under smoothness conditions on $\theta \mapsto P_\theta$, the expectations of $2p + 1$ “generic” continuous, bounded functions $\varphi : \mathbb{R}^{(m+1) \times d} \rightarrow \mathbb{R}^{2p+1}$ will be C^1 smooth in θ
- ▶ We define a C^1 smooth function $\Phi : \mathbb{R}^p \rightarrow \mathbb{R}^{2p+1}$ by *extending*

$$\theta \mapsto \mathbb{E}_{P_\theta} \left[\varphi \left([X_t(\theta), \dots, X_{t-m}(\theta)]^\top \right) \right]$$

from the compact subset $\Theta \subset \mathbb{R}^p$ to $\mathbb{R}^p$

- ▶ Unless we are very unlucky in our selection of φ , then Φ will be one-to-one on Θ and have a nice derivative
- ▶ If physical dependence measure (Wu 2005) decays uniformly over Θ , we can estimate $\Phi(\theta)$, $\theta \in \Theta$, via time-averages of $2p + 1$ random features

$$\frac{1}{n-m} \sum_{t=m+1}^n \varphi \left([X_t(\theta), \dots, X_{t-m}(\theta)]^\top \right)$$

Example: Mean of Gaussian

- **Goal:** Estimate $\theta_0 = 0$ given $X_1^{\text{obs}}, \dots, X_n^{\text{obs}} \stackrel{\text{iid}}{\sim} \mathcal{N}(\theta_0, 1)$ with $\Theta = [-3, 3]$

Example: Mean of Gaussian

- ▶ **Goal:** Estimate $\theta_0 = 0$ given $X_1^{\text{obs}}, \dots, X_n^{\text{obs}} \stackrel{\text{iid}}{\sim} \mathcal{N}(\theta_0, 1)$ with $\Theta = [-3, 3]$
- ▶ Use time-averages of $2p + 1 = 3$ random Fourier features with $m = 0$ lags

Example: Mean of Gaussian

- ▶ **Goal:** Estimate $\theta_0 = 0$ given $X_1^{\text{obs}}, \dots, X_n^{\text{obs}} \stackrel{\text{iid}}{\sim} \mathcal{N}(\theta_0, 1)$ with $\Theta = [-3, 3]$
- ▶ Use time-averages of $2p + 1 = 3$ random Fourier features with $m = 0$ lags
- ▶ For $i = 1, \dots, 2p + 1$, randomly draw frequencies $w_i \stackrel{\text{iid}}{\sim} \mathcal{N}(0, 1)$ and phases $b_i \stackrel{\text{iid}}{\sim} \text{Unif}(-\pi, \pi)$, and define the i -th feature as $\varphi_i(x) = \cos(w_i x + b_i)$

Example: Mean of Gaussian

- ▶ **Goal:** Estimate $\theta_0 = 0$ given $X_1^{\text{obs}}, \dots, X_n^{\text{obs}} \stackrel{\text{iid}}{\sim} \mathcal{N}(\theta_0, 1)$ with $\Theta = [-3, 3]$
- ▶ Use time-averages of $2p + 1 = 3$ random Fourier features with $m = 0$ lags
- ▶ For $i = 1, \dots, 2p + 1$, randomly draw frequencies $w_i \stackrel{\text{iid}}{\sim} \mathcal{N}(0, 1)$ and phases $b_i \stackrel{\text{iid}}{\sim} \text{Unif}(-\pi, \pi)$, and define the i -th feature as $\varphi_i(x) = \cos(w_i x + b_i)$
- ▶ Denoting $\varphi = (\varphi_1, \dots, \varphi_{2p+1})$, our estimator is given by

$$\hat{\theta} = \underset{\theta \in \Theta}{\operatorname{argmin}} \left\| \frac{1}{n} \sum_{t=1}^n \varphi(X_t^{\text{obs}}) - \frac{1}{s} \sum_{r=1}^s \left(\frac{1}{n} \sum_{t=1}^n \varphi(X_t^{(r)}(\theta)) \right) \right\|_2$$


Example: Mean of Gaussian

- ▶ **Goal:** Estimate $\theta_0 = 0$ given $X_1^{\text{obs}}, \dots, X_n^{\text{obs}} \stackrel{\text{iid}}{\sim} \mathcal{N}(\theta_0, 1)$ with $\Theta = [-3, 3]$
- ▶ Use time-averages of $2p + 1 = 3$ random Fourier features with $m = 0$ lags
- ▶ For $i = 1, \dots, 2p + 1$, randomly draw frequencies $w_i \stackrel{\text{iid}}{\sim} \mathcal{N}(0, 1)$ and phases $b_i \stackrel{\text{iid}}{\sim} \text{Unif}(-\pi, \pi)$, and define the i -th feature as $\varphi_i(x) = \cos(w_i x + b_i)$
- ▶ Denoting $\varphi = (\varphi_1, \dots, \varphi_{2p+1})$, our estimator is given by


$$\hat{\theta} = \underset{\theta \in \Theta}{\operatorname{argmin}} \left\| \frac{1}{n} \sum_{t=1}^n \varphi(X_t^{\text{obs}}) - \frac{1}{s} \sum_{r=1}^s \left(\frac{1}{n} \sum_{t=1}^n \varphi(X_t^{(r)}(\theta)) \right) \right\|_2$$

- ▶ We will compare MSE our random feature estimator $\hat{\theta}$ with $\hat{\theta}^{\text{MLE}} = \frac{1}{n} \sum_{t=1}^n X_t$


Time-average of 3 features, colored by parameter value ($n = 1000$, $s = 10$)



Density plot based on 1,000 independent estimates



Theoretical MSE of MLE vs MSE of our estimator (100 replicates per n)



Theoretical guarantees for random feature estimators

- We prove consistency $\hat{\theta} \xrightarrow{P} \theta_0$ and asymptotic normality

$$\sqrt{n}(\hat{\theta} - \theta_0) \Rightarrow N\left(0, (\Gamma^\top \Gamma)^{-1} \Gamma^\top V \Gamma (\Gamma^\top \Gamma)^{-1}\right)$$

where Γ is the $(2p+1) \times p$ Jacobian matrix of Φ evaluated at θ_0 , and V is the $(2p+1) \times (2p+1)$ asymptotic covariance matrix of $\frac{1}{n} \sum_{t=1}^n \varphi(X_t^{\text{obs}})$

Theoretical guarantees for random feature estimators

- We prove consistency $\hat{\theta} \xrightarrow{P} \theta_0$ and asymptotic normality

$$\sqrt{n}(\hat{\theta} - \theta_0) \Rightarrow N\left(0, (\Gamma^\top \Gamma)^{-1} \Gamma^\top V \Gamma (\Gamma^\top \Gamma)^{-1}\right)$$

where Γ is the $(2p+1) \times p$ Jacobian matrix of Φ evaluated at θ_0 , and V is the $(2p+1) \times (2p+1)$ asymptotic covariance matrix of $\frac{1}{n} \sum_{t=1}^n \varphi(X_t^{\text{obs}})$

- Rolling-window estimator also introduced

Theoretical guarantees for random feature estimators

- ▶ We prove consistency $\hat{\theta} \xrightarrow{P} \theta_0$ and asymptotic normality

$$\sqrt{n}(\hat{\theta} - \theta_0) \Rightarrow N\left(0, (\Gamma^\top \Gamma)^{-1} \Gamma^\top V \Gamma (\Gamma^\top \Gamma)^{-1}\right)$$

where Γ is the $(2p+1) \times p$ Jacobian matrix of Φ evaluated at θ_0 , and V is the $(2p+1) \times (2p+1)$ asymptotic covariance matrix of $\frac{1}{n} \sum_{t=1}^n \varphi(X_t^{\text{obs}})$

- ▶ Rolling-window estimator also introduced
- ▶ Analogous guarantees using framework for nonstationary time series (Mies 2023)

Theoretical guarantees for random feature estimators

- ▶ We prove consistency $\hat{\theta} \xrightarrow{P} \theta_0$ and asymptotic normality

$$\sqrt{n} \left(\hat{\theta} - \theta_0 \right) \Rightarrow N \left(0, (\Gamma^\top \Gamma)^{-1} \Gamma^\top V \Gamma (\Gamma^\top \Gamma)^{-1} \right)$$

where Γ is the $(2p+1) \times p$ Jacobian matrix of Φ evaluated at θ_0 , and V is the $(2p+1) \times (2p+1)$ asymptotic covariance matrix of $\frac{1}{n} \sum_{t=1}^n \varphi(X_t^{\text{obs}})$

- ▶ Rolling-window estimator also introduced
- ▶ Analogous guarantees using framework for nonstationary time series (Mies 2023)
- ▶ Examples: ODEs with dependent observational noise, nonstationary state-space models, and structural time series models

Paradigm shift in prediction: Time series foundation models (TSFMs)

- ▶ Inspired by success of transformer-based foundation models in language and vision

Paradigm shift in prediction: Time series foundation models (TSFMs)

- ▶ Inspired by success of transformer-based foundation models in language and vision
- ▶ Many TSFMs are *pre-trained* on synthetic time series generated by *parametric* time series models or *convex combinations* of such models

Paradigm shift in prediction: Time series foundation models (TSFMs)

- ▶ Inspired by success of transformer-based foundation models in language and vision
- ▶ Many TSFMs are *pre-trained* on synthetic time series generated by *parametric* time series models or *convex combinations* of such models
- ▶ Statistical foundations are unclear, especially for nonstationary time series

Paradigm shift in prediction: Time series foundation models (TSFMs)

- ▶ Inspired by success of transformer-based foundation models in language and vision
- ▶ Many TSFMs are *pre-trained* on synthetic time series generated by *parametric* time series models or *convex combinations* of such models
- ▶ Statistical foundations are unclear, especially for nonstationary time series
- ▶ We develop theory, and propose methods for smoothing and forecasting

Simulate to get pre-training data for forecasting with TSFMs

- For forecasting with TSFMs, Dooley et al. (2023) suggest training on simulations from *multiplicative seasonality-trend-noise decomposition models* (many $\theta \in \Theta$)

Simulate to get pre-training data for forecasting with TSFMs

- ▶ For forecasting with TSFMs, Dooley et al. (2023) suggest training on simulations from *multiplicative seasonality-trend-noise decomposition models* (many $\theta \in \Theta$)
- ▶ We consider an extension with multiplicative autoregressive noise

$$Y_t = \text{Se} \left(t, \theta^{\text{Se}} \right) \text{Tr} \left(t, \theta^{\text{Tr}} \right) \exp \left(\kappa X_t - \frac{\kappa^2 \sigma_X^2}{2(1 - \phi^2)} \right) \exp \left(\varepsilon_t^Y - \frac{\sigma_Y^2}{2} \right)$$
$$X_t = \phi X_{t-1} + \varepsilon_t^X$$

for some scaling factor $\kappa > 0$ and autoregressive coefficient $\phi \in (-1, 1)$, with $X_0 \sim N(0, \frac{\sigma_X^2}{1 - \phi^2})$ and $\varepsilon_t^X \stackrel{\text{iid}}{\sim} N(0, \sigma_X^2)$, $\varepsilon_t^Y \stackrel{\text{iid}}{\sim} N(0, \sigma_Y^2)$ for some $\sigma_X, \sigma_Y > 0$

Simulate to get pre-training data for forecasting with TSFMs

- ▶ For forecasting with TSFMs, Dooley et al. (2023) suggest training on simulations from *multiplicative seasonality-trend-noise decomposition models* (many $\theta \in \Theta$)
- ▶ We consider an extension with multiplicative autoregressive noise

$$Y_t = \text{Se} \left(t, \theta^{\text{Se}} \right) \text{Tr} \left(t, \theta^{\text{Tr}} \right) \exp \left(\kappa X_t - \frac{\kappa^2 \sigma_X^2}{2(1 - \phi^2)} \right) \exp \left(\varepsilon_t^Y - \frac{\sigma_Y^2}{2} \right)$$
$$X_t = \phi X_{t-1} + \varepsilon_t^X$$

for some scaling factor $\kappa > 0$ and autoregressive coefficient $\phi \in (-1, 1)$, with $X_0 \sim N(0, \frac{\sigma_X^2}{1 - \phi^2})$ and $\varepsilon_t^X \stackrel{\text{iid}}{\sim} N(0, \sigma_X^2)$, $\varepsilon_t^Y \stackrel{\text{iid}}{\sim} N(0, \sigma_Y^2)$ for some $\sigma_X, \sigma_Y > 0$

- ▶ Seasonality: linear combination of sines, cosines (weekly, monthly, yearly periods)

Simulate to get pre-training data for forecasting with TSFMs

- ▶ For forecasting with TSFMs, Dooley et al. (2023) suggest training on simulations from *multiplicative seasonality-trend-noise decomposition models* (many $\theta \in \Theta$)
- ▶ We consider an extension with multiplicative autoregressive noise

$$Y_t = \text{Se} \left(t, \theta^{\text{Se}} \right) \text{Tr} \left(t, \theta^{\text{Tr}} \right) \exp \left(\kappa X_t - \frac{\kappa^2 \sigma_X^2}{2(1 - \phi^2)} \right) \exp \left(\varepsilon_t^Y - \frac{\sigma_Y^2}{2} \right)$$
$$X_t = \phi X_{t-1} + \varepsilon_t^X$$

for some scaling factor $\kappa > 0$ and autoregressive coefficient $\phi \in (-1, 1)$, with $X_0 \sim N(0, \frac{\sigma_X^2}{1 - \phi^2})$ and $\varepsilon_t^X \stackrel{\text{iid}}{\sim} N(0, \sigma_X^2)$, $\varepsilon_t^Y \stackrel{\text{iid}}{\sim} N(0, \sigma_Y^2)$ for some $\sigma_X, \sigma_Y > 0$

- ▶ Seasonality: linear combination of sines, cosines (weekly, monthly, yearly periods)
- ▶ Trend: product of linear trend and exponential trend

Simulate to get pre-training data for forecasting with TSFMs

- ▶ For forecasting with TSFMs, Dooley et al. (2023) suggest training on simulations from *multiplicative seasonality-trend-noise decomposition models* (many $\theta \in \Theta$)
- ▶ We consider an extension with multiplicative autoregressive noise

$$Y_t = \text{Se} \left(t, \theta^{\text{Se}} \right) \text{Tr} \left(t, \theta^{\text{Tr}} \right) \exp \left(\kappa X_t - \frac{\kappa^2 \sigma_X^2}{2(1 - \phi^2)} \right) \exp \left(\varepsilon_t^Y - \frac{\sigma_Y^2}{2} \right)$$
$$X_t = \phi X_{t-1} + \varepsilon_t^X$$

for some scaling factor $\kappa > 0$ and autoregressive coefficient $\phi \in (-1, 1)$, with $X_0 \sim N(0, \frac{\sigma_X^2}{1 - \phi^2})$ and $\varepsilon_t^X \stackrel{\text{iid}}{\sim} N(0, \sigma_X^2)$, $\varepsilon_t^Y \stackrel{\text{iid}}{\sim} N(0, \sigma_Y^2)$ for some $\sigma_X, \sigma_Y > 0$

- ▶ Seasonality: linear combination of sines, cosines (weekly, monthly, yearly periods)
- ▶ Trend: product of linear trend and exponential trend
- ▶ Our rolling-window estimator can be used (consistent and asymptotically normal)

Simulation-and-regression with parameter uncertainty

- We propose minimizing the Monte Carlo approximation to

$$R_{P_{\widehat{\text{mix}}}}(f) = \mathbb{E}_{P_{\widehat{\text{mix}}}} \left[\sum_{t \in \mathcal{T}} L(X_t - [f(Y)]_t) \right] = \mathbb{E}_{\theta \sim \widehat{\text{CD}}_n} \left[\mathbb{E}_{P_\theta} \left[\sum_{t \in \mathcal{T}} L(X_t - [f(Y)]_t) \right] \right]$$

for some loss L , where the expectation is w.r.t. the mixture distribution $P_{\widehat{\text{mix}}} = \int_{\Theta} P_\theta d\widehat{\text{CD}}_n(\theta)$ of sequences X, Y corresponding to some *estimated confidence distribution* $\widehat{\text{CD}}_n$ over $\Theta \subset \mathbb{R}^p$ for θ_0

Simulation-and-regression with parameter uncertainty

- ▶ We propose minimizing the Monte Carlo approximation to

$$R_{P_{\widehat{\text{mix}}}}(f) = \mathbb{E}_{P_{\widehat{\text{mix}}}} \left[\sum_{t \in \mathcal{T}} L(X_t - [f(Y)]_t) \right] = \mathbb{E}_{\theta \sim \widehat{\text{CD}}_n} \left[\mathbb{E}_{P_\theta} \left[\sum_{t \in \mathcal{T}} L(X_t - [f(Y)]_t) \right] \right]$$

for some loss L , where the expectation is w.r.t. the mixture distribution $P_{\widehat{\text{mix}}} = \int_{\Theta} P_\theta d\widehat{\text{CD}}_n(\theta)$ of sequences X, Y corresponding to some *estimated confidence distribution* $\widehat{\text{CD}}_n$ over $\Theta \subset \mathbb{R}^p$ for θ_0

- ▶ Examples of confidence distributions:
 1. $N(\hat{\theta}, \hat{\Sigma}/n)$ where $\hat{\Sigma}$ is the estimate of the asymptotic covariance
 2. Bootstrap distribution, i.e. sampling distribution of re-estimates

Simulation-and-regression with parameter uncertainty

- ▶ We propose minimizing the Monte Carlo approximation to

$$R_{P_{\widehat{\text{mix}}}}(f) = \mathbb{E}_{P_{\widehat{\text{mix}}}} \left[\sum_{t \in \mathcal{T}} L(X_t - [f(Y)]_t) \right] = \mathbb{E}_{\theta \sim \widehat{\text{CD}}_n} \left[\mathbb{E}_{P_\theta} \left[\sum_{t \in \mathcal{T}} L(X_t - [f(Y)]_t) \right] \right]$$

for some loss L , where the expectation is w.r.t. the mixture distribution $P_{\widehat{\text{mix}}} = \int_{\Theta} P_\theta d\widehat{\text{CD}}_n(\theta)$ of sequences X, Y corresponding to some *estimated confidence distribution* $\widehat{\text{CD}}_n$ over $\Theta \subset \mathbb{R}^p$ for θ_0

- ▶ Examples of confidence distributions:
 1. $N(\hat{\theta}, \hat{\Sigma}/n)$ where $\hat{\Sigma}$ is the estimate of the asymptotic covariance
 2. Bootstrap distribution, i.e. sampling distribution of re-estimates
- ▶ Theory of confidence distributions from D. Liu, R. Y. Liu, and Xie (2022)

Connections to TSFMs and classical simulation-and-regression

- ▶ When $\widehat{CD}_n$ is uniform over Θ , this is like TSFM pre-training (no information)

Connections to TSFMs and classical simulation-and-regression

- ▶ When $\widehat{\text{CD}}_n$ is uniform over Θ , this is like TSFM pre-training (no information)
- ▶ When $\widehat{\text{CD}}_n$ is Dirac delta at $\hat{\theta}$, this is like classical simulation-and-regression, which aims to minimize the Monte Carlo approximation to

$$R_{P_{\hat{\theta}}}(f) = \mathbb{E}_{P_{\hat{\theta}}} \left[\sum_{t \in \mathcal{T}} L(X_t - [f(Y)]_t) \right]$$

where expectation is w.r.t. distribution $P_{\hat{\theta}}$ of sequences X, Y corresponding to $\hat{\theta}$

Questions we aim to answer

- ▶ **Calibration:** Does accounting for parameter uncertainty improve the calibration of pre-trained/untrained transformers that output multiple conditional quantiles?

Questions we aim to answer

- ▶ **Calibration:** Does accounting for parameter uncertainty improve the calibration of pre-trained/untrained transformers that output multiple conditional quantiles?
- ▶ **Pre-training:** When does a pre-trained transformer fine-tuned with our method outperform a transformer trained from scratch with our method?

Questions we aim to answer

- ▶ **Calibration:** Does accounting for parameter uncertainty improve the calibration of pre-trained/untrained transformers that output multiple conditional quantiles?
- ▶ **Pre-training:** When does a pre-trained transformer fine-tuned with our method outperform a transformer trained from scratch with our method?
- ▶ **Improving fine-tuning and simulation-and-regression:** When does accounting for parameter uncertainty improve methods based on parameter point estimates?

Thank you!

Questions or comments?

You can reach me at:
`mwiecksosa@cmu.edu`

Does accounting for parameter uncertainty improve the calibration of pre-trained/untrained transformers that output conditional quantiles?

- For untrained transformers, see answer to question: “In the context of simulation-and-regression, does accounting for parameter uncertainty achieve better performance than using parameter estimate?”
- For pre-trained transformers, see answer to question: “When does a pre-trained transformer fine-tuned with our method outperform a transformer trained from scratch with our method?”

In the context of simulation-and-regression, does accounting for parameter uncertainty achieve better performance than using parameter estimate?

– When does this inequality hold:

$$\mathbb{E}_{P_{\theta_0}} \left[\sum_{t \in \mathcal{T}} L \left(X_t - [f_{\widehat{P_{\theta_0}}_{\text{mix}}}^*(Y)]_t \right) \right] < \mathbb{E}_{P_{\theta_0}} \left[\sum_{t \in \mathcal{T}} L \left(X_t - [f_{\widehat{P_{\hat{\theta}}}}^*(Y)]_t \right) \right]$$

where both expectations are with respect to the distribution P_{θ_0} of the sequences corresponding to the unknown true parameter θ_0 , $f_{\widehat{P_{\theta_0}}_{\text{mix}}}^*$ is the minimizer in the function class of $R_{\widehat{P_{\theta_0}}_{\text{mix}}}(f)$, and $f_{\widehat{P_{\hat{\theta}}}}^*$ is the minimizer in the function class of $R_{\widehat{P_{\hat{\theta}}}}(f)$

When does a pre-trained transformer fine-tuned with our method outperform a transformer trained from scratch with our method?

- Intuitively, this occurs when the following conditions hold
- Large parameter uncertainty, e.g., when the sample size of the observed time series is small relative to the noise
- The distributions used in the pre-training phase are close to the true distribution, e.g., when we have domain knowledge or we have previously observed time series with similar distributions

What do we mean by “almost every” function? Prevalent subset

Definition (J. Yorke and Ott (2005))

Let V be a completely metrizable topological vector space. A Borel set $E \subset V$ is said to be **prevalent** if there exists some Borel measure μ on V such that:

1. $0 < \mu(C) < \infty$ for some compact subset C of V .
2. The set $E + v$ has full μ -measure (that is, the complement of $E + v$ has measure zero) for all $v \in V$.

– In this literature, it is common to say that, if $E \subset V$ contains a prevalent Borel set, then “almost every” element of V lies in E . In this sense, “almost every” smooth map from $\mathbb{R}^p$ to $\mathbb{R}^{2p+1}$ has the following properties.

What is a smooth manifold? (University of Toronto Lecture Notes) I

- A C^∞ -manifold is a Hausdorff topological space M , with countable basis, together with a C^∞ -atlas.
- An n -dimensional manifold is a space that locally looks like $\mathbb{R}^n$. To give a precise meaning to this idea, our space first of all has to come equipped with some topology (so that the word “local” makes sense).

What is a smooth manifold? (University of Toronto Lecture Notes) II

- Recall that a *topological space* is a set M , together with a collection of subsets of M , called *open subsets*, satisfying the following three axioms: (i) the empty set $\emptyset$ and the space M itself are both open, (ii) the intersection of any finite collection of open subsets is open, (iii) the union of any collection of open subsets is open.
- The collection of open subsets of M is also called the topology of M . A map $f : M_1 \rightarrow M_2$ between topological spaces is called *continuous* if the pre-image of any open subset in M_2 is open in M_1 . A continuous map with a continuous inverse is called a *homeomorphism*.

What is a smooth manifold? (University of Toronto Lecture Notes) III

- One ingredient in the definition of a manifold is that our topological space comes equipped with a covering by open sets which are homeomorphic to open subsets of $\mathbb{R}^n$.
- Let M be a topological space. An n -dimensional *chart* for M is a pair (U, ϕ) consisting of an open subset $U \subset M$ and a continuous map $\phi : U \rightarrow \mathbb{R}^n$ such that ϕ is a homeomorphism onto its image $\phi(U)$. Two such charts (U_α, ϕ_α) , (U_β, ϕ_β) are C^∞ -compatible if the transition map

$$\phi_\beta \circ \phi_\alpha^{-1} : \phi_\alpha(U_\alpha \cap U_\beta) \rightarrow \phi_\beta(U_\alpha \cap U_\beta)$$

is a diffeomorphism (a smooth map with smooth inverse). A covering $\mathcal{A} = (U_\alpha)_{\alpha \in A}$ of M by pairwise C^∞ -compatible charts is called a C^∞ -atlas.

What is a smooth manifold? (University of Toronto Lecture Notes) IV

– We recall that a topological space is called *Hausdorff* if any two points have disjoint open neighborhoods. A *basis* for a topological space M is a collection $\mathcal{B}$ of open subsets of M such that every open subset of M is a union of open subsets in the collection $\mathcal{B}$. For example, the collection of open balls $B_\varepsilon(x)$ in $\mathbb{R}^n$ define a basis. But one already has a basis if one takes only all balls $B_\varepsilon(x)$ with $x \in \mathbb{Q}^n$ and $\varepsilon \in \mathbb{Q}_{>0}$; this defines a countable basis. A topological space with countable basis is also called *second countable*.

References I

-  Bodik, Juraj and Olivier C. Pasche (2024). “Granger causality in extremes”. *arXiv preprint arXiv: 2407.09632*.
-  Dooley, Samuel et al. (2023). “Forecastpfn: Synthetically-trained zero-shot forecasting”. In: *Advances in Neural Information Processing Systems 36*, pp. 2403–2426.
-  Liu, Dungang, Regina Y. Liu, and Min-ge Xie (2022). “Nonparametric fusion learning for multiparameters: Synthesize inferences from diverse sources using data depth and confidence distribution”. In: *Journal of the American Statistical Association* 117.540, pp. 2086–2104.
-  Mies, Fabian (2023). “Functional estimation and change detection for nonstationary time series”. In: *Journal of the American Statistical Association* 118.542, pp. 1011–1022.
-  Sauer, Tim, James A. Yorke, and Martin Casdagli (1991). “Embedology”. In: *Journal of Statistical Physics* 65.3, pp. 579–616.

References II

-  Shah, Rajen D. and Jonas Peters (2020). “The hardness of conditional independence testing and the generalised covariance measure”. In: *Annals of Statistics* 48.3, pp. 1514–1538.
-  Whitney, Hassler (1936). “Differentiable manifolds”. In: *Annals of Mathematics* 37.3, pp. 645–680.
-  Wu, Wei Biao (2005). “Nonlinear system theory: another look at dependence”. In: *Proceedings of the National Academy of Sciences* 102.40, pp. 14150–14154.
-  Yorke, James and William Ott (2005). “Prevalence”. In: *Bulletin of the American Mathematical Society* 42.3, pp. 263–290.