

Simulation-based Time Series Analysis:

Testing for Structure, Likelihood-free Estimation, and Prediction with Transformers using Synthetic Data

Michael Wieck-Sosa

Department of Statistics & Data Science, Carnegie Mellon University

April 3, 2026

Overview of presentation

Overview of presentation

Part I

Testing

Economics application

Hardness result

Physical dependence

Type I error control

Overview of presentation

Part I

Testing

Economics application

Hardness result

Physical dependence

Type I error control

Part II

Estimation

Random features

Embedding results

Consistency

Asymptotic normality

Overview of presentation

Part I

Testing

Economics application

Hardness result

Physical dependence

Type I error control

Part II

Estimation

Random features

Embedding results

Consistency

Asymptotic normality

Part III

Prediction

Foundation models

Sim-and-regression

Conf. distributions

Calibration

Time series setting

- ▶ We observe X_t , $t = 1, \dots, n$, for some sample size $n \in \mathbb{N}$

Time series setting

- ▶ We observe X_t , $t = 1, \dots, n$, for some sample size $n \in \mathbb{N}$
- ▶ Each observation X_t takes values in $\mathbb{R}^d$ for some dimension $d \in \mathbb{N}$

Time series setting

- ▶ We observe X_t , $t = 1, \dots, n$, for some sample size $n \in \mathbb{N}$
- ▶ Each observation X_t takes values in $\mathbb{R}^d$ for some dimension $d \in \mathbb{N}$
- ▶ **Temporal dependence:** X_t may depend on the past $X_{t-1}, X_{t-2}, \dots$

Time series setting

- ▶ We observe X_t , $t = 1, \dots, n$, for some sample size $n \in \mathbb{N}$
- ▶ Each observation X_t takes values in $\mathbb{R}^d$ for some dimension $d \in \mathbb{N}$
- ▶ **Temporal dependence:** X_t may depend on the past $X_{t-1}, X_{t-2}, \dots$
- ▶ **Nonstationarity:** the joint distribution of $X_{t_1+\tau}, \dots, X_{t_m+\tau}$ at $m \in \mathbb{N}$ distinct time points $t_1, \dots, t_m$ may change with the time-offset τ

Conditional independence testing

- ▶ Let $X_{1:n}$, $Y_{1:n}$, $Z_{1:n}$ be nonstationary time series taking values in $\mathbb{R}^{d_x}$, $\mathbb{R}^{d_y}$, $\mathbb{R}^{d_z}$

Conditional independence testing

- ▶ Let $X_{1:n}$, $Y_{1:n}$, $Z_{1:n}$ be nonstationary time series taking values in $\mathbb{R}^{d_x}$, $\mathbb{R}^{d_y}$, $\mathbb{R}^{d_z}$
- ▶ Let these variables incorporate any desired leads and lags (to simplify notation)

Conditional independence testing

- ▶ Let $X_{1:n}$, $Y_{1:n}$, $Z_{1:n}$ be nonstationary time series taking values in $\mathbb{R}^{d_x}$, $\mathbb{R}^{d_y}$, $\mathbb{R}^{d_z}$
- ▶ Let these variables incorporate any desired leads and lags (to simplify notation)
- ▶ The conditioning set of variables Z_t must be known at time t

Conditional independence testing

- ▶ Let $X_{1:n}$, $Y_{1:n}$, $Z_{1:n}$ be nonstationary time series taking values in $\mathbb{R}^{d_x}$, $\mathbb{R}^{d_y}$, $\mathbb{R}^{d_z}$
- ▶ Let these variables incorporate any desired leads and lags (to simplify notation)
- ▶ The conditioning set of variables Z_t must be known at time t
- ▶ We present a test for the null hypothesis that

$$X_t \perp\!\!\!\perp Y_t \mid Z_t \text{ for all } t \in [n]$$

Conditional independence testing

- ▶ Let $X_{1:n}$, $Y_{1:n}$, $Z_{1:n}$ be nonstationary time series taking values in $\mathbb{R}^{d_x}$, $\mathbb{R}^{d_y}$, $\mathbb{R}^{d_z}$
- ▶ Let these variables incorporate any desired leads and lags (to simplify notation)
- ▶ The conditioning set of variables Z_t must be known at time t
- ▶ We present a test for the null hypothesis that

$$X_t \perp\!\!\!\perp Y_t \mid Z_t \text{ for all } t \in [n]$$

- ▶ Motivations: understanding structure, variable selection, and causal discovery

Conditional independence testing


- ▶ Let $X_{1:n}$, $Y_{1:n}$, $Z_{1:n}$ be nonstationary time series taking values in $\mathbb{R}^{d_X}$, $\mathbb{R}^{d_Y}$, $\mathbb{R}^{d_Z}$
- ▶ Let these variables incorporate any desired leads and lags (to simplify notation)
- ▶ The conditioning set of variables Z_t must be known at time t
- ▶ We present a test for the null hypothesis that

$$X_t \perp\!\!\!\perp Y_t \mid Z_t \text{ for all } t \in [n]$$

- ▶ Motivations: understanding structure, variable selection, and causal discovery
- ▶ **Example:** links in the global economic system

Real data application: links in the global economic system

- ▶ Four stock market indices:
 - ① S&P 500 (US)
 - ② FTSE 100 (UK)
 - ③ Hang Seng (HK)
 - ④ Nikkei 225 (JP)
- ▶ Daily returns based on closing prices
- ▶ Trading days Jan. 2022–Apr. 2025
- ▶ Sample size: $n = 860$
- ▶ Results based on BH-adjusted p-values



Conditional independence testing reveals structure

- For each pair $X, Y \in \{\text{S\&P}, \text{FTSE}, \text{HangSeng}, \text{Nikkei}\}$, test

$$X_t \perp\!\!\!\perp Y_t \text{ for all times } t \in [n]$$

Conditional independence testing reveals structure

- ▶ For each pair $X, Y \in \{\text{S\&P}, \text{FTSE}, \text{HangSeng}, \text{Nikkei}\}$, test

$$X_t \perp\!\!\!\perp Y_t \text{ for all times } t \in [n]$$

- ▶ We reject every null hypothesis at the significance level $\alpha = 0.05$

Conditional independence testing reveals structure

- ▶ For each pair $X, Y \in \{\text{S\&P}, \text{FTSE}, \text{HangSeng}, \text{Nikkei}\}$, test

$$X_t \perp\!\!\!\perp Y_t \text{ for all times } t \in [n]$$

- ▶ We reject every null hypothesis at the significance level $\alpha = 0.05$
- ▶ For each triplet $X, Y, Z \in \{\text{S\&P}, \text{FTSE}, \text{HangSeng}, \text{Nikkei}\}$, test

$$X_t \perp\!\!\!\perp Y_t \mid Z_t \text{ for all times } t \in [n]$$

Conditional independence testing reveals structure

- ▶ For each pair $X, Y \in \{\text{S\&P}, \text{FTSE}, \text{HangSeng}, \text{Nikkei}\}$, test

$$X_t \perp\!\!\!\perp Y_t \text{ for all times } t \in [n]$$

- ▶ We reject every null hypothesis at the significance level $\alpha = 0.05$
- ▶ For each triplet $X, Y, Z \in \{\text{S\&P}, \text{FTSE}, \text{HangSeng}, \text{Nikkei}\}$, test

$$X_t \perp\!\!\!\perp Y_t \mid Z_t \text{ for all times } t \in [n]$$

- ▶ **Retain:** S\&P(t) is independent of HangSeng(t) and Nikkei(t) given FTSE(t)

Conditional independence testing reveals structure

- ▶ For each pair $X, Y \in \{\text{S\&P}, \text{FTSE}, \text{HangSeng}, \text{Nikkei}\}$, test

$$X_t \perp\!\!\!\perp Y_t \text{ for all times } t \in [n]$$

- ▶ We reject every null hypothesis at the significance level $\alpha = 0.05$
- ▶ For each triplet $X, Y, Z \in \{\text{S\&P}, \text{FTSE}, \text{HangSeng}, \text{Nikkei}\}$, test

$$X_t \perp\!\!\!\perp Y_t \mid Z_t \text{ for all times } t \in [n]$$

- ▶ **Retain:** S\&P(t) is independent of HangSeng(t) and Nikkei(t) given FTSE(t)
- ▶ **Reject:** S\&P(t) is independent of FTSE(t) given HangSeng(t) or Nikkei(t)

Conditional independence testing reveals structure

- ▶ For each pair $X, Y \in \{\text{S\&P}, \text{FTSE}, \text{HangSeng}, \text{Nikkei}\}$, test

$$X_t \perp\!\!\!\perp Y_t \text{ for all times } t \in [n]$$

- ▶ We reject every null hypothesis at the significance level $\alpha = 0.05$
- ▶ For each triplet $X, Y, Z \in \{\text{S\&P}, \text{FTSE}, \text{HangSeng}, \text{Nikkei}\}$, test

$$X_t \perp\!\!\!\perp Y_t \mid Z_t \text{ for all times } t \in [n]$$

- ▶ **Retain:** S&P(t) is independent of HangSeng(t) and Nikkei(t) given FTSE(t)
- ▶ **Reject:** S&P(t) is independent of FTSE(t) given HangSeng(t) or Nikkei(t)
- ▶ UK has info about Asia-Pacific markets on the same day, but not vice versa

Conditional independence testing reveals structure

- ▶ For each pair $X, Y \in \{\text{S\&P}, \text{FTSE}, \text{HangSeng}, \text{Nikkei}\}$, test

$$X_t \perp\!\!\!\perp Y_t \text{ for all times } t \in [n]$$

- ▶ We reject every null hypothesis at the significance level $\alpha = 0.05$
- ▶ For each triplet $X, Y, Z \in \{\text{S\&P}, \text{FTSE}, \text{HangSeng}, \text{Nikkei}\}$, test

$$X_t \perp\!\!\!\perp Y_t \mid Z_t \text{ for all times } t \in [n]$$

- ▶ **Retain:** S&P(t) is independent of HangSeng(t) and Nikkei(t) given FTSE(t)
- ▶ **Reject:** S&P(t) is independent of FTSE(t) given HangSeng(t) or Nikkei(t)
- ▶ UK has info about Asia-Pacific markets on the same day, but not vice versa
- ▶ Why? Earliest to latest closing exchanges: Tokyo, Hong Kong, London, New York

Conditional independence testing reveals structure

- ▶ For each pair $X, Y \in \{\text{S\&P}, \text{FTSE}, \text{HangSeng}, \text{Nikkei}\}$, test

$$X_t \perp\!\!\!\perp Y_t \text{ for all times } t \in [n]$$

- ▶ We reject every null hypothesis at the significance level $\alpha = 0.05$
- ▶ For each triplet $X, Y, Z \in \{\text{S\&P}, \text{FTSE}, \text{HangSeng}, \text{Nikkei}\}$, test

$$X_t \perp\!\!\!\perp Y_t \mid Z_t \text{ for all times } t \in [n]$$

- ▶ **Retain:** S&P(t) is independent of HangSeng(t) and Nikkei(t) given FTSE(t)
- ▶ **Reject:** S&P(t) is independent of FTSE(t) given HangSeng(t) or Nikkei(t)
- ▶ UK has info about Asia-Pacific markets on the same day, but not vice versa
- ▶ Why? Earliest to latest closing exchanges: Tokyo, Hong Kong, London, New York
- ▶ Next, we present the testing procedure

Algorithm for conditional independence test (1/2)

- **Inputs:** Data $X_{1:n}$, $Y_{1:n}$, $Z_{1:n}$, significance level α , number of simulations s , $p \in [2, \infty]$ for norm in test statistic, and method for selecting the window size L

Algorithm for conditional independence test (1/2)

- ▶ **Inputs:** Data $X_{1:n}$, $Y_{1:n}$, $Z_{1:n}$, significance level α , number of simulations s , $p \in [2, \infty]$ for norm in test statistic, and method for selecting the window size L
- ▶ Regress X_t on Z_t , $t \in [n]$, and obtain residuals $\hat{\varepsilon}_t$, $t \in [n]$

Algorithm for conditional independence test (1/2)

- ▶ **Inputs:** Data $X_{1:n}$, $Y_{1:n}$, $Z_{1:n}$, significance level α , number of simulations s , $p \in [2, \infty]$ for norm in test statistic, and method for selecting the window size L
- ▶ Regress X_t on Z_t , $t \in [n]$, and obtain residuals $\hat{\varepsilon}_t$, $t \in [n]$
- ▶ Regress Y_t on Z_t , $t \in [n]$, and obtain residuals $\hat{\xi}_t$, $t \in [n]$

Algorithm for conditional independence test (1/2)

- ▶ **Inputs:** Data $X_{1:n}$, $Y_{1:n}$, $Z_{1:n}$, significance level α , number of simulations s , $p \in [2, \infty]$ for norm in test statistic, and method for selecting the window size L
- ▶ Regress X_t on Z_t , $t \in [n]$, and obtain residuals $\hat{\varepsilon}_t$, $t \in [n]$
- ▶ Regress Y_t on Z_t , $t \in [n]$, and obtain residuals $\hat{\xi}_t$, $t \in [n]$
- ▶ Obtain residual products $\hat{R}_t = \text{vec}(\hat{\varepsilon}_t \hat{\xi}_t^\top)$, $t \in [n]$

Algorithm for conditional independence test (1/2)

- ▶ **Inputs:** Data $X_{1:n}$, $Y_{1:n}$, $Z_{1:n}$, significance level α , number of simulations s , $p \in [2, \infty]$ for norm in test statistic, and method for selecting the window size L
- ▶ Regress X_t on Z_t , $t \in [n]$, and obtain residuals $\hat{\varepsilon}_t$, $t \in [n]$
- ▶ Regress Y_t on Z_t , $t \in [n]$, and obtain residuals $\hat{\xi}_t$, $t \in [n]$
- ▶ Obtain residual products $\hat{R}_t = \text{vec}(\hat{\varepsilon}_t \hat{\xi}_t^\top)$, $t \in [n]$
- ▶ Calculate test statistic on time series of residual products

$$T(\hat{R}_{1:n}) = \max_{j \in [n]} \left\| \frac{1}{\sqrt{n}} \sum_{t \leq j} \hat{R}_t \right\|_p$$

Algorithm for conditional independence test (2/2)

- ▶ Select the window size L for covariance estimation using $\hat{R}_{1:n}$

Algorithm for conditional independence test (2/2)

- ▶ Select the window size L for covariance estimation using $\hat{R}_{1:n}$
- ▶ For $t \in [n]$, obtain estimates of the time-varying covariances as

$$\hat{\Sigma}_t = \frac{1}{L} \left(\sum_{j=t-L+1}^t \hat{R}_j \right) \left(\sum_{j=t-L+1}^t \hat{R}_j \right)^{\top}$$

Algorithm for conditional independence test (2/2)

- ▶ Select the window size L for covariance estimation using $\hat{R}_{1:n}$
- ▶ For $t \in [n]$, obtain estimates of the time-varying covariances as

$$\hat{\Sigma}_t = \frac{1}{L} \left(\sum_{j=t-L+1}^t \hat{R}_j \right) \left(\sum_{j=t-L+1}^t \hat{R}_j \right)^\top$$

- ▶ For each $r \in [s]$ and $t \in [n]$, simulate *independent* Gaussians $\tilde{R}_t^{(r)} \sim N(0, \hat{\Sigma}_t)$

Algorithm for conditional independence test (2/2)

- ▶ Select the window size L for covariance estimation using $\hat{R}_{1:n}$
- ▶ For $t \in [n]$, obtain estimates of the time-varying covariances as

$$\hat{\Sigma}_t = \frac{1}{L} \left(\sum_{j=t-L+1}^t \hat{R}_j \right) \left(\sum_{j=t-L+1}^t \hat{R}_j \right)^\top$$

- ▶ For each $r \in [s]$ and $t \in [n]$, simulate *independent* Gaussians $\tilde{R}_t^{(r)} \sim N(0, \hat{\Sigma}_t)$
- ▶ For each $r \in [s]$, calculate test statistic on simulated Gaussian process

$$T(\tilde{R}_{1:n}^{(r)}) = \max_{j \in [n]} \left\| \frac{1}{\sqrt{n}} \sum_{t \leq j} \tilde{R}_t^{(r)} \right\|_p$$

Algorithm for conditional independence test (2/2)

- ▶ Select the window size L for covariance estimation using $\hat{R}_{1:n}$
- ▶ For $t \in [n]$, obtain estimates of the time-varying covariances as

$$\hat{\Sigma}_t = \frac{1}{L} \left(\sum_{j=t-L+1}^t \hat{R}_j \right) \left(\sum_{j=t-L+1}^t \hat{R}_j \right)^\top$$

- ▶ For each $r \in [s]$ and $t \in [n]$, simulate *independent* Gaussians $\tilde{R}_t^{(r)} \sim N(0, \hat{\Sigma}_t)$
- ▶ For each $r \in [s]$, calculate test statistic on simulated Gaussian process

$$T(\tilde{R}_{1:n}^{(r)}) = \max_{j \in [n]} \left\| \frac{1}{\sqrt{n}} \sum_{t \leq j} \tilde{R}_t^{(r)} \right\|_p$$

- ▶ Calculate the $1 - \alpha$ empirical quantile $\hat{q}_{1-\alpha}^{\text{boot}}$ of $T(\tilde{R}_{1:n}^{(1)}), \dots, T(\tilde{R}_{1:n}^{(s)})$

Algorithm for conditional independence test (2/2)

- ▶ Select the window size L for covariance estimation using $\hat{R}_{1:n}$
- ▶ For $t \in [n]$, obtain estimates of the time-varying covariances as

$$\hat{\Sigma}_t = \frac{1}{L} \left(\sum_{j=t-L+1}^t \hat{R}_j \right) \left(\sum_{j=t-L+1}^t \hat{R}_j \right)^\top$$

- ▶ For each $r \in [s]$ and $t \in [n]$, simulate *independent* Gaussians $\tilde{R}_t^{(r)} \sim N(0, \hat{\Sigma}_t)$
- ▶ For each $r \in [s]$, calculate test statistic on simulated Gaussian process

$$T(\tilde{R}_{1:n}^{(r)}) = \max_{j \in [n]} \left\| \frac{1}{\sqrt{n}} \sum_{t \leq j} \tilde{R}_t^{(r)} \right\|_p$$

- ▶ Calculate the $1 - \alpha$ empirical quantile $\hat{q}_{1-\alpha}^{\text{boot}}$ of $T(\tilde{R}_{1:n}^{(1)}), \dots, T(\tilde{R}_{1:n}^{(s)})$
- ▶ Reject null hypothesis at level α if $T(\hat{R}_{1:n}) > \hat{q}_{1-\alpha}^{\text{boot}}$, else retain

The hardness of conditional independence (CI) testing

No-free-lunch in CI testing (Shah and Peters 2020)

The hardness of conditional independence (CI) testing

No-free-lunch in CI testing (Shah and Peters 2020)

- ▶ If a CI test has Type I error control (i.e. size α), then it cannot have power against any alternative (i.e. power against each alternative $Q \in \mathcal{Q}$ is at most α)

The hardness of conditional independence (CI) testing

No-free-lunch in CI testing (Shah and Peters 2020)

- ▶ If a CI test has Type I error control (i.e. size α), then it cannot have power against any alternative (i.e. power against each alternative $Q \in \mathcal{Q}$ is at most α)
- ▶ If a CI test has power β against an alternative $Q \in \mathcal{Q}$, then there exists a null distribution $P \in \mathcal{P}$ s.t. the test will reject with prob. greater than or equal to β

The hardness of conditional independence (CI) testing

No-free-lunch in CI testing (Shah and Peters 2020)

- ▶ If a CI test has Type I error control (i.e. size α), then it cannot have power against any alternative (i.e. power against each alternative $Q \in \mathcal{Q}$ is at most α)
- ▶ If a CI test has power β against an alternative $Q \in \mathcal{Q}$, then there exists a null distribution $P \in \mathcal{P}$ s.t. the test will reject with prob. greater than or equal to β
- ▶ This result has been extended to the time series setting (Bodik and Pasche 2024)

The hardness of conditional independence (CI) testing

No-free-lunch in CI testing (Shah and Peters 2020)

- ▶ If a CI test has Type I error control (i.e. size α), then it cannot have power against any alternative (i.e. power against each alternative $Q \in \mathcal{Q}$ is at most α)
- ▶ If a CI test has power β against an alternative $Q \in \mathcal{Q}$, then there exists a null distribution $P \in \mathcal{P}$ s.t. the test will reject with prob. greater than or equal to β
- ▶ This result has been extended to the time series setting (Bodik and Pasche 2024)
- ▶ **Implication:** We can only hope to achieve Type I error control on a subset of null distributions $\mathcal{P}_n^{\text{CI}}$, so must restrict the null hypothesis to make CI testing feasible

The hardness of conditional independence (CI) testing

No-free-lunch in CI testing (Shah and Peters 2020)

- ▶ If a CI test has Type I error control (i.e. size α), then it cannot have power against any alternative (i.e. power against each alternative $Q \in \mathcal{Q}$ is at most α)
- ▶ If a CI test has power β against an alternative $Q \in \mathcal{Q}$, then there exists a null distribution $P \in \mathcal{P}$ s.t. the test will reject with prob. greater than or equal to β
- ▶ This result has been extended to the time series setting (Bodik and Pasche 2024)
- ▶ **Implication:** We can only hope to achieve Type I error control on a subset of null distributions $\mathcal{P}_n^{\text{CI}}$, so must restrict the null hypothesis to make CI testing feasible
- ▶ **Key assumption:** uniform decay of physical dependence measure (Wu 2005)

The hardness of conditional independence (CI) testing

No-free-lunch in CI testing (Shah and Peters 2020)

- ▶ If a CI test has Type I error control (i.e. size α), then it cannot have power against any alternative (i.e. power against each alternative $Q \in \mathcal{Q}$ is at most α)
- ▶ If a CI test has power β against an alternative $Q \in \mathcal{Q}$, then there exists a null distribution $P \in \mathcal{P}$ s.t. the test will reject with prob. greater than or equal to β
- ▶ This result has been extended to the time series setting (Bodik and Pasche 2024)
- ▶ **Implication:** We can only hope to achieve Type I error control on a subset of null distributions $\mathcal{P}_n^{\text{CI}}$, so must restrict the null hypothesis to make CI testing feasible
- ▶ **Key assumption:** uniform decay of physical dependence measure (Wu 2005)
- ▶ Assume this for all time series $X_{1:n}, Y_{1:n}, Z_{1:n}$

Key assumption 1: Representation in terms of noise inputs (Wu 2005)

- ▶ Let $(\varepsilon_i)_{i \in \mathbb{Z}}$ be an iid sequence of random variables

Key assumption 1: Representation in terms of noise inputs (Wu 2005)

- ▶ Let $(\varepsilon_i)_{i \in \mathbb{Z}}$ be an iid sequence of random variables
- ▶ Define the sequence of noise inputs up to time t by $\varepsilon_t = (\varepsilon_t, \varepsilon_{t-1}, \dots)$

Key assumption 1: Representation in terms of noise inputs (Wu 2005)

- ▶ Let $(\varepsilon_i)_{i \in \mathbb{Z}}$ be an iid sequence of random variables
- ▶ Define the sequence of noise inputs up to time t by $\varepsilon_t = (\varepsilon_t, \varepsilon_{t-1}, \dots)$
- ▶ Let Θ be a parameter space, which can be infinite-dimensional

Key assumption 1: Representation in terms of noise inputs (Wu 2005)

- ▶ Let $(\varepsilon_i)_{i \in \mathbb{Z}}$ be an iid sequence of random variables
- ▶ Define the sequence of noise inputs up to time t by $\varepsilon_t = (\varepsilon_t, \varepsilon_{t-1}, \dots)$
- ▶ Let Θ be a parameter space, which can be infinite-dimensional
- ▶ For each $n \in \mathbb{N}$, let the mappings $G_t^{(n)} : \mathbb{R}^\infty \times \Theta \rightarrow \mathbb{R}^d$, $t = 1, \dots, n$, define a *generative model* for time series

Key assumption 1: Representation in terms of noise inputs (Wu 2005)

- ▶ Let $(\varepsilon_i)_{i \in \mathbb{Z}}$ be an iid sequence of random variables
- ▶ Define the sequence of noise inputs up to time t by $\varepsilon_t = (\varepsilon_t, \varepsilon_{t-1}, \dots)$
- ▶ Let Θ be a parameter space, which can be infinite-dimensional
- ▶ For each $n \in \mathbb{N}$, let the mappings $G_t^{(n)} : \mathbb{R}^\infty \times \Theta \rightarrow \mathbb{R}^d$, $t = 1, \dots, n$, define a *generative model* for time series

Assumption

Given a parameter $\theta \in \Theta$, index $n \in \mathbb{N}$ (linked to sample size and approximation), and noise inputs ε_t , a time series $X_{1:n}$ is generated by nature as

$$X_t = G_t^{(n)}(\varepsilon_t, \theta), \quad t = 1, \dots, n$$

Key assumption 1: Representation in terms of noise inputs (Wu 2005)

- ▶ Let $(\varepsilon_i)_{i \in \mathbb{Z}}$ be an iid sequence of random variables
- ▶ Define the sequence of noise inputs up to time t by $\varepsilon_t = (\varepsilon_t, \varepsilon_{t-1}, \dots)$
- ▶ Let Θ be a parameter space, which can be infinite-dimensional
- ▶ For each $n \in \mathbb{N}$, let the mappings $G_t^{(n)} : \mathbb{R}^\infty \times \Theta \rightarrow \mathbb{R}^d$, $t = 1, \dots, n$, define a *generative model* for time series

Assumption

Given a parameter $\theta \in \Theta$, index $n \in \mathbb{N}$ (linked to sample size and approximation), and noise inputs ε_t , a time series $X_{1:n}$ is generated by nature as

$$X_t = G_t^{(n)}(\varepsilon_t, \theta), \quad t = 1, \dots, n$$

- ▶ **We do not know** the generative model $G_t^{(n)}$, noise inputs ε_t , or parameter θ

Key assumption 2: Decaying influence of noise inputs (Wu 2005)

- Let $(\varepsilon_i^*)_{i \in \mathbb{Z}}$ be an iid copy of $(\varepsilon_i)_{i \in \mathbb{Z}}$. For $\ell \in \mathbb{Z}$, $j \in \mathbb{N}_0$, denote

$$\tilde{\varepsilon}_{\ell,j} = (\varepsilon_\ell, \dots, \varepsilon_{\ell-j+1}, \varepsilon_{\ell-j}^*, \varepsilon_{\ell-j-1}, \dots)$$

Key assumption 2: Decaying influence of noise inputs (Wu 2005)

- ▶ Let $(\varepsilon_i^*)_{i \in \mathbb{Z}}$ be an iid copy of $(\varepsilon_i)_{i \in \mathbb{Z}}$. For $\ell \in \mathbb{Z}$, $j \in \mathbb{N}_0$, denote

$$\tilde{\varepsilon}_{\ell,j} = (\varepsilon_\ell, \dots, \varepsilon_{\ell-j+1}, \varepsilon_{\ell-j}^*, \varepsilon_{\ell-j-1}, \dots)$$

- ▶ We assume the influence of any one past noise input $\varepsilon_{\ell-j}^*$ *decays* as j grows

Key assumption 2: Decaying influence of noise inputs (Wu 2005)

- Let $(\varepsilon_i^*)_{i \in \mathbb{Z}}$ be an iid copy of $(\varepsilon_i)_{i \in \mathbb{Z}}$. For $\ell \in \mathbb{Z}$, $j \in \mathbb{N}_0$, denote

$$\tilde{\varepsilon}_{\ell,j} = (\varepsilon_\ell, \dots, \varepsilon_{\ell-j+1}, \varepsilon_{\ell-j}^*, \varepsilon_{\ell-j-1}, \dots)$$

- We assume the influence of any one past noise input $\varepsilon_{\ell-j}^*$ decays as j grows

Assumption

There exist $\Psi, \rho > 0$ such that, for some $q > 2$ and all $j \in \mathbb{N}$, we have

$$\sup_{i,t,\ell,\theta} \left\| G_t^{(i)}(\varepsilon_\ell, \theta) - G_t^{(i)}(\tilde{\varepsilon}_{\ell,j}, \theta) \right\|_{L^q(\theta)} \leq \Psi(j \vee 1)^{-\rho}$$

where $\|\cdot\|_{L^q(\theta)} = \mathbb{E}_\theta(\|\cdot\|_2^q)^{1/q}$, $\mathbb{E}_\theta(\cdot)$ is expectation w.r.t. θ -dependent distribution

We prove our test has uniformly asymptotic Type I error control

- For all $\theta \in \Theta$ and $n \in \mathbb{N}$, denote the distribution of $(G_t^{(n)}(\varepsilon_\ell, \theta))_{t \in [n], \ell \in \mathbb{Z}}$ by $P_{\theta, n}$
 - Note: $(G_t^{(n)}(\varepsilon_t, \theta))_{t \in [n]}$ is nonstationary, and $(G_t^{(n)}(\varepsilon_\ell, \theta))_{\ell \in \mathbb{Z}}$ is stationary $\forall t$

We prove our test has uniformly asymptotic Type I error control

- ▶ For all $\theta \in \Theta$ and $n \in \mathbb{N}$, denote the distribution of $(G_t^{(n)}(\varepsilon_\ell, \theta))_{t \in [n], \ell \in \mathbb{Z}}$ by $P_{\theta, n}$
 - Note: $(G_t^{(n)}(\varepsilon_t, \theta))_{t \in [n]}$ is nonstationary, and $(G_t^{(n)}(\varepsilon_\ell, \theta))_{\ell \in \mathbb{Z}}$ is stationary $\forall t$
- ▶ Let $\mathcal{P}_n^* \subset \mathcal{P}_n^{\text{CI}}$ be the collection of null distributions $\mathcal{P}_n^{\text{CI}}$ s.t. *all assumptions hold*

We prove our test has uniformly asymptotic Type I error control

- ▶ For all $\theta \in \Theta$ and $n \in \mathbb{N}$, denote the distribution of $(G_t^{(n)}(\varepsilon_\ell, \theta))_{t \in [n], \ell \in \mathbb{Z}}$ by $P_{\theta, n}$
 - Note: $(G_t^{(n)}(\varepsilon_t, \theta))_{t \in [n]}$ is nonstationary, and $(G_t^{(n)}(\varepsilon_\ell, \theta))_{\ell \in \mathbb{Z}}$ is stationary $\forall t$
- ▶ Let $\mathcal{P}_n^* \subset \mathcal{P}_n^{\text{CI}}$ be the collection of null distributions $\mathcal{P}_n^{\text{CI}}$ s.t. *all assumptions hold*
- ▶ Need slow convergence rates for regression estimators (nonstationary time series)

We prove our test has uniformly asymptotic Type I error control

- ▶ For all $\theta \in \Theta$ and $n \in \mathbb{N}$, denote the distribution of $(G_t^{(n)}(\varepsilon_\ell, \theta))_{t \in [n], \ell \in \mathbb{Z}}$ by $P_{\theta, n}$
 - Note: $(G_t^{(n)}(\varepsilon_t, \theta))_{t \in [n]}$ is nonstationary, and $(G_t^{(n)}(\varepsilon_\ell, \theta))_{\ell \in \mathbb{Z}}$ is stationary $\forall t$
- ▶ Let $\mathcal{P}_n^* \subset \mathcal{P}_n^{\text{CI}}$ be the collection of null distributions $\mathcal{P}_n^{\text{CI}}$ s.t. *all assumptions hold*
- ▶ Need slow convergence rates for regression estimators (nonstationary time series)
- ▶ For all distributions $P \in \mathcal{P}_n^*$, the **product** of $L^2(P)$ norms of the estimation errors must be $o\left(\frac{1}{\sqrt{n} \text{polylog}(n)}\right)$, but each $L^2(P)$ norm only needs to be $o\left(\frac{1}{\text{polylog}(n)}\right)$

We prove our test has uniformly asymptotic Type I error control

- ▶ For all $\theta \in \Theta$ and $n \in \mathbb{N}$, denote the distribution of $(G_t^{(n)}(\varepsilon_\ell, \theta))_{t \in [n], \ell \in \mathbb{Z}}$ by $P_{\theta, n}$
 - Note: $(G_t^{(n)}(\varepsilon_t, \theta))_{t \in [n]}$ is nonstationary, and $(G_t^{(n)}(\varepsilon_\ell, \theta))_{\ell \in \mathbb{Z}}$ is stationary $\forall t$
- ▶ Let $\mathcal{P}_n^* \subset \mathcal{P}_n^{\text{CI}}$ be the collection of null distributions $\mathcal{P}_n^{\text{CI}}$ s.t. *all assumptions hold*
- ▶ Need slow convergence rates for regression estimators (nonstationary time series)
- ▶ For all distributions $P \in \mathcal{P}_n^*$, the **product** of $L^2(P)$ norms of the estimation errors must be $o\left(\frac{1}{\sqrt{n} \text{polylog}(n)}\right)$, but each $L^2(P)$ norm only needs to be $o\left(\frac{1}{\text{polylog}(n)}\right)$

Theorem (Informal)

Recall the test statistic $T(\hat{R}_{1:n})$ and estimated quantile $\hat{q}_{1-\alpha}^{\text{boot}}$. We have

$$\limsup_{n \rightarrow \infty} \sup_{P \in \mathcal{P}_n^*} \mathbb{P}_P \left(T(\hat{R}_{1:n}) > \hat{q}_{1-\alpha}^{\text{boot}} \right) \leq \alpha.$$

Simulation-based estimation and inference

- ▶ Suppose we observe a time series $X_{1:n}^{\text{obs}}$ taking values in $\mathbb{R}^d$

Simulation-based estimation and inference

- ▶ Suppose we observe a time series $X_{1:n}^{\text{obs}}$ taking values in $\mathbb{R}^d$
- ▶ We assume the observed time series has been *generated by nature* as

$$X_t^{\text{obs}} = G_t^{(n)}(\varepsilon_t, \theta_0), \quad t = 1, \dots, n$$

at some unknown parameter $\theta_0 \in \Theta \subset \mathbb{R}^p$ with some unknown noise inputs ε_t

Simulation-based estimation and inference

- ▶ Suppose we observe a time series $X_{1:n}^{\text{obs}}$ taking values in $\mathbb{R}^d$
- ▶ We assume the observed time series has been *generated by nature* as

$$X_t^{\text{obs}} = G_t^{(n)}(\varepsilon_t, \theta_0), \quad t = 1, \dots, n$$

at some unknown parameter $\theta_0 \in \Theta \subset \mathbb{R}^p$ with some unknown noise inputs ε_t

- ▶ We want to estimate θ_0 , get a confidence set, or test goodness-of-fit

Simulation-based estimation and inference

- ▶ Suppose we observe a time series $X_{1:n}^{\text{obs}}$ taking values in $\mathbb{R}^d$
- ▶ We assume the observed time series has been *generated by nature* as

$$X_t^{\text{obs}} = G_t^{(n)}(\varepsilon_t, \theta_0), \quad t = 1, \dots, n$$

at some unknown parameter $\theta_0 \in \Theta \subset \mathbb{R}^p$ with some unknown noise inputs ε_t

- ▶ We want to estimate θ_0 , get a confidence set, or test goodness-of-fit
- ▶ The likelihood is intractable, so we use a simulation-based approach

Simulation-based estimation and inference

- ▶ Suppose we observe a time series $X_{1:n}^{\text{obs}}$ taking values in $\mathbb{R}^d$
- ▶ We assume the observed time series has been *generated by nature* as

$$X_t^{\text{obs}} = G_t^{(n)}(\varepsilon_t, \theta_0), \quad t = 1, \dots, n$$

at some unknown parameter $\theta_0 \in \Theta \subset \mathbb{R}^p$ with some unknown noise inputs ε_t

- ▶ We want to estimate θ_0 , get a confidence set, or test goodness-of-fit
- ▶ The likelihood is intractable, so we use a simulation-based approach
- ▶ **Key assumption:** We know $G_t^{(n)}$ and can simulate from it at any value of $\theta \in \Theta$

How simulation-based estimation is usually done

- ▶ **Rough idea:** Tune the parameters θ until some features F “look right”

How simulation-based estimation is usually done

- ▶ **Rough idea:** Tune the parameters θ until some features F “look right”
- ▶ For (asymptotically) stationary time series, often use minimum distance estimator

$$\hat{\theta} = \operatorname{argmin}_{\theta \in \Theta} \left\| F^{\text{obs}} - \frac{1}{s} \sum_{r=1}^s F^{(r)}(\theta) \right\|$$

How simulation-based estimation is usually done

- ▶ **Rough idea:** Tune the parameters θ until some features F “look right”
- ▶ For (asymptotically) stationary time series, often use minimum distance estimator

$$\hat{\theta} = \operatorname{argmin}_{\theta \in \Theta} \left\| F^{\text{obs}} - \frac{1}{s} \sum_{r=1}^s F^{(r)}(\theta) \right\|$$

- ▶ For consistency $\hat{\theta} \xrightarrow{P} \theta_0$, need $F(\theta) \xrightarrow{P} \mathbb{E}[F(\theta)] \equiv \Phi(\theta)$ for all θ (ULLN)

How simulation-based estimation is usually done

- ▶ **Rough idea:** Tune the parameters θ until some features F “look right”
- ▶ For (asymptotically) stationary time series, often use minimum distance estimator

$$\hat{\theta} = \operatorname{argmin}_{\theta \in \Theta} \left\| F^{\text{obs}} - \frac{1}{s} \sum_{r=1}^s F^{(r)}(\theta) \right\|$$

- ▶ For consistency $\hat{\theta} \xrightarrow{P} \theta_0$, need $F(\theta) \xrightarrow{P} \mathbb{E}[F(\theta)] \equiv \Phi(\theta)$ for all θ (ULLN)
- ▶ Hope Φ^{-1} exists and is nice, so we can translate from $\Phi(\theta)$ back to θ

How simulation-based estimation is usually done

- ▶ **Rough idea:** Tune the parameters θ until some features F “look right”
- ▶ For (asymptotically) stationary time series, often use minimum distance estimator

$$\hat{\theta} = \operatorname{argmin}_{\theta \in \Theta} \left\| F^{\text{obs}} - \frac{1}{s} \sum_{r=1}^s F^{(r)}(\theta) \right\|$$

- ▶ For consistency $\hat{\theta} \xrightarrow{P} \theta_0$, need $F(\theta) \xrightarrow{P} \mathbb{E}[F(\theta)] \equiv \Phi(\theta)$ for all θ (ULLN)
- ▶ Hope Φ^{-1} exists and is nice, so we can translate from $\Phi(\theta)$ back to θ
- ▶ For nonstationary time series, need rolling-window estimator (not discussed here)

How simulation-based estimation is usually done

- ▶ **Rough idea:** Tune the parameters θ until some features F “look right”
- ▶ For (asymptotically) stationary time series, often use minimum distance estimator

$$\hat{\theta} = \operatorname{argmin}_{\theta \in \Theta} \left\| F^{\text{obs}} - \frac{1}{s} \sum_{r=1}^s F^{(r)}(\theta) \right\|$$

- ▶ For consistency $\hat{\theta} \xrightarrow{P} \theta_0$, need $F(\theta) \xrightarrow{P} \mathbb{E}[F(\theta)] \equiv \Phi(\theta)$ for all θ (ULLN)
- ▶ Hope Φ^{-1} exists and is nice, so we can translate from $\Phi(\theta)$ back to θ
- ▶ For nonstationary time series, need rolling-window estimator (not discussed here)
- ▶ *Many* choices for F , e.g., moment-based statistics, parameter from auxiliary model

Why not use random features?

- ▶ Good features are sensitive to small changes in θ , but takes time to craft features

Why not use random features?

- ▶ Good features are sensitive to small changes in θ , but takes time to craft features
- ▶ Can we just use *random features*? No need to craft features, computationally efficient, little knowledge about model is needed

Why not use random features?

- ▶ Good features are sensitive to small changes in θ , but takes time to craft features
- ▶ Can we just use *random features*? No need to craft features, computationally efficient, little knowledge about model is needed
- ▶ Typically, need $2p + 1$ random features for a p -dimensional parameter space Θ

Why not use random features?

- ▶ Good features are sensitive to small changes in θ , but takes time to craft features
- ▶ Can we just use *random features*? No need to craft features, computationally efficient, little knowledge about model is needed
- ▶ Typically, need $2p + 1$ random features for a p -dimensional parameter space Θ
- ▶ Probabilistic embedding theorem due to Sauer, J. A. Yorke, and Casdagli (1991)

Why not use random features?

- ▶ Good features are sensitive to small changes in θ , but takes time to craft features
- ▶ Can we just use *random features*? No need to craft features, computationally efficient, little knowledge about model is needed
- ▶ Typically, need $2p + 1$ random features for a p -dimensional parameter space Θ
- ▶ Probabilistic embedding theorem due to Sauer, J. A. Yorke, and Casdagli (1991)
- ▶ For “almost every” C^1 smooth function Φ from $\mathbb{R}^p$ to $\mathbb{R}^{2p+1}$:
 1. Φ **is one-to-one** on compact subsets of $\mathbb{R}^p$
 2. Φ **has a nice derivative** on certain compact subsets of $\mathbb{R}^p$
 - Details: the $(2p + 1) \times p$ Jacobian matrix of Φ has full rank p on each compact subset of a smooth manifold contained in $\mathbb{R}^p$

How many random features?

- ▶ Let P_θ be the joint distribution of $[X_t(\theta), \dots, X_{t-m}(\theta)]^\top$ for some $m \in \mathbb{N}$

How many random features?

- ▶ Let P_θ be the joint distribution of $[X_t(\theta), \dots, X_{t-m}(\theta)]^\top$ for some $m \in \mathbb{N}$
- ▶ Under smoothness conditions on $\theta \mapsto P_\theta$, the expectations of $2p + 1$ “generic” continuous, bounded functions $\varphi : \mathbb{R}^{(m+1) \times d} \rightarrow \mathbb{R}^{2p+1}$ will be C^1 smooth in θ

How many random features?

- ▶ Let P_θ be the joint distribution of $[X_t(\theta), \dots, X_{t-m}(\theta)]^\top$ for some $m \in \mathbb{N}$
- ▶ Under smoothness conditions on $\theta \mapsto P_\theta$, the expectations of $2p+1$ “generic” continuous, bounded functions $\varphi : \mathbb{R}^{(m+1) \times d} \rightarrow \mathbb{R}^{2p+1}$ will be C^1 smooth in θ
- ▶ For some φ , define the C^1 smooth function $\Phi : \mathbb{R}^p \rightarrow \mathbb{R}^{2p+1}$ by extending the map

$$\theta \mapsto \mathbb{E}_{P_\theta} \left[\varphi \left([X_t(\theta), \dots, X_{t-m}(\theta)]^\top \right) \right]$$

from the domain $\Theta \subset \mathbb{R}^p$ to the domain $\mathbb{R}^p$

How many random features?

- ▶ Let P_θ be the joint distribution of $[X_t(\theta), \dots, X_{t-m}(\theta)]^\top$ for some $m \in \mathbb{N}$
- ▶ Under smoothness conditions on $\theta \mapsto P_\theta$, the expectations of $2p+1$ “generic” continuous, bounded functions $\varphi : \mathbb{R}^{(m+1) \times d} \rightarrow \mathbb{R}^{2p+1}$ will be C^1 smooth in θ
- ▶ For some φ , define the C^1 smooth function $\Phi : \mathbb{R}^p \rightarrow \mathbb{R}^{2p+1}$ by extending the map

$$\theta \mapsto \mathbb{E}_{P_\theta} \left[\varphi \left([X_t(\theta), \dots, X_{t-m}(\theta)]^\top \right) \right]$$

from the domain $\Theta \subset \mathbb{R}^p$ to the domain $\mathbb{R}^p$

- ▶ Unless we are very unlucky in our selection of φ , then Φ will be one-to-one on Θ and have a nice derivative

How many random features?

- ▶ Let P_θ be the joint distribution of $[X_t(\theta), \dots, X_{t-m}(\theta)]^\top$ for some $m \in \mathbb{N}$
- ▶ Under smoothness conditions on $\theta \mapsto P_\theta$, the expectations of $2p + 1$ “generic” continuous, bounded functions $\varphi : \mathbb{R}^{(m+1) \times d} \rightarrow \mathbb{R}^{2p+1}$ will be C^1 smooth in θ
- ▶ For some φ , define the C^1 smooth function $\Phi : \mathbb{R}^p \rightarrow \mathbb{R}^{2p+1}$ by extending the map

$$\theta \mapsto \mathbb{E}_{P_\theta} \left[\varphi \left([X_t(\theta), \dots, X_{t-m}(\theta)]^\top \right) \right]$$

from the domain $\Theta \subset \mathbb{R}^p$ to the domain $\mathbb{R}^p$

- ▶ Unless we are very unlucky in our selection of φ , then Φ will be one-to-one on Θ and have a nice derivative
- ▶ If physical dependence measure (Wu 2005) decays uniformly over Θ , we can estimate $\Phi(\theta)$, $\theta \in \Theta$, via time-averages of $2p + 1$ random features

$$\frac{1}{n-m} \sum_{t=m+1}^n \varphi \left([X_t(\theta), \dots, X_{t-m}(\theta)]^\top \right)$$

Example: Mean of Gaussian

- **Goal:** Estimate $\theta_0 = \mu_0 = 0.5$ given $X_1^{\text{obs}}, \dots, X_n^{\text{obs}} \stackrel{\text{iid}}{\sim} \mathcal{N}(\theta_0, 1)$ with $\Theta = [-3, 3]$

Example: Mean of Gaussian

- ▶ **Goal:** Estimate $\theta_0 = \mu_0 = 0.5$ given $X_1^{\text{obs}}, \dots, X_n^{\text{obs}} \stackrel{\text{iid}}{\sim} \mathcal{N}(\theta_0, 1)$ with $\Theta = [-3, 3]$
- ▶ Use time-averages of $2p + 1 = 3$ random Fourier features with $m = 0$ lags

Example: Mean of Gaussian

- ▶ **Goal:** Estimate $\theta_0 = \mu_0 = 0.5$ given $X_1^{\text{obs}}, \dots, X_n^{\text{obs}} \stackrel{\text{iid}}{\sim} \mathcal{N}(\theta_0, 1)$ with $\Theta = [-3, 3]$
- ▶ Use time-averages of $2p + 1 = 3$ random Fourier features with $m = 0$ lags
- ▶ For $i = 1, \dots, 2p + 1$, draw frequencies $w_i \stackrel{\text{iid}}{\sim} \mathcal{N}(0, 1)$, phases $b_i \stackrel{\text{iid}}{\sim} \text{Unif}(-\pi, \pi)$

Example: Mean of Gaussian

- ▶ **Goal:** Estimate $\theta_0 = \mu_0 = 0.5$ given $X_1^{\text{obs}}, \dots, X_n^{\text{obs}} \stackrel{\text{iid}}{\sim} \mathcal{N}(\theta_0, 1)$ with $\Theta = [-3, 3]$
- ▶ Use time-averages of $2p + 1 = 3$ random Fourier features with $m = 0$ lags
- ▶ For $i = 1, \dots, 2p + 1$, draw frequencies $w_i \stackrel{\text{iid}}{\sim} \mathcal{N}(0, 1)$, phases $b_i \stackrel{\text{iid}}{\sim} \text{Unif}(-\pi, \pi)$
- ▶ Define the i -th feature as $\varphi_i(x) = \cos(w_i x + b_i)$

Example: Mean of Gaussian

- ▶ **Goal:** Estimate $\theta_0 = \mu_0 = 0.5$ given $X_1^{\text{obs}}, \dots, X_n^{\text{obs}} \stackrel{\text{iid}}{\sim} \mathcal{N}(\theta_0, 1)$ with $\Theta = [-3, 3]$
- ▶ Use time-averages of $2p + 1 = 3$ random Fourier features with $m = 0$ lags
- ▶ For $i = 1, \dots, 2p + 1$, draw frequencies $w_i \stackrel{\text{iid}}{\sim} \mathcal{N}(0, 1)$, phases $b_i \stackrel{\text{iid}}{\sim} \text{Unif}(-\pi, \pi)$
- ▶ Define the i -th feature as $\varphi_i(x) = \cos(w_i x + b_i)$
- ▶ Denoting $\varphi = (\varphi_1, \dots, \varphi_{2p+1})$, our estimator is given by

$$\hat{\theta} = \underset{\theta \in \Theta}{\operatorname{argmin}} \left\| \frac{1}{n} \sum_{t=1}^n \varphi(X_t^{\text{obs}}) - \frac{1}{s} \sum_{r=1}^s \left(\frac{1}{n} \sum_{t=1}^n \varphi(X_t^{(r)}(\theta)) \right) \right\|$$


Example: Mean of Gaussian

- ▶ **Goal:** Estimate $\theta_0 = \mu_0 = 0.5$ given $X_1^{\text{obs}}, \dots, X_n^{\text{obs}} \stackrel{\text{iid}}{\sim} \mathcal{N}(\theta_0, 1)$ with $\Theta = [-3, 3]$
- ▶ Use time-averages of $2p + 1 = 3$ random Fourier features with $m = 0$ lags
- ▶ For $i = 1, \dots, 2p + 1$, draw frequencies $w_i \stackrel{\text{iid}}{\sim} \mathcal{N}(0, 1)$, phases $b_i \stackrel{\text{iid}}{\sim} \text{Unif}(-\pi, \pi)$
- ▶ Define the i -th feature as $\varphi_i(x) = \cos(w_i x + b_i)$
- ▶ Denoting $\varphi = (\varphi_1, \dots, \varphi_{2p+1})$, our estimator is given by


$$\hat{\theta} = \operatorname{argmin}_{\theta \in \Theta} \left\| \frac{1}{n} \sum_{t=1}^n \varphi(X_t^{\text{obs}}) - \frac{1}{s} \sum_{r=1}^s \left(\frac{1}{n} \sum_{t=1}^n \varphi(X_t^{(r)}(\theta)) \right) \right\|$$

- ▶ Let's visualize an embedding and check the estimation results


Each point is a time-average of 3 features, colored by value of parameter μ



Density plot from 1,000 independent estimates (different random features)



Theoretical MSE of $\hat{\theta}^{\text{MLE}}$ ($1/n$) vs MSE of our $\hat{\theta}$ (100 replicates per n)



Theoretical guarantees for time-average and rolling-window estimators

- ▶ Let $Q_n(\cdot)$ be sample discrepancy function s.t. $\hat{\theta} = \arg \min_{\theta \in \Theta} \|Q_n(\theta)\|$

Theoretical guarantees for time-average and rolling-window estimators

- ▶ Let $Q_n(\cdot)$ be sample discrepancy function s.t. $\hat{\theta} = \arg \min_{\theta \in \Theta} \|Q_n(\theta)\|$
- ▶ Let V be the asymptotic covariance matrix from $\sqrt{n}Q_n(\theta_0) \Rightarrow N(0, V)$

Theoretical guarantees for time-average and rolling-window estimators

- ▶ Let $Q_n(\cdot)$ be sample discrepancy function s.t. $\hat{\theta} = \arg \min_{\theta \in \Theta} \|Q_n(\theta)\|$
- ▶ Let V be the asymptotic covariance matrix from $\sqrt{n}Q_n(\theta_0) \Rightarrow N(0, V)$
- ▶ Let $Q(\cdot)$ be population discrepancy function s.t. $Q_n(\theta) \xrightarrow{P} Q(\theta)$ uniformly over Θ

Theoretical guarantees for time-average and rolling-window estimators

- ▶ Let $Q_n(\cdot)$ be sample discrepancy function s.t. $\hat{\theta} = \arg \min_{\theta \in \Theta} \|Q_n(\theta)\|$
- ▶ Let V be the asymptotic covariance matrix from $\sqrt{n}Q_n(\theta_0) \Rightarrow N(0, V)$
- ▶ Let $Q(\cdot)$ be population discrepancy function s.t. $Q_n(\theta) \xrightarrow{P} Q(\theta)$ uniformly over Θ
- ▶ Let Γ be the Jacobian matrix of $Q(\cdot)$ at θ_0

Theoretical guarantees for time-average and rolling-window estimators

- ▶ Let $Q_n(\cdot)$ be sample discrepancy function s.t. $\hat{\theta} = \arg \min_{\theta \in \Theta} \|Q_n(\theta)\|$
- ▶ Let V be the asymptotic covariance matrix from $\sqrt{n}Q_n(\theta_0) \Rightarrow N(0, V)$
- ▶ Let $Q(\cdot)$ be population discrepancy function s.t. $Q_n(\theta) \xrightarrow{P} Q(\theta)$ uniformly over Θ
- ▶ Let Γ be the Jacobian matrix of $Q(\cdot)$ at θ_0
- ▶ We prove consistency $\hat{\theta} \xrightarrow{P} \theta_0$ and asymptotic normality

$$\sqrt{n}(\hat{\theta} - \theta_0) \Rightarrow N\left(0, (\Gamma^\top \Gamma)^{-1} \Gamma^\top V \Gamma (\Gamma^\top \Gamma)^{-1}\right)$$

Theoretical guarantees for time-average and rolling-window estimators

- ▶ Let $Q_n(\cdot)$ be sample discrepancy function s.t. $\hat{\theta} = \arg \min_{\theta \in \Theta} \|Q_n(\theta)\|$
- ▶ Let V be the asymptotic covariance matrix from $\sqrt{n}Q_n(\theta_0) \Rightarrow N(0, V)$
- ▶ Let $Q(\cdot)$ be population discrepancy function s.t. $Q_n(\theta) \xrightarrow{P} Q(\theta)$ uniformly over Θ
- ▶ Let Γ be the Jacobian matrix of $Q(\cdot)$ at θ_0
- ▶ We prove consistency $\hat{\theta} \xrightarrow{P} \theta_0$ and asymptotic normality

$$\sqrt{n}(\hat{\theta} - \theta_0) \Rightarrow N\left(0, (\Gamma^\top \Gamma)^{-1} \Gamma^\top V \Gamma (\Gamma^\top \Gamma)^{-1}\right)$$

- ▶ Guarantees for time-average and rolling-window estimators

Theoretical guarantees for time-average and rolling-window estimators

- ▶ Let $Q_n(\cdot)$ be sample discrepancy function s.t. $\hat{\theta} = \arg \min_{\theta \in \Theta} \|Q_n(\theta)\|$
- ▶ Let V be the asymptotic covariance matrix from $\sqrt{n}Q_n(\theta_0) \Rightarrow N(0, V)$
- ▶ Let $Q(\cdot)$ be population discrepancy function s.t. $Q_n(\theta) \xrightarrow{P} Q(\theta)$ uniformly over Θ
- ▶ Let Γ be the Jacobian matrix of $Q(\cdot)$ at θ_0
- ▶ We prove consistency $\hat{\theta} \xrightarrow{P} \theta_0$ and asymptotic normality

$$\sqrt{n}(\hat{\theta} - \theta_0) \Rightarrow N\left(0, (\Gamma^\top \Gamma)^{-1} \Gamma^\top V \Gamma (\Gamma^\top \Gamma)^{-1}\right)$$

- ▶ Guarantees for time-average and rolling-window estimators
- ▶ Examples of nonstationary time series:
 1. ODEs with *temporally dependent* observational noise
 2. State-space models with controls
 3. Multiplicative trend-seasonality-noise decomposition with latent AR process

Paradigm shift in prediction: Time series foundation models (TSFMs)

- ▶ Inspired by success of transformer-based foundation models in language and vision

Paradigm shift in prediction: Time series foundation models (TSFMs)

- ▶ Inspired by success of transformer-based foundation models in language and vision
- ▶ Many TSFMs are *pre-trained* on synthetic data generated by *parametric* models

Paradigm shift in prediction: Time series foundation models (TSFMs)

- ▶ Inspired by success of transformer-based foundation models in language and vision
- ▶ Many TSFMs are *pre-trained* on synthetic data generated by *parametric* models
- ▶ Renewed interest in finding parametric models that generate realistic time series

Paradigm shift in prediction: Time series foundation models (TSFMs)

- ▶ Inspired by success of transformer-based foundation models in language and vision
- ▶ Many TSFMs are *pre-trained* on synthetic data generated by *parametric* models
- ▶ Renewed interest in finding parametric models that generate realistic time series
- ▶ X. Liu et al. (2025) argue that “adaptively fine-tuning” TSFMs on “carefully designed synthetic data” remains a “missed opportunity”

Paradigm shift in prediction: Time series foundation models (TSFMs)

- ▶ Inspired by success of transformer-based foundation models in language and vision
- ▶ Many TSFMs are *pre-trained* on synthetic data generated by *parametric* models
- ▶ Renewed interest in finding parametric models that generate realistic time series
- ▶ X. Liu et al. (2025) argue that “adaptively fine-tuning” TSFMs on “carefully designed synthetic data” remains a “missed opportunity”
- ▶ Statistical foundations are unclear, especially for nonstationary time series

Paradigm shift in prediction: Time series foundation models (TSFMs)

- ▶ Inspired by success of transformer-based foundation models in language and vision
- ▶ Many TSFMs are *pre-trained* on synthetic data generated by *parametric* models
- ▶ Renewed interest in finding parametric models that generate realistic time series
- ▶ X. Liu et al. (2025) argue that “adaptively fine-tuning” TSFMs on “carefully designed synthetic data” remains a “missed opportunity”
- ▶ Statistical foundations are unclear, especially for nonstationary time series
- ▶ We develop a method for adaptively fine-tuning TSFMs with synthetic data

Paradigm shift in prediction: Time series foundation models (TSFMs)

- ▶ Inspired by success of transformer-based foundation models in language and vision
- ▶ Many TSFMs are *pre-trained* on synthetic data generated by *parametric* models
- ▶ Renewed interest in finding parametric models that generate realistic time series
- ▶ X. Liu et al. (2025) argue that “adaptively fine-tuning” TSFMs on “carefully designed synthetic data” remains a “missed opportunity”
- ▶ Statistical foundations are unclear, especially for nonstationary time series
- ▶ We develop a method for adaptively fine-tuning TSFMs with synthetic data
- ▶ To begin, we explain classical simulation-and-regression methods

Example to keep in mind: State-space model

- ▶ Decomposes observed time series into latent states, and noise

Example to keep in mind: State-space model

- ▶ Decomposes observed time series into latent states, and noise
- ▶ For times $t = 1, \dots, n$ and some $\theta \in \Theta \subset \mathbb{R}^p$, let

$$\begin{aligned} Y_t &= M_t(X_t, \varepsilon_t^Y, \theta), & (\varepsilon_t^X, \varepsilon_t^Y) &\stackrel{\text{iid}}{\sim} P_\theta(\varepsilon_t^X, \varepsilon_t^Y), \\ X_t &= S_t(X_{t-1}, \varepsilon_t^X, \theta), & X_0 &\sim P_\theta(X_0), \end{aligned}$$

where M_t is a measurement function and S_t is a state transition function

Example to keep in mind: State-space model

- ▶ Decomposes observed time series into latent states, and noise
- ▶ For times $t = 1, \dots, n$ and some $\theta \in \Theta \subset \mathbb{R}^p$, let

$$\begin{aligned} Y_t &= M_t(X_t, \varepsilon_t^Y, \theta), & (\varepsilon_t^X, \varepsilon_t^Y) &\stackrel{\text{iid}}{\sim} P_\theta(\varepsilon_t^X, \varepsilon_t^Y), \\ X_t &= S_t(X_{t-1}, \varepsilon_t^X, \theta), & X_0 &\sim P_\theta(X_0), \end{aligned}$$

where M_t is a measurement function and S_t is a state transition function

- ▶ **Smoothing:** estimate past latent states with current information

Example to keep in mind: State-space model

- ▶ Decomposes observed time series into latent states, and noise
- ▶ For times $t = 1, \dots, n$ and some $\theta \in \Theta \subset \mathbb{R}^p$, let

$$\begin{aligned} Y_t &= M_t(X_t, \varepsilon_t^Y, \theta), & (\varepsilon_t^X, \varepsilon_t^Y) &\stackrel{\text{iid}}{\sim} P_\theta(\varepsilon_t^X, \varepsilon_t^Y), \\ X_t &= S_t(X_{t-1}, \varepsilon_t^X, \theta), & X_0 &\sim P_\theta(X_0), \end{aligned}$$

where M_t is a measurement function and S_t is a state transition function

- ▶ **Smoothing:** estimate past latent states with current information
- ▶ **Forecasting:** predict future observations or states with current information

What simulation-and-regression methods usually do

- 1 Estimate $\hat{\theta}$ from Y^{obs} , then simulate s time series $\{(X^{(r)}(\hat{\theta}), Y^{(r)}(\hat{\theta}))\}_{r \in [s]}$

What simulation-and-regression methods usually do

- ① Estimate $\hat{\theta}$ from Y^{obs} , then simulate s time series $\{(X^{(r)}(\hat{\theta}), Y^{(r)}(\hat{\theta}))\}_{r \in [s]}$
- ② Choose hyperparameters and learn $\hat{f}$ by minimizing

$$\frac{1}{s} \sum_{r \in [s]} \left[\sum_{t \in \mathcal{T}} L(X_t^{(r)} - [f(Y^{(r)})]_t) \right]$$

What simulation-and-regression methods usually do

- ① Estimate $\hat{\theta}$ from Y^{obs} , then simulate s time series $\{(X^{(r)}(\hat{\theta}), Y^{(r)}(\hat{\theta}))\}_{r \in [s]}$
- ② Choose hyperparameters and learn $\hat{f}$ by minimizing

$$\frac{1}{s} \sum_{r \in [s]} \left[\sum_{t \in \mathcal{T}} L(X_t^{(r)} - [f(Y^{(r)})]_t) \right]$$

- ③ Apply $\hat{f}(Y^{\text{obs}})$ to estimate the latent states (e.g., conditional quantiles)

What simulation-and-regression methods usually do

- 1 Estimate $\hat{\theta}$ from Y^{obs} , then simulate s time series $\{(X^{(r)}(\hat{\theta}), Y^{(r)}(\hat{\theta}))\}_{r \in [s]}$
- 2 Choose hyperparameters and learn $\hat{f}$ by minimizing

$$\frac{1}{s} \sum_{r \in [s]} \left[\sum_{t \in \mathcal{T}} L(X_t^{(r)} - [f(Y^{(r)})]_t) \right]$$

- 3 Apply $\hat{f}(Y^{\text{obs}})$ to estimate the latent states (e.g., conditional quantiles)
- Aims to minimize Monte Carlo approximation to

$$R_{P_{\hat{\theta}}}(f) = \mathbb{E}_{P_{\hat{\theta}}} \left[\sum_{t \in \mathcal{T}} L(X_t - [f(Y)]_t) \right]$$

where expectation is w.r.t. distribution $P_{\hat{\theta}}$ of time series X, Y corresponding to $\hat{\theta}$

What simulation-and-regression methods usually do

- 1 Estimate $\hat{\theta}$ from Y^{obs} , then simulate s time series $\{(X^{(r)}(\hat{\theta}), Y^{(r)}(\hat{\theta}))\}_{r \in [s]}$
- 2 Choose hyperparameters and learn $\hat{f}$ by minimizing

$$\frac{1}{s} \sum_{r \in [s]} \left[\sum_{t \in \mathcal{T}} L(X_t^{(r)} - [f(Y^{(r)})]_t) \right]$$

- 3 Apply $\hat{f}(Y^{\text{obs}})$ to estimate the latent states (e.g., conditional quantiles)
- Aims to minimize Monte Carlo approximation to

$$R_{P_{\hat{\theta}}}(f) = \mathbb{E}_{P_{\hat{\theta}}} \left[\sum_{t \in \mathcal{T}} L(X_t - [f(Y)]_t) \right]$$

where expectation is w.r.t. distribution $P_{\hat{\theta}}$ of time series X, Y corresponding to $\hat{\theta}$

- But $\hat{\theta} \neq \theta_0$, so the minimizer $f_{P_{\hat{\theta}}}^*$ in the function class of $R_{P_{\hat{\theta}}}(f)$ will be “wrong”

Our method: Simulation-and-regression with parameter uncertainty

- ▶ Let $\widehat{\text{CD}}_n$ be an estimated *confidence distribution* over Θ for $\theta_0 \in \mathbb{R}^p$

Our method: Simulation-and-regression with parameter uncertainty

- ▶ Let $\widehat{\text{CD}}_n$ be an estimated *confidence distribution* over Θ for $\theta_0 \in \mathbb{R}^p$
- ▶ *Frequentist* theory from D. Liu, R. Y. Liu, and Min-ge Xie (2022) covers:
 1. $N(\hat{\theta}, \hat{\Sigma}/n)$ where $\hat{\Sigma}$ is the estimate of the asymptotic covariance
 2. Bootstrap distribution, i.e. sampling distribution of re-estimates

Our method: Simulation-and-regression with parameter uncertainty

- ▶ Let $\widehat{\text{CD}}_n$ be an estimated *confidence distribution* over Θ for $\theta_0 \in \mathbb{R}^p$
- ▶ *Frequentist* theory from D. Liu, R. Y. Liu, and Min-ge Xie (2022) covers:
 1. $N(\hat{\theta}, \hat{\Sigma}/n)$ where $\hat{\Sigma}$ is the estimate of the asymptotic covariance
 2. Bootstrap distribution, i.e. sampling distribution of re-estimates
- ▶ We propose minimizing the Monte Carlo approximation to

$$R_{P_{\widehat{\text{mix}}}}(f) = \mathbb{E}_{\theta \sim \widehat{\text{CD}}_n} \left[\mathbb{E}_{P_\theta} \left[\sum_{t \in \mathcal{T}} L(X_t - [f(Y)]_t) \right] \right] = \mathbb{E}_{P_{\widehat{\text{mix}}}} \left[\sum_{t \in \mathcal{T}} L(X_t - [f(Y)]_t) \right]$$

for some loss L , where the expectation is w.r.t. the mixture distribution $P_{\widehat{\text{mix}}} = \int_{\Theta} P_\theta d\widehat{\text{CD}}_n(\theta)$ of time series X, Y yielded by $\widehat{\text{CD}}_n$

Our method: Simulation-and-regression with parameter uncertainty

- ▶ Let $\widehat{\text{CD}}_n$ be an estimated *confidence distribution* over Θ for $\theta_0 \in \mathbb{R}^p$
- ▶ *Frequentist* theory from D. Liu, R. Y. Liu, and Min-ge Xie (2022) covers:
 1. $N(\hat{\theta}, \hat{\Sigma}/n)$ where $\hat{\Sigma}$ is the estimate of the asymptotic covariance
 2. Bootstrap distribution, i.e. sampling distribution of re-estimates
- ▶ We propose minimizing the Monte Carlo approximation to

$$R_{P_{\widehat{\text{mix}}}}(f) = \mathbb{E}_{\theta \sim \widehat{\text{CD}}_n} \left[\mathbb{E}_{P_\theta} \left[\sum_{t \in \mathcal{T}} L(X_t - [f(Y)]_t) \right] \right] = \mathbb{E}_{P_{\widehat{\text{mix}}}} \left[\sum_{t \in \mathcal{T}} L(X_t - [f(Y)]_t) \right]$$

for some loss L , where the expectation is w.r.t. the mixture distribution

$P_{\widehat{\text{mix}}} = \int_{\Theta} P_\theta d\widehat{\text{CD}}_n(\theta)$ of time series X, Y yielded by $\widehat{\text{CD}}_n$

- ▶ When $\widehat{\text{CD}}_n$ is uniform over Θ , this is like TSFM pre-training with synthetic data

Our method: Simulation-and-regression with parameter uncertainty

- ▶ Let $\widehat{\text{CD}}_n$ be an estimated *confidence distribution* over Θ for $\theta_0 \in \mathbb{R}^p$
- ▶ *Frequentist* theory from D. Liu, R. Y. Liu, and Min-ge Xie (2022) covers:
 1. $N(\hat{\theta}, \hat{\Sigma}/n)$ where $\hat{\Sigma}$ is the estimate of the asymptotic covariance
 2. Bootstrap distribution, i.e. sampling distribution of re-estimates
- ▶ We propose minimizing the Monte Carlo approximation to

$$R_{P_{\widehat{\text{mix}}}}(f) = \mathbb{E}_{\theta \sim \widehat{\text{CD}}_n} \left[\mathbb{E}_{P_\theta} \left[\sum_{t \in \mathcal{T}} L(X_t - [f(Y)]_t) \right] \right] = \mathbb{E}_{P_{\widehat{\text{mix}}}} \left[\sum_{t \in \mathcal{T}} L(X_t - [f(Y)]_t) \right]$$

for some loss L , where the expectation is w.r.t. the mixture distribution $P_{\widehat{\text{mix}}} = \int_{\Theta} P_\theta d\widehat{\text{CD}}_n(\theta)$ of time series X, Y yielded by $\widehat{\text{CD}}_n$

- ▶ When $\widehat{\text{CD}}_n$ is uniform over Θ , this is like TSFM pre-training with synthetic data
- ▶ When $\widehat{\text{CD}}_n$ is Dirac measure at $\hat{\theta}$, this is like classical simulation-and-regression

Future directions

- ▶ Forecasting and smoothing with transformers trained on synthetic time series

Future directions

- ▶ Forecasting and smoothing with transformers trained on synthetic time series
- ▶ Minimize sum of multiple quantile losses (with crossing penalty), so that the transformer outputs multiple quantile predictions

Future directions

- ▶ Forecasting and smoothing with transformers trained on synthetic time series
- ▶ Minimize sum of multiple quantile losses (with crossing penalty), so that the transformer outputs multiple quantile predictions
- ▶ **Research question:** Does our method (1) improve the calibration of pre-trained transformers, and (2) yield well-calibrated transformers when trained from scratch?

Future directions

- ▶ Forecasting and smoothing with transformers trained on synthetic time series
- ▶ Minimize sum of multiple quantile losses (with crossing penalty), so that the transformer outputs multiple quantile predictions
- ▶ **Research question:** Does our method (1) improve the calibration of pre-trained transformers, and (2) yield well-calibrated transformers when trained from scratch?
- ▶ **Primary objectives until defense:** Develop method, evaluate performance with extensive forecasting and smoothing experiments, write paper, and submit paper

Future directions

- ▶ Forecasting and smoothing with transformers trained on synthetic time series
- ▶ Minimize sum of multiple quantile losses (with crossing penalty), so that the transformer outputs multiple quantile predictions
- ▶ **Research question:** Does our method (1) improve the calibration of pre-trained transformers, and (2) yield well-calibrated transformers when trained from scratch?
- ▶ **Primary objectives until defense:** Develop method, evaluate performance with extensive forecasting and smoothing experiments, write paper, and submit paper
- ▶ *After submitting the main paper*, if time allows, some possible extensions include:
 - Pre-training vs training from scratch
 - Parameter confidence distribution vs parameter point estimate
 - Application to epidemiology

Thank you!

Questions and comments are welcome
You can reach me at: mwiecksosa@cmu.edu

Michael Wieck-Sosa
Department of Statistics & Data Science
Carnegie Mellon University

Main ideas

- ▶ **Research question:** Does incorporating parameter uncertainty improve the calibration of pre-trained/untrained transformers that output quantile predictions?
- ▶ We would like the quantile predictions to have the correct coverage
- ▶ That is, the probability that the time series, at a particular time, is less than or equal to the predicted α -quantile, at that same time, is *actually* equal to α
- ▶ For smoothing, we assess this with synthetic experiments
- ▶ For forecasting, we assess this with real data and synthetic experiments

Multi-quantile loss function

- Denote the times by $\mathcal{T}$, and number of times by $T = |\mathcal{T}|$, dimensions by d , quantile levels by K , and simulations by s (in practice, use batches)
- Multi-quantile regression heads by minimizing sum of multiple quantile losses

$$\frac{1}{s} \sum_{r \in [s]} \left[\sum_{t \in \mathcal{T}} \left(\sum_{j \in [d]} \left[\sum_{k \in [K]} \rho_{\tau_k} \left(X_{t,j}^{(r)} - [f(Y^{(r)})]_{t,j,\tau_k} \right) + Q_{\lambda}(r, t, j) \right] \right) \right]$$

where $\rho_{\tau_k}(u) = (\tau_k - \mathbb{I}[u < 0])u$, each $\tau_k \in (0, 1)$, and $\mathbb{I}$ is the indicator function

- Want predicted quantiles at lower levels to be less than those at higher levels, i.e. denoting $\hat{q}_{t,j}^{\text{obs}}(\tau_k) \equiv [\hat{f}(Y^{\text{obs}})]_{t,j,\tau_k}$ we want $\hat{q}_{t,j}^{\text{obs}}(\tau_1) \leq \hat{q}_{t,j}^{\text{obs}}(\tau_2) \leq \dots \leq \hat{q}_{t,j}^{\text{obs}}(\tau_K)$, so we penalize **quantile crossing** by adding the penalty

$$Q_{\lambda}(r, t, j) = \lambda \sum_{k=1}^{K-1} \max \left(0, [f(Y^{(r)})]_{t,j,\tau_k} - [f(Y^{(r)})]_{t,j,\tau_{k+1}} \right)$$

for some tuning parameter $\lambda > 0$, where the levels are ordered $\tau_1 \leq \tau_2 \leq \dots \leq \tau_K$

Algorithm for calibration experiments

- 1 Estimate confidence distribution $\widehat{\text{CD}}_n$ using the observed time series Y^{obs}
- 2 For $r \in [s]$, draw $\theta^{(r)} \sim \widehat{\text{CD}}_n$ and simulate $(X^{(r)}, Y^{(r)})_{r \in [s]}$ at $\theta^{(r)}$
- 3 Select hyperparameters (e.g., K -fold or 80/20 train/val over simulations)
- 4 For multi-quantile loss L , obtain predictor $\hat{f}$ by minimizing

$$\frac{1}{s} \sum_{r \in [s]} \left[\sum_{t \in \mathcal{T}} L \left(X_t^{(r)} - [f(Y^{(r)})]_t \right) \right]$$

over some seq2seq function class (e.g., transformers)

- 5 Apply seq2seq function $\hat{f}(Y^{\text{obs}})$ to predict latent states at all times based on the multiple predicted quantiles

Setup for calibration metrics

- ▶ To assess calibration in the context of quantile forecasting nonstationary time series, we adapt the metrics from Adler et al. (2025), which use time-averages
- ▶ **Main difference:** We average each metric over s simulations from a model
- ▶ When forecasting observations of real time series, we can use rolling averages
- ▶ When forecasting observations, just let X denote the forecasting target
- ▶ We present the metrics for the univariate setting, i.e. $d = 1$, to simplify the notation, and for $d \in \mathbb{N}$, we can aggregate metrics by averaging over dimensions
- ▶ Let $\mathcal{T}$ to be the times with $T = |\mathcal{T}|$, Y be the observed time series,
 $\hat{q}_t^{(r)}(\tau_k) \equiv [\hat{f}(Y^{(r)})]_{t, \tau_k}$ be the r -th predicted quantile at time t for level $\tau_k \in \{0.1, 0.2, \dots, 0.9\}$, and $\Delta = \{(0.1, 0.9), (0.2, 0.8), (0.3, 0.7), (0.4, 0.6)\}$

Calibration metric 1

- For $t \in \mathcal{T}$, define the *Probabilistic Calibration Error (PCE)* as:

$$\text{PCE}(t) = \frac{1}{K} \sum_{k \in [K]} \left| \tau_k - \frac{1}{s} \sum_{r=1}^s \mathbb{I} \left[X_t^{(r)} \leq \hat{q}_t^{(r)}(\tau_k) \right] \right|$$

- Lower values indicate better-calibrated models:
- 1 PCE is lower-bounded by 0 (i.e. predicted quantile above)
 - 2 PCE is upper-bounded by 0.5 (i.e. predicted quantile below)

Calibration metric 2

- ▶ Well-calibrated models do not necessarily yield useful forecasts
- ▶ A model can predict very large quantiles and be calibrated
- ▶ Need metric that captures sharpness (i.e. concentration of predictive distribution)
- ▶ Scale width of prediction interval by width of corresponding interval using unconditional quantiles $\bar{q}_t^{(r)}(\cdot)$ estimated via order statistics of s simulations
- ▶ For $t \in \mathcal{T}$, define the *Scaled Interval Width (SIW)* as:

$$\text{SIW}(t) = \frac{1}{|\Delta|} \sum_{(\tau_\ell, \tau_h) \in \Delta} \left[\frac{1}{s} \sum_{r=1}^s \frac{\hat{q}_t^{(r)}(\tau_h) - \hat{q}_t^{(r)}(\tau_\ell)}{\bar{q}_t^{(r)}(\tau_h) - \bar{q}_t^{(r)}(\tau_\ell)} \right]$$

- ▶ Lower (resp. higher) values correspond to tighter (resp. wider) prediction intervals

Calibration metric 3

- ▶ We introduce a metric to help check whether a model is under or over confident
- ▶ For $t \in \mathcal{T}$, define the *Centered Calibration Error (CCE)* as:

$$\text{CCE}(t) = \frac{1}{|\Delta|} \sum_{(\tau_\ell, \tau_h) \in \Delta} \left((\tau_h - \tau_\ell) - \frac{1}{s} \sum_{r=1}^s \mathbb{I} \left[\hat{q}_t^{(r)}(\tau_\ell) \leq X_t^{(r)} \leq \hat{q}_t^{(r)}(\tau_h) \right] \right)$$

- ▶ Positive CCE values mean that more observations are outside the prediction interval than expected
- ▶ When combined with a low SIW value, this suggests a model is overconfident, i.e. that the prediction intervals are too narrow

Calibration

- ▶ For improving calibration with $\widehat{CD}_n$ *beyond* doing simulation-and-regression and fine-tuning based on the parameter estimate $\hat{\theta}$, we can reduce this to the other questions by instantiating the loss function L as the multi-quantile loss
- ▶ For untrained transformers, refer to *simulation-based training*
- ▶ For pre-trained transformers, refer to *pre-training*

Simulation-based training

- ▶ **Question:** In the context of simulation-and-regression, does accounting for parameter uncertainty achieve better performance than using parameter estimate?
- ▶ When does this inequality hold:

$$\mathbb{E}_{P_{\theta_0}} \left[\sum_{t \in \mathcal{T}} L \left(X_t - [f_{\widehat{P}_{\text{mix}}}^*(Y)]_t \right) \right] < \mathbb{E}_{P_{\theta_0}} \left[\sum_{t \in \mathcal{T}} L \left(X_t - [f_{\widehat{P}_{\hat{\theta}}}^*(Y)]_t \right) \right]$$

where both expectations are with respect to the distribution P_{θ_0} of the time series corresponding to the unknown true parameter θ_0 , $f_{\widehat{P}_{\text{mix}}}^*$ is the minimizer in the function class of $R_{\widehat{P}_{\text{mix}}}(f)$, and $f_{\widehat{P}_{\hat{\theta}}}^*$ is the minimizer in the function class of $R_{\widehat{P}_{\hat{\theta}}}(f)$

Pre-training

- ▶ **Question:** When does a pre-trained transformer fine-tuned with our method outperform a transformer trained from scratch with our method?
- ▶ Intuitively, this occurs when the following conditions hold
- ▶ Large parameter uncertainty, e.g., when the sample size of the observed time series is small relative to the noise
- ▶ The distributions used in the pre-training phase are close to the true distribution, e.g., when we have domain knowledge or we have previously observed time series with similar distributions

Connections of proposed method to Bayesian prediction

- ▶ When using the Gaussian depth-CD $N(\hat{\theta}, \hat{\Sigma}/n)$, our proposed method can be viewed as a frequentist, likelihood-free analogue of the Laplace approximation to the posterior distribution in Bayesian statistics
- ▶ Specifically, when using the Laplace approximation to the posterior distribution to obtain an approximate posterior predictive distribution
- ▶ Recall that the Laplace approximation uses a Gaussian distribution centered at the maximum a posteriori (MAP) estimate with covariance given by the inverse observed Fisher information to approximate the posterior distribution
- ▶ The Bernstein–von Mises theorem is used to justify the Laplace approximation
- ▶ In contrast, to justify our approach, we use asymptotic theory for likelihood-free parameter estimators for parametric time series models, nonstationary time series theory (Mies 2023), and the theory of confidence distributions

Definition of confidence distribution

- ▶ In general, our frequentist procedure can be used with any *confidence distribution* over $\Theta \subset \mathbb{R}^p$ which expresses our uncertainty about the $\theta_0 \in \mathbb{R}^p$
- ▶ Ideally, this distribution is easy to simulate from
- ▶ We give the definition of confidence distributions for the $p = 1$ case, and the generalization to the $p \in \mathbb{N}$ case
- ▶ In the definition, $\Theta \subset \mathbb{R}^p$ is the parameter space and $\mathcal{X}_n$ is the sample space corresponding to the data $X_{1:n}$

Definition of confidence distribution

Definition (Confidence Distribution, $p = 1$, (Min-ge Xie and Singh 2013))

A function $\widehat{\text{CD}}_n(\cdot) = \widehat{\text{CD}}_n(X_{1:n}, \cdot)$ on $\mathcal{X}_n \times \Theta \rightarrow [0, 1]$ is called a confidence distribution (CD) for a parameter θ , if:

- ① For each given $X_{1:n}$, $\widehat{\text{CD}}_n(\cdot) = \widehat{\text{CD}}_n(X_{1:n}, \cdot)$ is a CDF on Θ
 - ② At the true parameter value $\theta = \theta_0$, $\widehat{\text{CD}}_n(\theta_0) \equiv \widehat{\text{CD}}_n(X_{1:n}, \theta_0)$, as a function of the sample $X_{1:n}$, follows the uniform distribution $U[0, 1]$
- Also, the function $\widehat{\text{CD}}_n(\cdot)$ is an asymptotic confidence distribution (aCD), if the $U[0, 1]$ requirement is true only asymptotically

Definition of confidence distribution

- ▶ The definition of confidence distributions for the general $p \in \mathbb{N}$ case involves depth functions
- ▶ As discussed in R. Y. Liu, Parelius, and Singh (1999), a data depth is a way of measuring how deep or central a point $x \in \mathbb{R}^d$ is with respect to a probability distribution P on $\mathbb{R}^d$, $d \in \mathbb{N}$, or with respect to a given data cloud $X_{1:n}$
- ▶ In general, a point with a larger depth value means that it is closer to the center of the distribution or more central in a data cloud, while points with smaller depth values are farther away from the center, so data depth is a center-outward ordering of multivariate observations

Definition of confidence distribution

- ▶ For example, the Mahalanobis depth (D_M) (Mahalanobis 1936) at $x \in \mathbb{R}^d$ with respect to P is defined as

$$D_M(P, x) = \frac{1}{1 + (x - \mu_P)^\top \Sigma_P^{-1} (x - \mu_P)},$$

where μ_P and Σ_P are the mean vector and the dispersion matrix of P , respectively

- ▶ We obtain the sample version of D_M by replacing μ_P and Σ_P by the corresponding sample estimates $\hat{\mu}_P$ and $\hat{\Sigma}_P$
- ▶ See R. Y. Liu, Parelius, and Singh (1999) for more examples of data depth

Definition of confidence distribution

Definition (Confidence Distribution, $p \in \mathbb{N}$, (D. Liu, R. Y. Liu, and Min-ge Xie 2022))

A function $\widehat{\text{CD}}_n(\cdot) \equiv \widehat{\text{CD}}_n(X_{1:n}, \cdot)$ on $\Theta \subset \mathbb{R}^p$ is called a depth confidence distribution (depth-CD) associated with depth function D for a vector-valued parameter θ , if:

- ① It is a distribution function on the parameter space Θ for any fixed sample set $X_{1:n}$
- ② The $(1 - \alpha)$ “central region” of the distribution $\widehat{\text{CD}}_n(\cdot)$,

$$\mathcal{R}_{1-\alpha}(\widehat{\text{CD}}_n) = \{\theta \in \Theta : C_{\widehat{\text{CD}}_n}(\theta) \geq \alpha\},$$

is a confidence region for θ with a coverage probability of $(1 - \alpha)$. Here, the centrality function associated with depth D and confidence distribution $\widehat{\text{CD}}_n(\cdot)$, that is, $C_{\widehat{\text{CD}}_n}(\theta)$, is also referred to as a depth confidence curve (depth-CV)

- If the statements in (2) hold only asymptotically, then we refer to $\widehat{\text{CD}}_n$ and $C_{\widehat{\text{CD}}_n}$ as asymptotic depth-CD and asymptotic depth-CV, respectively

Depth-CV for Gaussian depth-CD

- ▶ In this work, we focus on asymptotic Gaussian depth confidence distributions of the form $\widehat{\text{CD}}_n = N(\hat{\theta}, \hat{\Sigma}/n)$ for a multi-dimensional parameter $\theta \in \mathbb{R}^p$, $p \in \mathbb{N}$
- ▶ Let $q_{1-\alpha}$ be the $1 - \alpha$ quantile of χ_p^2 , i.e., $\mathbb{P}(\chi_p^2 \leq q_{1-\alpha}) = 1 - \alpha$
- ▶ As usual, the asymptotic $1 - \alpha$ confidence region for θ_0 is

$$\mathcal{R}_{1-\alpha}(\widehat{\text{CD}}_n) = \{\theta \in \Theta : (\theta - \hat{\theta})^\top (\hat{\Sigma}/n)^{-1} (\theta - \hat{\theta}) \leq q_{1-\alpha}\}$$

so applying the χ_p^2 CDF $F_{\chi_p^2}$, subtracting by 1, dividing by -1 , yields

$$\mathcal{R}_{1-\alpha}(\widehat{\text{CD}}_n) = \{\theta \in \Theta : C_{\widehat{\text{CD}}_n}(\theta) = 1 - F_{\chi_p^2}[(\theta - \hat{\theta})^\top (\hat{\Sigma}/n)^{-1} (\theta - \hat{\theta})] \geq \alpha\}$$

Result for confidence distributions

- We justify this Gaussianity through the asymptotic normality of the parameter estimator, nonstationary time series theory (Mies 2023), and the following result from D. Liu, R. Y. Liu, and Min-ge Xie (2022) (Theorem 2 therein)

Theorem (Depth-CDs from Pivot Statistics (D. Liu, R. Y. Liu, and Min-ge Xie 2022))

Assume that $A_n(X_{1:n})$ is a nonsingular matrix such that $A_n(X_{1:n})(\hat{\theta} - \theta)$ follows a distribution Q_n that is free of all unknown parameters. Also assume that η_n is a random vector, independent of the sample $X_{1:n}$, following the distribution Q_n . Then, conditional on $X_{1:n}$, the distribution function of $(\hat{\theta} - A_n(X_{1:n})^{-1}\eta_n)$ is a depth-CD for θ , under the following conditions:

- ① *The depth D is affine-invariant.*
- ② *The depth contours $\{\eta_n : D_{Q_n}(\eta_n) = \epsilon\}$ all have probability zero with respect to Q_n for any $\epsilon > 0$.*

Bootstrap distribution is a confidence distribution

- ▶ Alternatively, use parametric bootstrap (Efron 1979; Efron and Tibshirani 1994)
- ▶ To use a parametric bootstrap approach, we would estimate the parameter on the observed data, simulate from the model with the estimated parameter, then re-estimate the parameter for each simulated realization
- ▶ The sampling distribution of the re-estimates is used as the depth-CD
- ▶ Theorem 1 in D. Liu, R. Y. Liu, and Min-ge Xie (2022) shows that this bootstrap distribution is an asymptotic depth-CD under certain regularity conditions
- ▶ For more info about the connections between bootstrap distributions and confidence distributions, see Example 3 in Min-ge Xie and Singh (2013)

Reason we use Gaussian depth-CD instead of bootstrap depth-CD

- ▶ We opt for an asymptotic Gaussian depth-CD, rather than a bootstrap depth confidence distribution, to improve the computational efficiency of our method
- ▶ This computational efficiency is particularly important, because our proposed method is intended to be applicable to models with intractable likelihoods
- ▶ In this setting, parameter estimation typically relies on simulation-based procedures that are substantially more computationally demanding than methods like maximum likelihood estimation
- ▶ Therefore, our procedure shares some of the motivations behind the Krinsky and Robb (1986) method

Estimating the asymptotic covariance

- ▶ Let $Q_n(\cdot)$ be sample discrepancy function s.t. $\hat{\theta} = \arg \min_{\theta \in \Theta} \|Q_n(\theta)\|$
- ▶ Let V be the asymptotic covariance matrix from $\sqrt{n}Q_n(\theta_0) \Rightarrow N(0, V)$
- ▶ Let $Q(\cdot)$ be population discrepancy function s.t. $Q_n(\theta) \xrightarrow{p} Q(\theta)$ uniformly over Θ
- ▶ Let Γ be the Jacobian matrix of $Q(\cdot)$ at θ_0
- ▶ We prove consistency $\hat{\theta} \xrightarrow{p} \theta_0$ and asymptotic normality

$$\sqrt{n}(\hat{\theta} - \theta_0) \Rightarrow N\left(0, (\Gamma^\top \Gamma)^{-1} \Gamma^\top V \Gamma (\Gamma^\top \Gamma)^{-1}\right)$$

- ▶ We need to estimate $(\Gamma^\top \Gamma)^{-1} \Gamma^\top V \Gamma (\Gamma^\top \Gamma)^{-1}$
- ▶ First, we estimate Γ using numerical differentiation tools (to be determined)
- ▶ Second, we estimate V by (slightly) extending covariance estimation results from nonstationary time series theory Mies and Steland (2023) and Mies (2023)
- ▶ For both estimators, use plug-in approach with $\hat{\theta}$ in place of the unknown θ_0

How to get synthetic training data for forecasting with TSFMs

- ▶ For forecasting with TSFMs, Dooley et al. (2023) pre-train by simulating from *multiplicative seasonality-trend-noise decomposition models* with many $\theta \in \Theta$
- ▶ We consider an extension with multiplicative autoregressive noise

$$Y_t = \text{Se} \left(t, \theta^{\text{Se}} \right) \text{Tr} \left(t, \theta^{\text{Tr}} \right) \exp \left(\kappa X_t - \frac{\kappa^2 \sigma_X^2}{2(1 - \phi^2)} \right) \exp \left(\varepsilon_t^Y - \frac{\sigma_Y^2}{2} \right)$$
$$X_t = \phi X_{t-1} + \varepsilon_t^X$$

for some scaling factor $\kappa > 0$ and autoregressive coefficient $\phi \in (-1, 1)$, with $X_0 \sim N(0, \frac{\sigma_X^2}{1-\phi^2})$ and $\varepsilon_t^X \stackrel{\text{iid}}{\sim} N(0, \sigma_X^2)$, $\varepsilon_t^Y \stackrel{\text{iid}}{\sim} N(0, \sigma_Y^2)$ for some $\sigma_X, \sigma_Y > 0$

- ▶ **Seasonality:** linear combination of sines, cosines (weekly, monthly, yearly periods)
- ▶ **Trend:** product of linear trend and exponential trend
- ▶ Our rolling-window estimator can be used (consistent and asymptotically normal)


Model we use for smoothing application in epidemiology

- ▶ Discrete-valued process driven by a latent autoregressive process and controls
- ▶ We condition on the controls, so that we may view each control as a discretely observed function of continuous time $Z_t \equiv Z(t/n)$ that is nice (e.g., Hölder)
- ▶ We define a (conditional) state-space model for $(X_{1:n}, Y_{1:n})$ given $Z_{1:n}$ as follows

$$Y_t \mid X_t, Z_t \sim \text{Pois} \left(\exp \left(\kappa X_t + \beta^\top Z_t \right) \right)$$
$$X_t = \mu + \phi(X_{t-1} - \mu) + \varepsilon_t^X$$

- ▶ Effect sizes $\kappa \in \mathbb{R}$, $\beta \in \mathbb{R}^{dz}$, autoregressive coefficient $\phi \in (-1, 1)$, and mean $\mu \in \mathbb{R}$, with $X_0 \sim N(\mu, \frac{\sigma_X^2}{1-\phi^2})$, $\varepsilon_t^X \stackrel{\text{iid}}{\sim} N(0, \sigma_X^2)$ indep. of $Z_{1:n}$, for some $\sigma_X > 0$
- ▶ $Y_{1:n}$ are observed case counts, $X_{1:n}$ is a *latent infectivity* process, and $Z_{1:n}$ are covariates (e.g., vaccination coverage, variant relative frequency, mobility)
- ▶ Our rolling-window estimator can be used (consistent and asymptotically normal)

COVID-19 time series May 2020–2022 ($n=365 \times 2=730$, $d=4$)



Nonstationarity and infill asymptotics

- ▶ When there is strong nonstationarity (i.e. not asymptotically stationary), we introduce rescaled time $u \in [0, 1]$ and use infill asymptotics so observations at t/n
- ▶ **Intuition:** more observations at each local structure of the process as n grows
- ▶ Recall: $(\varepsilon_i)_{i \in \mathbb{Z}}$ is an iid sequence of random variables, we denote $\varepsilon_t = (\varepsilon_t, \varepsilon_{t-1}, \dots)$, and Θ is a parameter space
- ▶ For each $n \in \mathbb{N}$, let the mappings $G_t^{(n)} : \mathbb{R}^\infty \times \Theta \rightarrow \mathbb{R}^d$, $t = 1, \dots, n$, and $\tilde{G}_u^{(n)} : \mathbb{R}^\infty \times \Theta \rightarrow \mathbb{R}^d$, $u \in [0, 1]$, define a *generative model* for time series

Assumption

Given a parameter $\theta \in \Theta$, index $n \in \mathbb{N}$ (linked to sample size and approximation), and noise inputs ε_t , a time series $X_{1:n}$ is generated by nature as

$$X_t = G_t^{(n)}(\varepsilon_t, \theta) = \tilde{G}_{t/n}^{(n)}(\varepsilon_t, \theta), \quad t = 1, \dots, n$$

Control of nonstationarity

- We measure the nonstationarity using a p -variation-type quantity, define

$$\left\| \left(\tilde{G}_u^{(i)}(\varepsilon_0, \theta) \right)_u \right\|_{p\text{-var}, \mathcal{L}^q(\theta)} = \sup_{\substack{0=u_0 < \dots < u_\ell=1 \\ \ell \in \mathbb{N}}} \left(\sum_{j=1}^{\ell} \left\| \tilde{G}_{u_j}^{(i)}(\varepsilon_0, \theta) - \tilde{G}_{u_{j-1}}^{(i)}(\varepsilon_0, \theta) \right\|_{\mathcal{L}^q(\theta)}^p \right)^{\frac{1}{p}}$$

- We impose the following regularity condition to control the nonstationarity

Assumption

For some $q > 2$, some $p \in [1, 2)$, and all $i \in \mathbb{N}$, there exists a constant $\Lambda > 0$ such that

$$\sup_{\theta \in \Theta} \sup_{u \in [0,1]} \left\| \tilde{G}_u^{(i)}(\varepsilon_0, \theta) \right\|_{\mathcal{L}^q(\theta)} + \sup_{\theta \in \Theta} \left\| \left(\tilde{G}_u^{(i)}(\varepsilon_0, \theta) \right)_u \right\|_{p\text{-var}, \mathcal{L}^q(\theta)} \leq \Lambda$$

Example of nonstationary time series

- ▶ Time-varying autoregressive (tvAR(1)) process
- ▶ Given parameter curves $\beta : [0, 1] \rightarrow (-1, 1)$ and $\sigma : [0, 1] \rightarrow \mathbb{R}$ define

$$X_t^{(n)} = \beta\left(\frac{t}{n}\right) X_{t-1}^{(n)} + \sigma\left(\frac{t}{n}\right) \varepsilon_t, \quad t = 1, \dots, n$$

where $(\varepsilon_t)_t$ are iid random variables

- ▶ Extend β, σ to domain $(-\infty, 1]$ by setting $\beta(u) = \beta(0)$ and $\sigma(u) = \sigma(0)$ for $u < 0$
- ▶ Substituting yields the time-varying moving average (tvMA(∞)) representation

$$X_t^{(n)} = \sum_{j=0}^{\infty} \psi_{t,j}^{(n)} \varepsilon_{t-j},$$

where $\psi_{t,0}^{(n)} = \sigma\left(\frac{t}{n}\right)$ and for $j \geq 1$

$$\psi_{t,j}^{(n)} = \left(\prod_{r=0}^{j-1} \beta\left(\frac{t-r}{n}\right) \right) \sigma\left(\frac{t-j}{n}\right)$$

Replication and Rosenblatt transform

- ▶ We can take the random variables $(\varepsilon_\ell)_{\ell \in \mathbb{Z}}$ to be $U[0, 1]$
- ▶ Standard results in probability theory imply that the causal representations based on measurable functions of sequences of iid $U[0, 1]$ are already sufficiently general
- ▶ First, by Kallenberg (2021) Lemma 4.21, Lemma 4.22, and the surrounding discussion, we can replicate each of the iid $U[0, 1]$ random variables by taking each $G_t^{(n)}$ to include compositions with measurable functions to replicate the original $U[0, 1]$, which yields d iid $U[0, 1]$ random variables $U_1, \dots, U_d$
- ▶ Second, we can use the Rosenblatt transform (i.e. for each $j \in [d]$, sample $X_{t,j}$ from the conditional quantile given $X_{t,1}, \dots, X_{t,j-1}$ by transforming U_j), see Section 4 of Wu and Mielniczuk (2010)

Choices of features

- ▶ **Method of simulated moments:** F based on functions of moments
 - ▶ Sample mean, variance, skewness, autocorrelation, and so on
- ▶ **Indirect inference:** F from parameter estimates in auxiliary model
 - ▶ To get $\hat{\theta}$ for MA(1): $X_t = \varepsilon_t + \theta_0 \varepsilon_{t-1}$
 - ▶ Use $F = \hat{\beta}$ from AR(1): $X_t = \beta X_{t-1} + \xi_t$
- ▶ **Synthetic likelihood:** F from user-chosen summary statistics
 - ▶ Coefficients from polynomial regression of X_t on $X_{t-1}, \dots, X_{t-\ell}$
- ▶ **Neural summary statistics:** F from neural network
 - ▶ Extract early layer of neural network

Random Fourier features

- ▶ Let $x_1, \dots, x_{m+1}$ be vectors in $\mathbb{R}^d$, and define $x = (x_1, \dots, x_{m+1})^\top \in \mathbb{R}^{(m+1) \times d}$
- ▶ We focus on random Fourier features, though our framework can be used with different random features
- ▶ Specifically, we consider random Fourier features of the form

$$\varphi_i(x) = \cos \left(\sum_{j=1}^{m+1} \Omega_{i,j} \cdot x_j + \alpha_i \right)$$

where $\Omega_{i,j} \stackrel{\text{iid}}{\sim} N(0, I_d)$ and $\alpha_i \stackrel{\text{iid}}{\sim} U(-\pi, \pi)$ for $i = 1, \dots, k$ and $j = 1, \dots, m+1$

- ▶ Denote all k of these functions by

$$\varphi = (\varphi_1, \dots, \varphi_k)$$

so that $\varphi : \mathbb{R}^{(m+1) \times d} \rightarrow \mathbb{R}^k$

Random Fourier features

- Define the i -th random feature of the observed and r -th simulated time series at time $t = m + 1, \dots, n$ as

$$f_{t,i}^{\text{obs}} = \varphi_i(X_{t-m:t}^{\text{obs}}), \quad f_{t,i}^{(r)}(\theta) = \varphi_i\left(X_{t-m:t}^{(r)}(\theta)\right)$$

- Write all k random features at time t as

$$\mathbf{f}_t^{\text{obs}} = \left(f_{t,1}^{\text{obs}}, \dots, f_{t,k}^{\text{obs}}\right), \quad \mathbf{f}_t^{(r)}(\theta) = \left(f_{t,1}^{(r)}(\theta), \dots, f_{t,k}^{(r)}(\theta)\right)$$

Time-average estimator

- Define the observed and simulated time-averages of the random features by

$$F^{\text{obs}} = \frac{1}{n-m} \sum_{t=m+1}^n f_t^{\text{obs}}, \quad \bar{F}^{\text{sim}}(\theta) = \frac{1}{n-m} \sum_{t=m+1}^n \bar{f}_t^{\text{sim}}(\theta)$$

where $\bar{f}_t^{\text{sim}}(\theta) = \frac{1}{s} \sum_{r=1}^s f_t^{(r)}(\theta)$. The time-average estimator is given by

$$\hat{\theta}^{\text{TA}} = \underset{\theta \in \Theta}{\operatorname{argmin}} \left\| \hat{Q}_n^{\text{TA}}(\theta) \right\|$$

where the sample discrepancy function is defined as

$$\hat{Q}_n^{\text{TA}}(\theta) = F^{\text{obs}} - \bar{F}^{\text{sim}}(\theta)$$

- In practice, an optimization procedure is used to find the minimizer of $\left\| \hat{Q}_n^{\text{TA}}(\cdot) \right\|$

Rolling-window estimator

- Define the observed and simulated rolling-window random features at time t by

$$F_t^{\text{obs}} = \frac{1}{w \wedge t} \sum_{j=(t-w) \vee 1}^t f_j^{\text{obs}}, \quad \bar{F}_t^{\text{sim}}(\theta) = \frac{1}{w \wedge t} \sum_{j=(t-w) \vee 1}^t \bar{f}_j^{\text{sim}}(\theta),$$

where $\bar{f}_t^{\text{sim}}(\theta) = \frac{1}{s} \sum_{r=1}^s f_t^{(r)}(\theta)$ and $w \in \mathbb{N}$ is the window size

Rolling-window estimator

- The rolling-window estimator is given by

$$\hat{\theta}^{\text{RW}} = \underset{\theta \in \Theta}{\operatorname{argmin}} \left| \hat{Q}_n^{\text{RW}}(\theta) \right|,$$

where the sample discrepancy function is defined as

$$\begin{aligned} \hat{Q}_n^{\text{RW}}(\theta) &= \frac{1}{n-m} \sum_{t=m+\tau+L}^n \left[\left\| F_t^{\text{obs}} - \bar{F}_t^{\text{sim}}(\theta) \right\|^2 + K_t(\theta) \right], \\ K_t(\theta) &= 2 \left(F_{t-L}^{\text{obs}} - \bar{F}_{t-L}^{\text{sim}}(\theta) \right)^{\top} \left(\left[f_t^{\text{obs}} - \bar{f}_t^{\text{sim}}(\theta) \right] - \left[F_{t-L}^{\text{obs}} - \bar{F}_{t-L}^{\text{sim}}(\theta) \right] \right), \end{aligned}$$

where $\tau \in \mathbb{N}$ is an initial time-offset and $L \in \mathbb{N}$ is a lag

- In practice, an optimization procedure is used to find the minimizer of $\left| \hat{Q}_n^{\text{RW}}(\cdot) \right|$

What do we mean by “almost every” function? Prevalent subset

Definition (J. Yorke and Ott (2005))

Let V be a completely metrizable topological vector space. A Borel set $E \subset V$ is said to be **prevalent** if there exists some Borel measure μ on V such that:

- ① $0 < \mu(C) < \infty$ for some compact subset C of V
- ② The set $E + v$ has full μ -measure (that is, the complement of $E + v$ has measure zero) for all $v \in V$

- ▶ In this literature, it is common to say that, if $E \subset V$ contains a prevalent Borel set, then “almost every” element of V lies in E
- ▶ In this sense, “almost every” C^1 smooth map from $\mathbb{R}^p$ to $\mathbb{R}^{2p+1}$ has the properties from the embedding theorem of Sauer, J. A. Yorke, and Casdagli (1991)

Definition of smooth manifold (University of Toronto Lecture Notes)

Definition

A C^∞ -manifold (i.e. smooth manifold or differentiable manifold) is a Hausdorff topological space M , with countable basis, together with a C^∞ -atlas

- The next slides explain the definitions needed to understand this

Definition of smooth manifold (University of Toronto Lecture Notes)

- ▶ An n -dimensional manifold is a space that locally looks like $\mathbb{R}^n$. To give a precise meaning to this idea, our space first of all has to come equipped with some topology (so that the word “local” makes sense)
- ▶ Recall that a *topological space* is a set M , together with a collection of subsets of M , called *open subsets*, satisfying the following three axioms: (i) the empty set $\emptyset$ and the space M itself are both open, (ii) the intersection of any finite collection of open subsets is open, (iii) the union of any collection of open subsets is open
- ▶ The collection of open subsets of M is also called the topology of M . A map $f : M_1 \rightarrow M_2$ between topological spaces is called *continuous* if the pre-image of any open subset in M_2 is open in M_1 . A continuous map with a continuous inverse is called a *homeomorphism*

Definition of smooth manifold (University of Toronto Lecture Notes)

- ▶ One ingredient in the definition of a manifold is that our topological space comes equipped with a covering by open sets which are homeomorphic to open subsets of $\mathbb{R}^n$
- ▶ Let M be a topological space. An n -dimensional *chart* for M is a pair (U, ϕ) consisting of an open subset $U \subset M$ and a continuous map $\phi : U \rightarrow \mathbb{R}^n$ such that ϕ is a homeomorphism onto its image $\phi(U)$. Two such charts (U_α, ϕ_α) , (U_β, ϕ_β) are C^∞ -compatible if the transition map

$$\phi_\beta \circ \phi_\alpha^{-1} : \phi_\alpha(U_\alpha \cap U_\beta) \rightarrow \phi_\beta(U_\alpha \cap U_\beta)$$

is a diffeomorphism (a smooth map with smooth inverse). A covering $\mathcal{A} = (U_\alpha)_{\alpha \in A}$ of M by pairwise C^∞ -compatible charts is called a C^∞ -atlas

Definition of smooth manifold (University of Toronto Lecture Notes)

- ▶ We recall that a topological space is called *Hausdorff* if any two points have disjoint open neighborhoods.
- ▶ A *basis* for a topological space M is a collection $\mathcal{B}$ of open subsets of M such that every open subset of M is a union of open subsets in the collection $\mathcal{B}$.
- ▶ A topological space with countable basis is also called *second countable*
- ▶ For example, the collection of open balls $B_\varepsilon(x)$ in $\mathbb{R}^n$ define a basis. But one already has a basis if one takes only all balls $B_\varepsilon(x)$ with $x \in \mathbb{Q}^n$ and $\varepsilon \in \mathbb{Q}_{>0}$; this defines a countable basis.

References I

-  Adler, Coen et al. (2025). “Beyond accuracy: Are time series foundation models well-calibrated?” [arXiv preprint arXiv:2510.16060](#).
-  Bodik, Juraj and Olivier C. Pasche (2024). “Granger causality in extremes”. [arXiv preprint arXiv: 2407.09632](#).
-  Dooley, Samuel et al. (2023). “Forecastpfm: Synthetically-trained zero-shot forecasting”. In: *Advances in Neural Information Processing Systems 36*, pp. 2403–2426.
-  Efron, Bradley (1979). “Bootstrap methods: Another look at the jackknife”. In: *Annals of Statistics* 7.1, pp. 1–26.
-  Efron, Bradley and Robert J. Tibshirani (1994). *An introduction to the bootstrap*. Chapman and Hall/CRC.
-  Kallenberg, Olav (2021). *Foundations of modern probability*. Third edition. Springer.

References II

-  Krinsky, Itzhak and A. Leslie Robb (1986). “On approximating the statistical properties of elasticities”. In: *The Review of Economics and Statistics*, pp. 715–719.
-  Liu, Dungang, Regina Y. Liu, and Min-ge Xie (2022). “Nonparametric fusion learning for multiparameters: Synthesize inferences from diverse sources using data depth and confidence distribution”. In: *Journal of the American Statistical Association* 117.540, pp. 2086–2104.
-  Liu, Regina Y., Jesse M. Parelius, and Kesar Singh (1999). “Multivariate analysis by data depth: descriptive statistics, graphics and inference”. In: *The Annals of Statistics* 27.3, pp. 783–858.
-  Liu, Xu et al. (2025). “Empowering time series analysis with synthetic data: A survey and outlook in the era of foundation models”. [arXiv preprint arXiv:2503.11411](#).
-  Mahalanobis, Prasanta Chandra (1936). “On the generalized distance in statistics”. In: *Proceedings of the National Institute of Science of India* 2, pp. 49–55.

References III

-  Mies, Fabian (2023). “Functional estimation and change detection for nonstationary time series”. In: *Journal of the American Statistical Association* 118.542, pp. 1011–1022.
-  Mies, Fabian and Ansgar Steland (2023). “Sequential Gaussian approximation for nonstationary time series in high dimensions”. In: *Bernoulli* 29.4, pp. 3114–3140.
-  Sauer, Tim, James A. Yorke, and Martin Casdagli (1991). “Embedology”. In: *Journal of Statistical Physics* 65.3, pp. 579–616.
-  Shah, Rajen D. and Jonas Peters (2020). “The hardness of conditional independence testing and the generalised covariance measure”. In: *Annals of Statistics* 48.3, pp. 1514–1538.
-  Wu, Wei Biao (2005). “Nonlinear system theory: another look at dependence”. In: *Proceedings of the National Academy of Sciences* 102.40, pp. 14150–14154.
-  Wu, Wei Biao and Jan Mielniczuk (2010). “A new look at measuring dependence”. In: *Dependence in Probability and Statistics*, pp. 123–142.

References IV

-  Xie, Min-ge and Kesar Singh (2013). “Confidence distribution, the frequentist distribution estimator of a parameter: A review”. In: *International Statistical Review* 81.1, pp. 3–39.
-  Yorke, James and William Ott (2005). “Prevalence”. In: *Bulletin of the American Mathematical Society* 42.3, pp. 263–290.